


09 | Operating Systems | Multitasking | Processes

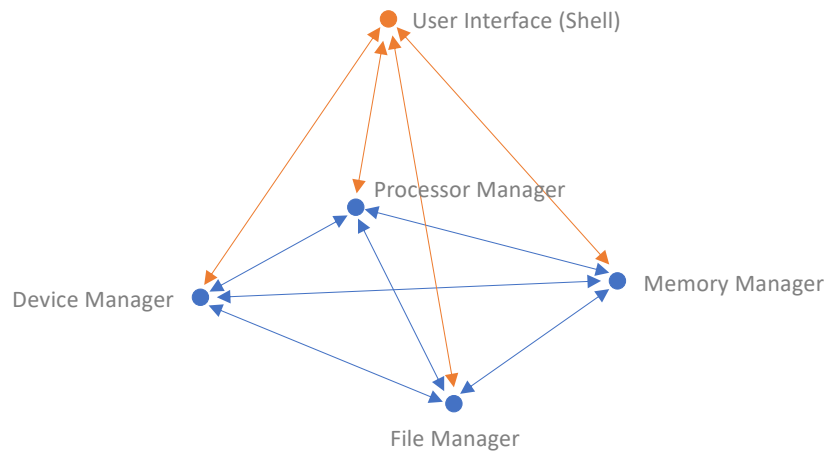
Dr Stuart Thomason

Operating Systems

- Operating systems serve two main purposes
 - Turn hardware components into a usable device
 - Make efficient use of resources (particularly when shared between processes)
- Common general-purpose operating systems...
 - Windows
 - Unix
 - Linux
 - MacOS
 - iOS
 - Android

Based on (or heavily influenced by) Unix
- Embedded operating systems inside home appliances, TV boxes, game consoles, etc.

OS Abstract View



OS Managers

- The base of the pyramid shows the four essential managers of every operating system
 - Processor manager
 - File Manager
 - Device Manager
 - Memory manager
- Each manager must perform certain tasks
 - Continuous monitoring of resources
 - Enforcement of policies (who gets resources, when, and how much)
 - Allocation of resources when appropriate
 - Deallocation of resources when no longer needed
- Managers must work in harmony with each other to complete tasks

Processes and Programs

- There is a difference between processes and programs
- A **program** is the code that performs some task (algorithm, etc.)
 - Source code (Java, C++, etc.)
 - Object code (compiled executable) stored on disk somewhere
 - Program is **static** (doesn't change) after it's compiled
- A **process** is the activity that the CPU performs when it executes a program
 - Code is loaded from disk into memory
 - Instruction pointer starts at first instruction in memory
 - Process is **dynamic** (execution branches depending on input)
- Not necessarily a one-to-one correspondence between programs and processes

Operating System Structure

- An operating system consists of various parts
- A central **kernel**...
 - Resides permanently in memory
 - Performs low-level, frequently needed activities (such as context switching)
 - Operates in **kernel mode** (also known as privileged mode)
- A **shell**...
 - Provides the user interface for the operating system
 - Allows the user to run and interact with programs (processes)
 - Operates in **user mode** (no special privileges, restricted access)
- A set of **processes**...
 - Might be created by the kernel to carry out its activities
 - Or executed when the user runs software
 - Can be either privileged or non-privileged depending on how they were started

System Boot (Initialisation)

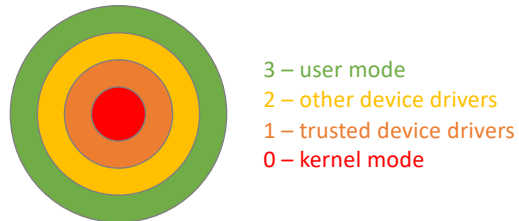
- When the device is first turned on...
 - Interrupt is sent to the CPU
 - Instruction pointer is set to the first address in ROM (read-only memory)
 - Fetch-execute cycle begins from this address
- The ROM contains a small **bootstrap** program
 - Performs basic system checks
 - Sets up system bus and I/O channels
 - Loads OS kernel from disk and passes control to it (via instruction pointer)
- Kernel performs further setup and system checks
 - Starts various processes to perform background tasks of the operating system
 - Starts the main shell process (from which further processes can be triggered)

Command Interpreter (Shell)

- The shell itself is a collection of processes that are spawned by the kernel
 - Allows user to interact with the operating system
 - Processes run in user mode
 - User can run more processes by typing a command or clicking an app icon
- In the early days...
 - Computers were big and expensive (as big as a room)
 - Users connected to a central mainframe via a text-based terminal
 - Much of the ASCII notation dates back to this era (character codes)
- Nowadays...
 - Linux, MacOS and Windows all provide a text-based view (terminal or DOS window)
 - But they also have graphical shells (desktop metaphor)

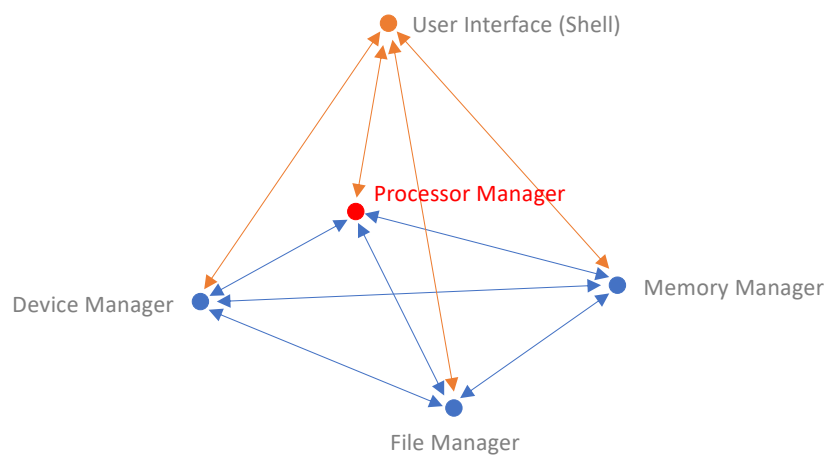
Protection Levels

- The kernel has access to all parts of the CPU, every component, and every I/O device
 - Must be protected from unauthorised access by users (hackers, malware, etc.)
 - Kernel mode is enforced via a **protection ring** (Intel x86 processors have four rings)
 - The CPU flags register stores the **privilege level** of each process



- Certain registers, instructions and memory addresses are protected
 - Can only be used by processes with the correct privilege level
 - Interrupt generated if privilege level doesn't match (**general protection fault**)

Processor Manager

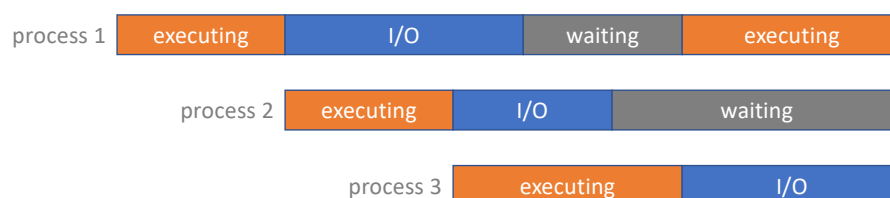


Processor Manager

- Decides how to allocate the CPU to waiting processes
- Driven by a desire to do something useful when a process cannot continue (ie. maximise [throughput](#))
- Processor manager performs various tasks for the operating system...
 - Creates processes when a program is executed
 - Initialises memory and stack for new processes
 - Keeps track of the status of processes
 - Assigns processes to the CPU when available ([context switch](#))
 - Changes process states as events occur
 - Handles termination of processes on completion or abort
 - Handles inter-process communication
 - Manages process queues and prioritisation ([scheduling](#))

Multiprogramming

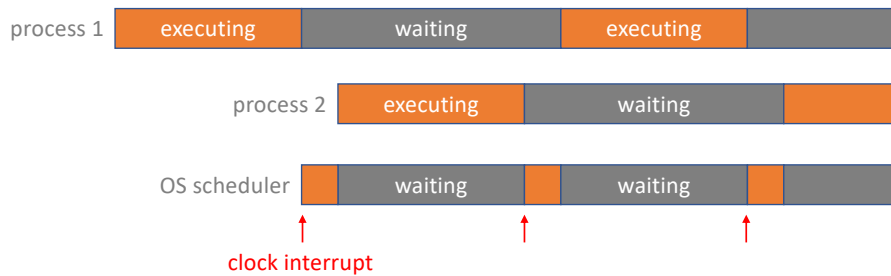
- Early operating systems facilitated multiprogramming
 - Load several processes into memory simultaneously (sharing the CPU)
 - When running process can't continue (eg. waiting for I/O), switch to another
 - Hence I/O and computation can overlap



- This means the CPU is always in use (executing something)
- Problems arise with [compute-bound](#) and [I/O-bound](#) processes

Multitasking (Time-Sharing)

- Extends the concept of multiprogramming by providing fair access to CPU
 - OS switches rapidly between processes to give illusion of uninterrupted execution in parallel ([multitasking](#))
 - Each running process is given a fixed time slice ([quantum](#)) on the CPU



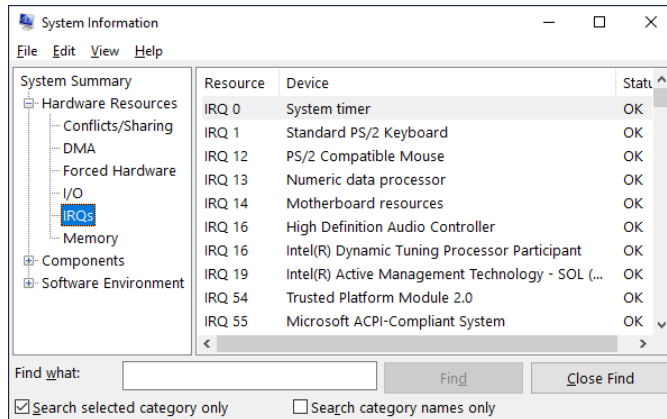
- The OS itself needs time on the CPU to perform its own scheduling tasks

Interrupt Handling

- Multitasking depends on the ability to interrupt the CPU at regular intervals
 - An [interrupt request \(IRQ\)](#) is a hardware signal
 - Usually occurs because something happened outside normal program execution
 - Can happen at any time, regardless of what the CPU is doing
 - Tells CPU to stop current process execution and load an [interrupt handler](#)
- CPU has an [interrupt vector](#)
 - Stores the memory address of handler for each type of interrupt
 - Populated by operating system when it first boots up
 - So the OS is responsible for handling and managing each interrupt
 - Interrupt handlers also known as [interrupt service routines \(ISR\)](#)

Inspecting Interrupts

- View interrupt vector easily in Windows or Linux
 - See the IRQ number and description
 - **IRQ0** is the clock interrupt (quantum tick)



```
[st0034@lxfarm02 ~]$ cat /proc/interrupts
0: IR-IQ-APIC 2-edge timer
1: IR-IQ-APIC 1-edge i8042
8: IR-IQ-APIC 8-edge rtc0
9: IR-IQ-APIC 9-fasteoi acpi
12: IR-IQ-APIC 12-edge i8042
16: IR-IQ-APIC 16-fasteoi ehci_hcd:usb1
18: IR-IQ-APIC 18-fasteoi i801_smbus
23: IR-IQ-APIC 23-fasteoi ehci_hcd:usb2
24: DMAR-MSI 0-edge dmar0
25: DMAR-MSI 1-edge dmar1
26: IR-PCI-MSI 327680-edge xhci_hcd
27: IR-PCI-MSI 409600-edge eno1
28: IR-PCI-MSI 512000-edge ahci[0000:00:1f.2]
29: IR-PCI-MSI 360448-edge mei_me
NMI: Non-maskable interrupts
LOC: Local timer interrupts
SPU: Spurious interrupts
PMI: Performance monitoring interrupts
IWI: IRQ work interrupts
RTR: APIC ICR read retries
RES: Rescheduling interrupts
CAL: Function call interrupts
TLB: TLB shootdowns
TRM: Thermal event interrupts
THR: Threshold APIC interrupts
DFR: Deferred Error APIC interrupts
MCE: Machine check exceptions
MCP: Machine check polls
HYP: Hypervisor callback interrupts
HRE: Hyper-V reenlightenment interrupts
HVS: Hyper-V stimer0 interrupts
ERR:
MIS:
PIN: Posted-interrupt notification event
NPI: Nested posted-interrupt event
PIW: Posted-interrupt wakeup event
```

Context Switch

- The **clock interrupt** is triggered at the end of each time slice (quantum)
 - Operating system runs its **scheduling algorithm** to choose next process
 - This is known as a **context switch**
- During the context switch, the current state of the CPU (registers) is saved in a special data structure called the **process control block (PCB)**
 - Operating system stores a PCB in memory for each process
 - When process is placed onto CPU, the state of its registers is restored from its PCB
- The context switch needs time on the CPU
 - Wastes a few CPU cycles to perform the switch
 - But gives us the benefit of multitasking and maximum CPU usage
 - Without the problems associated with multiprogramming

Process Control Block

- The kernel maintains a PCB for every process
 - Usually stored in memory
 - Can also be represented as a file (see Linux lectures)
- Contains information such as...
 - Unique process ID
 - User ID of the process owner
 - Process state
 - Memory address of process
 - Accounting statistics (time used, etc.)
 - Resources allocated to process (open files, network connections, devices, etc.)
 - Register values from context switch (so CPU can be restored exactly)

Process States

- A process goes through possibly many state changes during its lifetime
- The operating system must keep track of these and update the PCB accordingly
 - **Running** – Currently being executed by the CPU (interrupted at end of quantum)
 - **Ready** – Able to run, but waiting for CPU to become available
 - **Blocked** – Waiting for I/O to complete
- There could be several (or many) processes in each of the ready and blocked states
 - Blocked processes are unavailable for despatch to the CPU
 - Ready processes selected for despatch according to a **scheduling algorithm**
- Processor manager is responsible for creating and terminating processes
 - **Creation** – Reserve memory for the process and its stack, set up process control block, initialise I/O channels, place process into the ready state
 - **Termination** – Close any open I/O channels, remove process control block, deallocate memory

State Changes

