

11 | Priority Scheduling | Round Robin

Dr Stuart Thomason

Process Scheduling


- Recall the two types of scheduling algorithm
 - **Non-preemptive** – Process stays on the CPU until it terminates or blocks for I/O
 - **Preemptive** – Process can be interrupted and replaced with another process
- Remember...
 - The **burst time** is the time a process takes on the CPU before it blocks or terminates
 - The **arrival time** is the time a process is first put into the ready state
- Processes are being created and terminated all the time
 - These lecture examples are simplified
 - Arrival time is not always zero for all processes
 - A newly arrived process might slot into the schedule before an existing process
 - Calculations (eg. average wait time) are constantly changing as processes arrive

Priority Scheduling

- A preemptive algorithm that gives preferential treatment to important processes
 - Each process has a priority assigned to it
 - Highest priority in ready queue gets onto the CPU first
 - Processes with equal priority just use FCFS
 - If a new process arrives with higher priority than current process, interrupt it
- Priorities can be set by the user or assigned by the OS depending on whatever characteristic is most important
 - Memory usage
 - I/O throughput
 - Total CPU time
 - Time already elapsed

Priority Scheduling – Example

- We have five processes with the following burst times and priorities, and they arrive in this order...
 - P1 with burst of 9 ms, priority 3
 - P2 with burst of 2 ms, priority 2
 - P3 with burst of 1 ms, priority 5
 - P4 with burst of 5 ms, priority 3
 - P5 with burst of 6 ms, priority 4
- Higher numbers mean higher priority (zero is least important)



Priority Scheduling – Average Wait Time

- We can work out the average wait time as we did before...

- P1 waited for 7 ms
- P2 waited for 21 ms
- P3 waited for 0 ms
- P4 waited for 16 ms
- P5 waited for 1 ms

- Average wait time = $(7 + 21 + 0 + 16 + 1) / 5 = 9 \text{ ms}$



Priority Scheduling – Advantages & Disadvantages

- Advantages...
 - Smaller wait times and latency for higher priority processes
 - Important processes are dealt with quickly
- Disadvantages...
 - Bigger wait times and latency for lower priority processes
 - Process starvation is still possible
- Starvation can be solved by [process aging](#)
 - Increase priority for processes that have been waiting a long time
 - Eventually their priority will bring them to the front of the queue
- Commonly used in [real-time operating systems](#) where processes must meet deadlines

Round Robin Scheduling


- A preemptive algorithm that gives a set amount of CPU time to each process
- Time slice (**quantum**) is typically between **10 ms** and **100 ms** depending on OS and need
- Similar to FCFS but processes can be interrupted before they terminate or block
- Ready state is treated as a circular FIFO queue
 - New processes arrive at back of queue
 - Scheduler takes next process from front of queue
 - Process runs for a fixed amount of time
 - Context switch moves current process to back of queue and despatches next in line
 - Scheduler keeps going round and round all processes in the ready queue

Round Robin Scheduling

- When a process is placed on the CPU...
 - It will either terminate before the quantum expires
 - Or it will be interrupted (preempted) before it terminates
- If the burst time is less than the quantum time...
 - The process will terminate or block for I/O
 - Scheduler will immediately take next process from queue and place on the CPU
 - See **T** on next slide
- If the burst time is more than the quantum time...
 - Process will be interrupted and a context switch will happen
 - Interrupted process will be moved to the back of the ready queue
 - See **P** on next slide


Round Robin Scheduling – Example

- We have five processes with the following burst times, and they arrive in this order...
 - P1 with a burst of 14 ms
 - P2 with a burst of 2 ms
 - P3 with a burst of 5 ms
 - P4 with a burst of 4 ms
 - P5 with a burst of 9 ms
- The scheduler is using a fixed time quantum of 4 ms




Round Robin – Preemption & Termination

- In the example...
 - P1 had 4 ms on the CPU and was then preempted
 - P2 only needed 2 ms and terminated before the quantum
 - P3 had the next 4 ms and was preempted
 - P4 needed 4 ms so used all the quantum before terminating
 - P5 had the next 4 ms and was preempted
- Then the round robin started again at the front of the queue...



Round Robin – Turn Around & Wait Times

- For each process...
 - Its **turnaround time** is the time when it terminated minus the time when it arrived in the process queue
 - Its **wait time** is its turnaround time minus its burst time
- To simplify the example, assume that all processes arrived at time 0 (in reality processes are arriving all the time)
- Consider P2...
 - It arrived at time 0 and terminated at time 6
 - So its turnaround time was 6 ms
 - Its burst time was 2 ms
 - So its wait time was 4 ms




Round Robin – Average Time Calculations

- Average turnaround time is...

$$(34 + 6 + 23 + 14 + 32) / 5 = 21.8 \text{ ms}$$
- Average wait time is...

$$(20 + 4 + 18 + 10 + 23) / 5 = 15 \text{ ms}$$

Process	Turnaround Time	Wait Time
P1	$34 - 0 = 34$	$34 - 14 = 20$
P2	$6 - 0 = 6$	$6 - 2 = 4$
P3	$23 - 0 = 23$	$23 - 5 = 18$
P4	$14 - 0 = 14$	$14 - 4 = 10$
P5	$32 - 0 = 32$	$32 - 9 = 23$



Round Robin – Advantages & Disadvantages

- Advantages:
 - Easy to implement as a queue
 - Doesn't depend on guessing or estimating burst times
 - Response time based on number of processes instead of process burst times
 - No starvation (every process gets a fair turn on the CPU)
 - Balanced throughput (see below)
- Disadvantage:
 - Depends on selecting a good time quantum
 - Extensive overhead if the quantum is too short
- If quantum is too big, behaves like FCFS (but shorter processes executed faster)
- If quantum is too small, behaves like SJF (but longer processes executed faster)

Windows Priority Scheduling Algorithm

- Windows uses the **Priority Round Robin** algorithm
- Priority values range from 0 to 31
- Every new process is given one of five base priorities
 - **IDLE** (4)
 - **BELOW_NORMAL** (6)
 - **NORMAL** (8) – the default priority
 - **HIGH** (13)
 - **REALTIME** (24)
- Processes given a **priority boost** when they are in foreground (ie. they are the currently active window)
- Process can be boosted when it receives input or an event for which it has been waiting
- Boosted processes are reduced in priority by one level at the end of each time slice (quantum) until base level is reached

Linux Priority Scheduling Algorithm

- Linux also assigns a priority to each process
 - Kernel processes have a value from 1 to 99
 - User processes have a value from 100 to 139
- It also assigns a 'nice' value to each user
 - Ranges from -20 (highest priority) to +19 (lowest priority)
 - Default is 0 for every user
 - Users can change priorities with **nice** and **renice** commands
- Linux uses an algorithm called **Completely Fair Scheduling** (similar to Round Robin)
 - Ready processes are stored in a balanced tree (not a queue)
 - Gives credit for time spent in the blocked state
 - Time quantum varies depending on processor demand
 - Process priority and nice values combine to give overall priority