

Dr Stuart Thomason

- Processes are created (spawned) by other processes
- Original process is the **parent**
- New process is the **child**
- All running processes form a tree structure
- Linux **ps** command
- Everything has **systemd** as the top-level ancestor



Linux Process Creation


- Linux provides several system calls that allow a process to spawn a child
 - `exec()`
 - Allows the process to execute another process
 - Child replaces (overwrites) parent in memory and PCB
 - `fork()`
 - Spawns a new clone of the process
 - Both parent and child continue to run
 - `wait()`
 - Called by parent process
 - Blocks until child process terminates

Process Execution Flow


- P1 uses `exec()` to spawn P2



- P1 uses `fork()` to spawn P2



- P1 uses `fork()` to spawn P2 and then `wait()` to block



fork() – Return Value

- The `fork()` system call returns one of three possible values
 - `< 0` (negative) – If the child could not be created (failure)
 - `= 0` (zero) – In the child process
 - `> 0` (positive) – In the parent process (child process ID)

- Typical fragment of C code to spawn a new process...

```
int pid = fork();
// Call returns in each of the two processes
if(pid == 0) {
    printf("I'm the child process");
    // Will usually call exec() to load its own code
} else {
    printf("I'm the parent and my child's ID is %d", pid);
}
```

The First Process

- The ROM stores a small program that runs a bootloader when a system is first turned on (ie. booted)
 - Linux systems use GRUB – GNU Grand Unified Bootloader
 - Loads kernel image from disk and starts it executing
- The first process to run is called `systemd`
 - Its process ID (`PID`) is `1`
 - Spawns all the other processes required by the kernel
 - Can be configured for various targets (eg. server or desktop)
- Continues to run as a background process (daemon)
 - Offers on-demand spawning of other services
 - Maintains logfiles to record system activity
 - Keeps track of other processes and kernel settings

Shell (Terminal) Login

- The `sshd` daemon runs in the background, waiting for incoming connections (spawned by `systemd` when the system boots up)
 - We use a `ssh` client to connect to a Linux server
 - The `sshd` daemon uses `fork()` to spawn a child process
 - The child uses `exec()` to run a login process
 - The login process checks our credentials are correct
 - Then it uses `exec()` to run our preferred shell process
- Processes are being created all the time to handle whatever the system needs to do, and whatever users are trying to do
- Everything is done via the `fork()` and `exec()` system calls

Running a Shell Command

- When we type a command into the shell...
 - The shell uses `fork()` to create a child process
 - That child uses `exec()` to run the command we typed in
- You can see this with the `ps` command

```
[scap21@lxfarm03 ~]$ ps
  PID TTY          TIME CMD
   730 pts/0        00:00:00 bash
  4737 pts/0        00:00:00 ps
[scap21@lxfarm03 ~]$
```

- Both the shell (`bash`, PID 730) and the `ps` command (PID 4737) are running at the same time

Zombies and Orphans

- Parent processes usually wait for their children to die (!)
- If the death (termination) of a child is not acknowledged by the parent via the `wait()` system call...
 - The child becomes a **zombie**
 - It has finished but is still present in the process table (because clean-up doesn't happen until parent is ready)
- If a parent terminates before its children...
 - The children become **orphans**
 - They are adopted by the `systemd` process
- The `systemd` process periodically calls `wait()` on its children to clean up any zombies and orphans that are in the process table

Daemon Processes

- Already mentioned some daemon processes such as `systemd` and `sshd`, but there are many others
 - Usually the process name ends with 'd' to signify daemon
 - Not associated with any shell or any user
 - Run permanently in the background
- Perform the background operations of the operating system
 - Subsystem managers are daemon processes
 - Need their own time on the CPU, so must be scheduled!
 - Usually run with a higher priority
- Perform tasks requested by other processes (ie. they act as servers in a client-server relationship)

Processes in the Linux File System

- Kernel stores housekeeping information in the `/proc` directory
 - Dynamic details about the current state of the kernel
 - Virtual file system called `procfs` (ie. not all real files on disk)
 - Subdirectory for each running process

Subdirectory	Purpose
<code>/proc/PID/</code>	Stores all the details and status of process <code>PID</code>
<code>/proc/PID/cmdline</code>	The text that was typed to start the process
<code>/proc/PID/fdinfo</code>	The status of any open files used by the process
<code>/proc/PID/status</code>	The overall status of the process (lots of detail)
<code>/proc/cpuinfo</code>	Stores details about the physical CPU(s)
<code>/proc/modules</code>	Stores info about currently loaded kernel modules

(Note: There are many other subdirectories – these are just a few examples)

Looking at Linux Processes

- We can use the `top` command to see dynamic process details

```
top - 09:25:03 up 916 days, 18:14, 2 users, load average: 0.08, 0.06, 0.05
Tasks: 240 total, 1 running, 239 sleeping, 0 stopped, 0 zombie
%Cpu(s): 1.0 us, 0.0 sy, 0.0 ni, 99.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 16304004 total, 1259880 free, 3400064 used, 11644060 buff/cache
KiB Swap: 16777212 total, 10266540 free, 6510672 used. 11910512 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1717	gdm	20	0	1054492	263320	1276	S	7.0	1.6	65753.59	gsd-color
898	scap21	20	0	172960	2484	1636	R	0.3	0.0	0:00.33	top
11123	root	20	0	863580	12420	5512	S	0.3	0.1	25:18.90	docker-containe
15768	root	20	0	153536	19860	8028	S	0.3	0.1	945:52.83	gitlab-runner
31438	root	20	0	1218932	47472	33304	S	0.3	0.3	391:22.48	tail2ban_server
1	root	20	0	128584	7180	4088	S	0.0	0.0	192:31.17	systemd
2	root	20	0	0	0	0	S	0.0	0.0	2:14.13	uthread
3	root	20	0	0	0	0	S	0.0	0.0	3:28.69	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:0H
7	root	rt	0	0	0	0	S	0.0	0.0	1:35.39	migration/0
8	root	20	0	0	0	0	S	0.0	0.0	0:00.02	rcu_bh
9	root	20	0	0	0	0	S	0.0	0.0	443:45.73	rcu_sched
10	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	lru-add-drain

Linux Signals

- A process (running from the shell) can usually be terminated by typing **^C** (**CTRL + C**)
 - This sends an **interrupt** signal to the process
 - Process intercepts signal and responds by terminating
- Signals can be sent between processes with **signal()** system call
- We can send signals with the **kill** command at the shell prompt
 - There are various signals denoted by numbers and codes
 - For example, to terminate process 438

```
kill -s SIGKILL 438
```

Responding to Signals

- The process that receives a signal can respond in three ways
 - Perform the action requested
 - Ignore the signal completely
 - Catch the signal and run some other arbitrary code
- The only signal that cannot be ignored or caught is **SIGKILL** (9)

Code	Number	Meaning
SIGINT	2	Interrupted from keyboard (via CTRL+C)
SIGKILL	9	Request to terminate process (cannot be ignored)
SIGTERM	15	Request to terminate process (might be ignored)
SIGCHLD	17	Indicates that a child process has terminated
SIGIO	29	Indicates that input or output is ready

(Note: There are many other signals – these are just a few examples)

Terminating Zombie Processes

- Imagine a parent process is badly coded...
 - Spawns a child via `fork()`
 - But does not call `wait()` to clean up after its child terminates
- As we saw, the child becomes a zombie
 - Hangs around in the process table doing nothing
 - Will take up some (minimal) system resources
- We can try to send the `SIGCHLD` signal to the parent, but if it's badly coded it will probably just ignore it anyway
- So we will have to send the `SIGKILL` signal to the parent to kill it
 - Parent will terminate and `systemd` will adopt the zombie
 - Zombie will be cleaned up via the periodic call to `wait()`

Inter-Process Communication (IPC)

- Processes very often need to communicate with each other
 - To share, send and receive data
 - To provide services (servers) to other processes (clients)
- Types of IPC
 - Shared memory
 - Shared files
 - Pipes
 - Sockets
- Shared memory and shared files allow two processes to access the same memory location or file at the same time
 - Introduces synchronisation issues
 - Needs to be coordinated by `semaphores` and `locks` (which will be covered in a future bundle of videos)

Pipes

- A **pipe** is a form of IPC between two children of the same parent
- Usually triggered by typing a command at the prompt
 - Join two processes with the **|** (pipe) symbol
 - Output from the first becomes input to the second
- Example...
 - We can list all the kernel modules with **cat /proc/modules**
 - And we can count the lines in a file with **wc --lines filename** (but if we miss out the filename it counts lines of input)

```
[scap21@lxfarm03 ~]$ cat /proc/modules | wc --lines
149
[scap21@lxfarm03 ~]$
```

- So we can find out how many kernel modules are running

Sockets

- A **socket** is a form of IPC that can span multiple systems
 - One process is the **server** (daemon) listening for clients
 - Other process is the **client** that connects to the server
 - Communication is **bidirectional** (ie. both can send/receive)
 - Implemented as special files in the file system
- The processes don't need to be on the same machine
 - Don't need the same parent process (unlike pipes)
 - Typically provide internet services to the rest of the world
 - Server (eg. **httpd**, **maild**, **sshd**) is running as a daemon
 - Clients connect when they need to get/send data
- Server process uses the **listen()** system call to wait for clients
 - When a client connects, usually spawns child via **fork()**
 - Child process handles the communication then terminates

Summary

- Process creation via system calls
- Chain of processes involved in logging in and running commands
- Zombies and orphans
- Daemon processes
- Process housekeeping stored in the file system
- Signals
- Inter-process communication – pipes and sockets