


15 | File Manager | File Allocation | Disk Formats

Dr Stuart Thomason

File Manager




File Manager

- Keeps track of every file in the system
- Provides an abstracted view of files to the user ([logical file system](#)) and maps this on to the real organisation of data on disk ([physical file system](#))
- File manager performs various tasks for the operating system...
 - Organises all files into folders (directories)
 - Manages the locations of files on disk
 - Maps logical structure on to physical locations
 - Enforces restrictions on who can access files ([permission system](#))
 - Deals with standard file operations (open, close, read, write, delete...)
 - Provides standard system calls so other software can use files
 - Liaises with the device manager to access physical disks within the system

Disk Drives

- Mechanical drive ([HDD](#)) – magnetic platter with a read/write head that moves across it
- Solid state drive ([SSD](#)) – uses NAND flash chips to store data via electronic circuits
- The operating system provides an [abstraction layer](#) on top of the physical hardware



tinyurl.com/vhk4xbt5

HDD vs. SSD

- HDDs are mechanical devices with moving parts
 - Slower access speeds (disk needs to spin, head needs to move)
 - Shorter lifespan from wear and tear
 - Uses more power (around 8 watts)
 - Risk from vibration and damage in transit
- SSDs have no moving parts
 - Faster access speeds (electronic circuits via system bus)
 - Longer lifespan
 - More energy efficient (around 2 watts)
 - Impervious to knocks and vibrations
- However, HDDs can store much more data (many multiple terabytes) and are cheaper

Disk Blocks

- Disks are based around the idea of blocks
 - Each disk is split into blocks of a certain size
 - Operating system chooses the block size (typically 4K = 4096 bytes)
 - Files fill up one or more blocks
 - Only one file can exist in each block
 - Space is wasted when a file doesn't fully fill a block
- Each physical disk can be divided into separate partitions
 - Each partition will be formatted with a particular disk format (eg. FAT)
 - This splits the disk (partition) into blocks of equal size
 - Disk format and block usage is determined by the file allocation method

Free List & Metadata

- Parts of each disk are reserved to store data about the disk and the files on it
 - Not all blocks are used for data (so the total storage size is less than the disk size)
 - Some blocks store the [free list](#) and other housekeeping information ([metadata](#))
- The free list is a simple bit vector that records which blocks are free (ie. not used)
 - 1 means in use; 0 means free
 - 8 blocks can be represented by one byte
 - File manager can easily see which blocks are available for new files
 - Modern disk formats use a much more complex set of metadata (taking more room)
- Some disadvantages...
 - Can become out of sync with disk if held in memory (needs to be saved regularly)
 - Could be huge (160GB disk with 4K block size needs 5MB for the free list)

File Systems and System Calls


- Every file system has four parts
 - Physical storage on disk – Formatted as [FAT](#), [XFS](#), [EXT4](#), or one of many others
 - Logical naming scheme – File and directory names (and the way they are nested)
 - Permissions model – Access control such as the permissions used by Linux
 - System calls – Standardised programmatic access to files via the kernel
- System calls include [fopen](#), [fclose](#), [fread](#), [fwrite](#), [fseek](#), [fflush](#)
 - Any program can manipulate data in files through these standard calls
 - The same calls can be used regardless of the physical media or disk format
 - System calls provide an [abstraction](#) of the underlying disk hardware (might not even be a real disk at all - but looks and behaves like one)
- User programs (and code) see the [logical view](#) of the file system via these system calls

Logical and Physical Views

- Operating systems provide conveniences to make life easier for us
 - Filenames with extensions (eg. [.docx](#), [.pdf](#))
 - Folder hierarchies with nested folders and files
 - Shell with drag and drop to move or copy files
 - This is the logical view of a file system
- But the operating system doesn't follow (or need) these human conveniences
 - Directories and files are just bit patterns stored in blocks of the disk
 - OS needs to know which blocks belong to which files
 - Metadata about files and block usage is also stored on the disk
 - Arrangement and usage of blocks depends on the [file allocation method](#)
 - This is the physical view of a file system

Inodes and Metadata

- There must be some way to map the logical view on to physical locations on disk
 - We use [inodes](#) to point to physical blocks on disk
 - Each file has an inode that stores metadata about it
 - Also stores block information (eg. start block) according to allocation method
 - The [inode table](#) needs to be stored on the disk somewhere
 - Each disk format has its own [housekeeping](#) data that needs to be stored
 - The actual capacity of a disk will be reduced by all the space needed for metadata

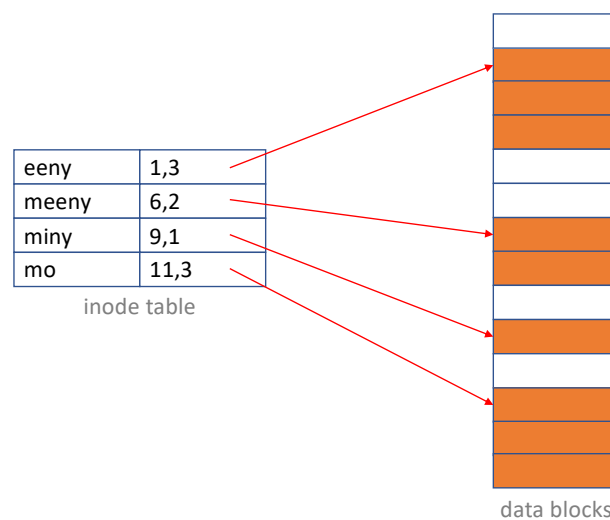


File Allocation and Access

- There are several ways to allocate files to blocks on a disk
 - Contiguous
 - Linked
 - Indexed
- The disk format (file allocation method) dictates how easy it is to access blocks in a file
- Sequential file access...
 - Starts with the first block in a file
 - Reads each block in turn until the required data is found
- Direct file access...
 - Goes directly to the block that contains the required data
 - Skips over earlier blocks in the file

Contiguous Allocation

- Each file is allocated across contiguous blocks on the disk (ie. next to each other)

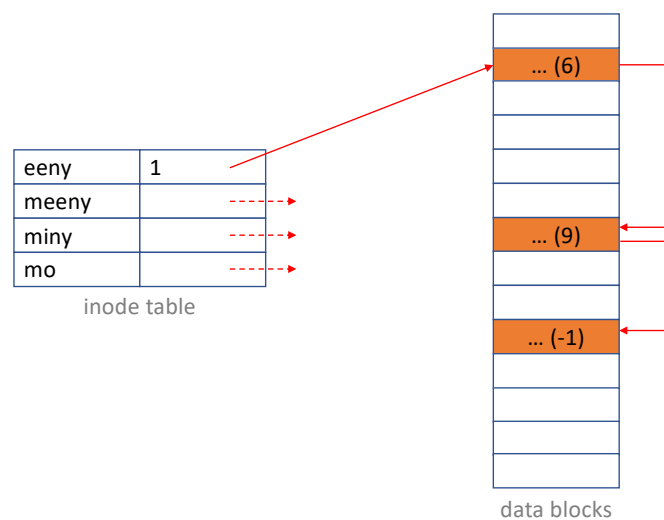


Contiguous Allocation

- The inode stores the start block number and number of blocks
- Advantages...
 - Fast for sequential access
 - Fast for direct access
- Disadvantages...
 - Fragmentation of free blocks
 - Might need regular compaction (defragmentation)
 - Needs policy for choosing which free blocks to allocate
 - Files don't (necessarily) have room to grow after initial allocation

Linked Allocation

- Each block contains a pointer to the next block (similar to a linked list)




Linked Allocation

- The inode stores the start block and each block ends with the number of next block
- Advantages...
 - Easy to grow and shrink files
 - No danger of fragmentation
- Disadvantages...
 - Blocks are widely dispersed (worse in HDDs)
 - Sequential access less efficient
 - Direct access even worse (requires **N** reads to get to block **N** in the chain)
 - Danger of pointer corruption

Indexed Allocation

- First block holds index to all other blocks in a file



Indexed Allocation

- The inode stores number of the index block, which contains a list of all blocks for the file
- Advantages...
 - Each file's block information is held in one place
 - Very efficient for both sequential and direct access
- Disadvantages...
 - Blocks can still be widely dispersed across disk
 - Wastes a block for very small files
 - Can run out of pointers for large files (might need to chain index blocks)
 - Danger of index block becoming corrupted

File Allocation Table (FAT)

- Many Windows installations use the FAT disk format
 - Splits a physical disk into distinct sections
 - Uses indexed allocation but the index blocks are all stored in one place (the FAT)



- The disk is formatted as follows...
 - Boot sector – Contains housekeeping data (and bootstrap code on primary disk)
 - FAT – Master copy of the index blocks for each file (inode table)
 - FAT (copy) – Backup copy of the FAT to help recover from corrupted disks
 - Data region – Contains blocks that store the actual content of files
- Many USB sticks (flash drives) use FAT because it can be understood by all major operating systems (Windows, MacOS, Linux)

File Allocation Table (FAT)

- Indexed allocation method where the inode table is stored in its own part of the disk
- Advantages...
 - All pointer information held in one place
 - Easier to protect
 - No need for a separate free list
 - Direct access much more efficient
- Disadvantages...
 - Requires HDD head to move constantly between FAT area and file area
 - Space for FAT must be reserved when the disk is formatted
 - FAT will become huge for large disks

New Technology File System (NTFS)

- Default disk format for new Windows installations since Windows 3.1 (1993)
- Uses a master file table (MFT) but the general principle is the same as FAT
- Disk is formatted differently to provide many useful new features
 - Access control lists (similar to Linux permissions)
 - Encryption
 - Journaling
 - Hard links
 - File compression
 - System compression
 - Sparse files
 - Shadow copies
 - Disk quotas

Popular File Systems

- Windows...
 - [FAT32](#)
 - [NTFS](#)
- MacOS...
 - [APFS](#) – Apple File System
 - Can also read and write FAT disks
 - Can read NTFS disks but writing to them requires non-trivial system changes
- Linux...
 - Most installations use [EXT4](#) (EXT2 and EXT3 are earlier versions)
 - EXT4 disks can't be read by MacOS or Windows (so use FAT for portable disks)
 - Other popular systems include [btrfs](#) and [ZFS](#)