


21 | Threads | Concurrency | Race Conditions

Dr Stuart Thomason

Threads and Processes

- A thread is like a mini-process within another process
 - Threads are executed independently from (and in parallel with) rest of process
 - Each thread has its **own** instruction pointer and stack space
 - But it **shares** code and data (variables) with other threads in the same process



- In most systems, the kernel is not involved in context switching of individual threads
 - Handled at a higher level within the process that the threads belong to
 - Java threads are handled by the Java virtual machine (JVM) during its own time slice
 - These are known as **user-level threads**
 - Important mechanism for event-driven programs (especially user interfaces)

Threads

- A thread can be seen as a **lightweight** process within a normal **heavyweight** process
- It will usually serve a specific purpose that helps the main process to manage its tasks
- Consider a web browser
 - One thread for retrieving data from internet
 - Another thread to render the page within a window
 - Another thread to wait for keypresses and clicks in the user interface
- Consider a word processor
 - One thread handles the rendering of text on screen
 - Another handles keypresses and mouse clicks
 - Another performs spellchecking in the background
 - Another performs grammar checking

Thread Benefits

- **Ease of programming** – Easier to create one thread for each distinct activity instead of trying to do everything in one linear piece of code
- **Responsiveness** – A multithreaded application can continue to run even if one of the threads is blocked
- **Resource sharing** – Threads can share memory and resources of the process they belong to (but we need to handle concurrency properly)
- **Economy** – Threads are easier to create and switch between because they are allocated and managed within a process that already exists
- **Parallel execution** – Each thread can be physically running at the same time if the CPU has multiple cores (but concurrent programming brings other issues)

Concurrent Programming

- A single thread would perform all the steps in a calculation in a sequence (usually based on the mathematical order of precedence)

$$x = (-b + \sqrt{b^2 - 4ac}) / 2a$$

- In a multi-threaded (or multi-core) system this can be split into two types of operation
 - **Concurrent** operations can be executed in parallel (independently)
 - **Serial** operations rely on the results of earlier operations

```
t1 = -b
t2 = b * b
t3 = 4 * a
t4 = 2 * a
```

Concurrent

```
t5 = t3 * c
t5 = t2 - t5
t5 = √t5
t5 = t1 + t5
x = t5 / t4
```

Serial

- Processing time is reduced from **9 units** to **6 units** (however long an operation takes)
- A good compiler will identify concurrent instructions and assign to different CPU cores

Java Threads

- Each Java program runs in its own heavyweight process
 - The main program code runs in one lightweight thread
 - Java virtual machine (JVM) starts other threads for garbage collection, event handling, screen rendering, and so on
- Programmer can spawn new threads within their code by extending the **Thread** class
`public class Thread extends Object implements Runnable`
- The **Runnable** interface defines a **run()** method that the programmer must implement
 - This code does not start executing until the main program requests it
 - Create an instance of the thread and then call its **start()** method
- Threads are managed internally by the JVM
 - When JVM gets time on the CPU it decides which thread to run (including its own)
 - So it has its own internal thread scheduling algorithm (similar to round robin)

Java Thread Example

- Threads don't start running until explicitly instructed to

```
class MyWorker extends Thread {
    public void run() {
        System.out.println("I am a worker thread");
    }
}

public class MyMain {
    public static void main (String[] args) {
        MyWorker runner = new MyWorker();
        runner.start();
        System.out.println("I am the main thread");
    }
}
```


- Think about which order the output will appear in – which line is printed first?

Thread Start

- The `MyWorker` class extends the `Thread` class
 - So the programmer must implement the `Runnable` interface
 - Define a `run()` method and put the code for the thread inside it
- In the main program...
 - Create a new object of type `MyWorker`
 - Call its `start()` method
- Within the JVM...
 - Memory is allocated for the new thread object
 - Its `run()` method is called and begins execution
- Now both the main program thread and the child thread are running in parallel

Java Thread States

- Each thread begins in the **new** state when its object is allocated memory
 - It moves to the **runnable** state when its **start()** method is called (but it might not run immediately depending on the JVM thread scheduler)
 - It moves to the **blocked** state when its **sleep()** method is called, or when it's waiting for I/O, or if it's waiting for another thread to finish
 - It moves to the **dead** state when it has terminated and is waiting for the JVM to clean up its memory via garbage collection



Synchronisation Problems

- Consider this thread object with a constructor that sets two internal character variables

```
class TwoChar extends Thread {  
    private char out1, out2;  
  
    public TwoChar(char first, char second) {  
        out1 = first;  
        out2 = second;  
    }  
  
    public void run() {  
        System.out.print(out1);  
        System.out.print(out2);  
    }  
}
```

- When the thread runs, it will output each character one after the other

Synchronisation Problems

- We can create and start multiple threads

```
public class ThreadExample {  
    public static void main(String[] args) {  
        TwoChar tc1 = new TwoChar('A', 'B');  
        TwoChar tc2 = new TwoChar('1', '2');  
  
        tc1.start();  
        tc2.start();  
    }  
}
```

- We will get different results when we run this, depending on JVM thread scheduler

AB12 A1B2 12AB 1A2B 1AB2 A12B

- We would probably see AB12 most often
 - The order that they appear in the program code
 - But there is **no guarantee** about the execution order of threads

Shared Variables

- Suppose we have an object that can be shared by multiple threads

```
class Something {  
    private int num = 0;  
    public void increase() {  
        num = num + 1;  
    }  
}
```

- This object is passed into the constructor of two threads

```
Something thing = new Something();  
MyThread t1 = new MyThread(thing);  
MyThread t2 = new MyThread(thing);
```

- At some point during the program execution, both threads access the shared object
`thing.increase();`
- We **cannot guarantee** that both increases will happen (due to thread scheduling)

Race Conditions

- Consider what happens at the register (assembly) level when the increase happens

```
mov eax, num
inc eax
mov num, eax
```
- Works perfectly fine for most of the time, but occasionally the instructions are executed in a way that causes a problem
- If the starting value is 0, we would expect it to be 2 after both threads have executed

T1	T2	NUM
MOV (eax = 0)		0
INC (eax = 1)		0
	MOV (eax = 0)	0
	INC (eax = 1)	0
	MOV (eax = 1)	1
MOV (eax = 1)		1

- This is known as a **race condition** and is very hard to debug or even notice it happened

Critical Regions

- A race condition could happen whenever a variable is shared between multiple threads
 - Most of the time it won't happen because the threads are not interrupted
 - But we can't ignore the problem because sometimes it will happen
- Whenever a process or thread is accessing a shared resource, its code is said to be in a **critical region** (or critical section)
- We must ensure only one thread (or process) can access the resource at any one time
 - The programmer must somehow 'lock' the resource so it gets exclusive access
 - Other threads must wait for the resource to be unlocked

Semaphores

- The concept of semaphores (or tokens) comes from single track railways in the 1800s
 - A train can only enter a single track if the driver picks up a physical token
 - Only one token exists, so only one train can be on the track
 - Driver leaves token at the other end when they have passed the single track
- Programming semaphores were proposed by Dijkstra in the 1960s
 - Critical region of code is like a single track railway
 - Thread must acquire the semaphore before it can enter the region
 - Other threads must wait (block) until the semaphore becomes available
- Use of semaphores requires careful programming
 - Must consider how each shared resource will be used by each thread
 - Critical regions must be kept as small as possible to reduce thread blocking