

## 22 | Producer-Consumer Problem | Semaphores

Dr Stuart Thomason

### Semaphore Notation

- The semaphore mechanism provides two operations
  - `wait()` and `signal()`
  - Sometimes called `P` and `V`, `test` and `inc`, or `up` and `down`
  - We will implement our own `boolean semaphore` over the course of this lecture
  - Java uses `different names` for its semaphore methods but they do the same thing
- A thread must call `wait()` before it enters a critical region
  - This will `block` if the region (semaphore) is in use by another thread
  - Otherwise the thread obtains the lock and can enter the critical region
- The thread must call `signal()` just before it exits the critical region
  - This releases (unlocks) the semaphore
  - Any other blocked threads will be notified that the region is now safe to enter (only one thread will be successful if there are multiple threads waiting)

## Producer-Consumer Problem

- The producer-consumer problem is a good way to illustrate how semaphores work
  - The producer and consumer are separate threads with different roles
  - Both the producer and consumer require access to a shared resource
  - There could be multiple producers and/or consumers accessing the same resource
- For example...
  - A secretary writes a letter (producer)
  - They put the letter into a tray (shared resource)
  - A manager takes the letter from the tray to sign it (consumer)
- In programming terms...
  - Threads are usually consumers or producers of some resource
  - The resource is a shared variable or object represented by a class (in Java)

## Buffer – First Attempt

- We can model the problem and solution as two threads trying to access a buffer
  - Producer cannot insert if the buffer is full
  - Consumer cannot remove if the buffer is empty
  - Buffer cannot be accessed by multiple producers and consumers at the same time
- Consider the code for the buffer object

```
class Buffer {  
    private int store;  
    public void insert(int item) {  
        store = item;  
    }  
    public int remove() {  
        return store;  
    }  
}
```

## Producer and Consumer

- Consider the code for the producer object

```
class Producer extends Thread {  
    private Buffer buff;  
    public Producer(Buffer b) {  
        buff = b;  
    }  
    public void run() {  
        int m;  
        ...  
        buff.insert(m);  
        ...  
    }  
}
```

- The consumer object looks the same but it removes from the buffer instead  
`n = buff.remove();`

## Main Program

- The code for the main program creates instances of the buffer, producer and consumer

```
public class ProdCon {  
    public static void main(String[] args) {  
        Buffer b = new Buffer();  
        Producer p = new Producer(b);  
        Consumer c = new Consumer(b);  
        p.start();  
        c.start();  
        ...  
    }  
}
```

- None of this code has protection for the constraints mentioned earlier...
  - Producer doesn't check to see if the buffer is full
  - Consumer doesn't check to see if the buffer is empty
  - Buffer is not protected from being used by both threads at the same time

## Buffer – Second Attempt

- Prevent threads trying to insert if the buffer is full (or remove if the buffer is empty)

```
class Buffer {
    private int store;
    private volatile boolean empty = true;
    public void insert(int item) {
        while(!empty);
        store = item;
        empty = false;
    }
    public int remove() {
        while(empty);
        empty = true;
        return store;
    }
}
```

- The **volatile** keyword prevents variables from being cached (always uses latest value)

## Spinlock

- Before inserting anything, the producer waits for the buffer to be empty

```
while(!empty);
empty = false;
```

- Before removing anything, the consumer waits for the buffer to be full

```
while(empty);
empty = true;
```

- At first it might seem like these loops will never terminate
  - But the status variable is shared between **multiple** threads
  - Another thread could change its value and then the loop will terminate
- This concept is known as **spinlock** or **busy waiting**
  - It works well, but it is inefficient because it wastes CPU cycles
  - Should only be used when the expected wait time is very short

## Yielding

- It could be more efficient for the waiting thread to give up control of the CPU

```
public void insert(int item) {
    while(!empty) {
        Thread.yield();
    }
    store = item;
    empty = false;
}
```

- The `yield()` method politely tells the JVM that the thread is willing to block (give up CPU)
  - Behaviour depends on JVM implementation
  - The JVM might ignore the request, so we still need to put it within a spinlock
  - The request is **non-deterministic** (might behave differently on different occasions)

## Buffer – Third Attempt

- Prevent multiple threads from trying to insert and remove at the same time

```
class Buffer {
    private int store;
    private volatile boolean empty = true;
    public synchronized void insert(int item) {
        while(!empty) { Thread.yield(); }
        store = item;
        empty = false;
    }
    public synchronized int remove() {
        while(empty) { Thread.yield(); }
        empty = true;
        return store;
    }
}
```

- Note the American spelling of the **synchronized** keyword (with a zed)

## Java Object Locks

- Every Java object has a **lock** associated with it
  - Locks are not shared across multiple instances of the same class
  - Each object instance has its own lock that works independently of other instances
- The **synchronized** keyword declares methods that contain **critical regions**
  - When a thread calls a synchronized method, it gets exclusive control of the lock
  - Other calls to the object's synchronized methods must wait (block) in an **entry set**
  - Note that non-synchronized methods can still be called without any issues
- When the thread with the lock exits the synchronized method, the lock is released
  - JVM will select an arbitrary thread from the entry set
  - That thread will be able to call the synchronized method it was waiting for

## Wait and Notify


- Java provides methods to implement semaphores (but it uses **notify** instead of **signal**)

```
public synchronized void insert(int item) {  
    while(!empty) {  
        try {  
            wait();  
        } catch (InterruptedException e) {}  
    }  
    store = item;  
    empty = false;  
    notify();  
}
```

- These methods will always work regardless of the JVM or platform
- Note that exception handling is outside the scope of this module
  - It has to be included so the code will compile
  - But we are not interested in what it does (see **COMP122** or online tutorials)

## Entry and Wait Sets

- The **wait()** call suspends current thread and moves it to the **wait set**
- The **notify()** call moves an arbitrary thread from the wait set to the **entry set** (the choice of thread depends on the JVM implementation)



- Can also call **notifyAll()** to allow all threads in the wait set to compete for the chance to resume (winner is the one that calls a synchronized method first)

## Java Semaphore Class

- The previous examples showed how we can implement our own semaphores
- Java provides a Semaphore class that we could use instead

```
import java.util.concurrent.*;
```
- Specify how many **permits** the semaphore allows (**counting semaphore**)

```
Semaphore sem = new Semaphore(1);
```
- When the **acquire()** method is called...
  - If the counter is **0**, the thread will block (added to a **wait queue**)
  - If the counter is greater than **0**, counter is decreased and the thread can continue
- When the **release()** method is called...
  - The counter is increased
  - The first blocked thread in the wait queue can acquire the permit
- Testing and setting the counter must be **atomic** operations to avoid race conditions

## Buffer – Semaphore Version

- Update the code to use the Java semaphore mechanism

```
class Buffer {
    private int store;
    private Semaphore sem = new Semaphore(1);
    public synchronized void insert(int item) {
        try { sem.acquire(); } catch (InterruptedException e) {}
        store = item;
        sem.release();
    }
    public synchronized int remove() {
        try { sem.acquire(); } catch (InterruptedException e) {}
        int item = store;
        sem.release();
        return item;
    }
}
```