

23 | Dining Philosophers Problem | Deadlock


Dr Stuart Thomason

Dining Philosophers Problem

- In the Producer-Consumer problem...
 - Threads have **distinct** roles (eg. manager and secretary)
 - The synchronisation environment involves a shared resource (eg. letter tray)
- In the Dining Philosophers problem...
 - Threads have **identical** roles
 - The synchronisation environment involves a shared resource
- This problem helps to visualise two common issues with concurrent systems
 - Deadlock
 - Starvation
- Operating system processes and resources are also affected by these issues


Dining Philosophers Problem

- There are **N** philosophers seated around a table
 - Each philosopher flips between eating and thinking (eat – think – eat – think – etc.)
 - Philosophers don't need any resources to think
 - But they need two chopsticks to eat
- There are only **N** chopsticks (one between each philosopher)
- If a philosopher wants to eat, they must wait for both neighbours to stop eating



Solution Approach

- Each philosopher has a unique index value in the set {0, 1, 2, 3, 4}
- And each chopstick has an index from the set {0, 1, 2, 3, 4}
- The philosopher with index **K** must wait for both chopsticks **K** and $(K + 1) \% 5$
- When a philosopher eats, they block two other philosophers (their neighbours)
 - When **P0** is eating, they are using sticks **S0** and **S1**
 - Which means **P1** and **P4** cannot eat until **P0** has stopped eating



Java Semaphore Solution

- We treat each chopstick as a shared resource with its own critical region

```
class Chopstick {
    private volatile boolean using = false;

    public synchronized void pickUp() {
        while(using) {
            try {
                wait();
            } catch (InterruptedException e) {}
        }
        using = true;
    }
    public synchronized void putDown() {
        using = false;
        notify();
    }
}
```

Java Philosopher Thread

- Each philosopher is a thread with a `run()` method (note that this is incomplete code)

```
public void run() {
    while(alive) {
        think();
        stick[i].pickUp();
        stick[(i + 1) % 5].pickUp(); } non-atomic statements
        eat();
        stick[i].putDown();
        stick[(i + 1) % 5].putDown();
    }
}
```

- Each chopstick has a semaphore
 - Two philosophers cannot pick it up simultaneously
 - If a philosopher needs the stick but it's already in use, philosopher moves to wait set
 - When chopstick is put down, waiting threads are notified and it can be picked up

Starvation

- We cannot guarantee...
 - The order in which threads will be executed
 - When threads will be interrupted during execution
- Some possible scenarios...
 - If P0 and P2 pick up first, the others (P1, P3, P4) must wait and cannot eat
 - If P0 and P2 put down and immediately pick up again, the others cannot eat
 - If P0 and P2 take turns to eat, P1 still cannot eat
- Without proper controls on when threads run, some philosophers could starve
- In general, starvation occurs when one or more of the threads (or processes) in a concurrent system is denied access to the resources it needs

Deadlock

- Suppose each philosopher picks up the chopstick to their right simultaneously
 - Nobody can pick up the chopstick to their left
 - All philosophers move into the waiting state
 - Nobody ever releases their chopsticks, so all philosophers wait forever
- This is known as **deadlock** (or **deadly embrace**)
- In general, deadlock occurs in a concurrent system when each thread (or process) is waiting for one of the others to do something (or release something)

“When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone”

Illogical law passed by the state of Kansas

Deadlock Solutions

- We could add some rules to the code to prevent deadlock
 - Allow only $N - 1$ philosophers to dine at the same time
 - Introduce asymmetry (so even indices pick up right chopstick first; and vice versa)
 - Insist that both chopsticks are picked up in one atomic operation (uninterruptable)
- All these solutions add extra complication to the code
- The dining philosophers problem helps us to understand the issues that can occur at the operating system level with processes and resources
 - Processes are the philosophers
 - Resources are the chopsticks

Resource Allocation & Deadlock


- The operating system must allocate and share resources sensibly
 - CPUs
 - Devices
 - Memory
 - Files
- In its simplest form, deadlock will occur in the following situation...
 - Process A is granted resource X and then requests resource Y
 - Process B is granted resource Y and then requests resource X
 - Both resources are non-shareable
 - Both resources are non-preemptible (ie. cannot be taken away from the process)
- Operating system can either prevent, detect, avoid or ignore deadlock

Deadlock Prevention

- We could force processes to claim all resources in one atomic operation
 - When process starts, immediately claim all resources it will need
 - Some resources will be under-utilised if not needed right away
 - Process might not know what future resources it will need as it executes
- We could number each resource and force processes to claim them in a specific order
 - Always request resource X before claiming resource Y
 - This avoids processes waiting for each other
- We could force processes to use only one resource at a time
 - Must release resource X before trying to claim resource Y
 - Prevents any process from needlessly hogging a resource
 - But this is very unrealistic in any modern operating system


Resource Allocation Graphs

- We can draw a graph that models the processes and resources in a system
 - Processes represented as circles
 - Resource types represented as rectangles with resource instances as dots
 - Directed edges to show resource requests and assignments (arrow changes direction when a request is satisfied)



Resource Allocation Graphs

- Suppose P3 requests access to R2
 - This introduces a cycle (in fact, two cycles)
 - Indicates that deadlock is possible
 - But does not necessarily mean there **is** or **will be** deadlock (it will only happen if the request is granted at the 'wrong' time)



Deadlock Detection and Recovery

- Resource allocation graph is monitored by the operating system
 - Very costly to do for every request
 - Might only check the graph every so often, such as when CPU usage deteriorates
- If a cycle is detected, the operating system could...
 - Take an existing resource away from a process to remove the cycle (preemption)
 - Kill one or more processes (a crude solution but it solves the problem)
 - Roll back to a previous state (requires OS to store checkpoints of process states)
- Each solution involves losing some processing time and even losing some data
 - Rollbacks will involve running the previous instructions again in a different order
 - Lots of memory will be used to store checkpoints
 - Calculations (including memory and disk writes) would need to be undone

Deadlock Avoidance

- Deadlock avoidance is very similar to deadlock prevention
 - Never grant access to a resource if it would lead to deadlock
 - Processes must declare up front what resources they will be using (might not know)
 - Processes will be forced to wait unnecessarily (because deadlock is rare)
- The **banker's algorithm** is a theoretical approach to deadlock avoidance
 - Devised by Dijkstra
 - Needs to know in advance what resources a process will request
 - Determines whether a static collection of processes is in a **safe** or **unsafe** state
 - Guarantees that processes will terminate but doesn't care how long it takes
- In practice, this algorithm is unsuitable because new processes arrive all the time in a dynamic system, and we cannot force processes to wait for an indefinite time

Ignoring Deadlock

- It is practically impossible to implement efficient deadlock detection or prevention
- The **ostrich algorithm** is a common approach in many modern systems
 - (The bird that sticks its head in the sand to pretend there are no problems)
 - More **cost-effective** to allow the problem to occur than try to prevent or avoid it
 - Detecting and preventing deadlock is expensive and adds complexity to code
 - In all practical situations, deadlock is **relatively rare**
 - The easiest way to recover from deadlock is to reboot the system when it happens
 - Rebooting every so often is cheaper than trying to prevent deadlock
 - This trades **correctness** for **convenience**
- Windows and Linux ignore deadlock and don't do anything to detect or prevent it
- The JVM also ignores deadlock when scheduling individual Java threads