

Distributed Systems

COMP 212

Lecture 6

Othon Michail

Leader Election

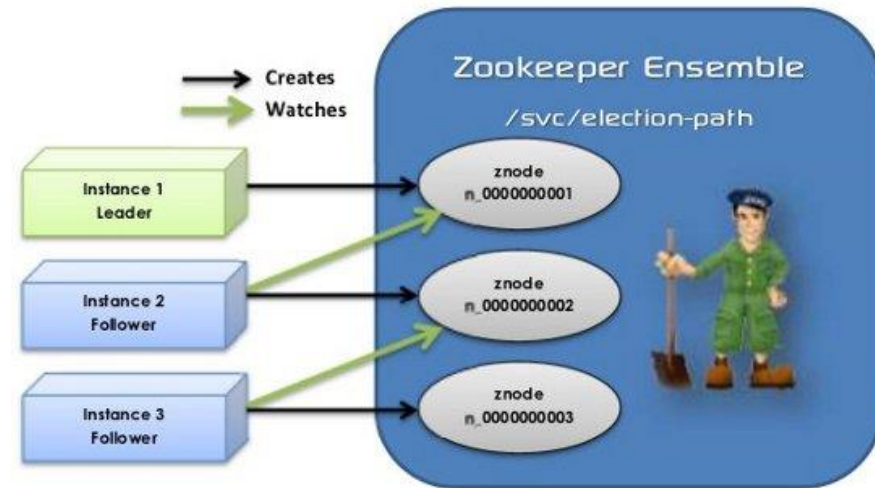
Election for Coordination

- Many Distributed Systems require a **single** process to act as the **coordinator** (for various reasons)
 - Time server
 - Coordinator in commit protocols
 - Master process in distributed computations
 - Master database server
- Coordinator may fail → the distributed group of processes must execute an **election algorithm** to determine a **new** coordinator process

Applications

- **Zookeeper**

- distributed, open-source **coordination service** for distributed applications
- provides a centralized infrastructure and services that enable synchronization across an Apache Hadoop cluster
- uses a **leader server**
- **elects** a new one if needed
- Hadoop: processing of datasets of big data using the MapReduce programming model



- **Chubby lock service**

- **Google File System** and **MapReduce** use it to **elect a master**

Problem Statement


- *Elect a **unique leader processor** from among all the processors in the distributed system*
- Leader to be interpreted as:
 - **coordinator**
 - **master processor**
- Special case of **consensus/agreement**
- Processors should agree eventually on who they elect

Variants of Leader Election

- General network or special type of network, e.g., a ring
- Processors can be **identical** or have pre-assigned **unique ids**
- All processors may be required to
 - **know** the **elected** processor
 - to **output** the fact that themselves were **not elected**
 - to **terminate**
- Processors may possess in advance some **information about the network**
 - e.g., the number of processors **n** in the network

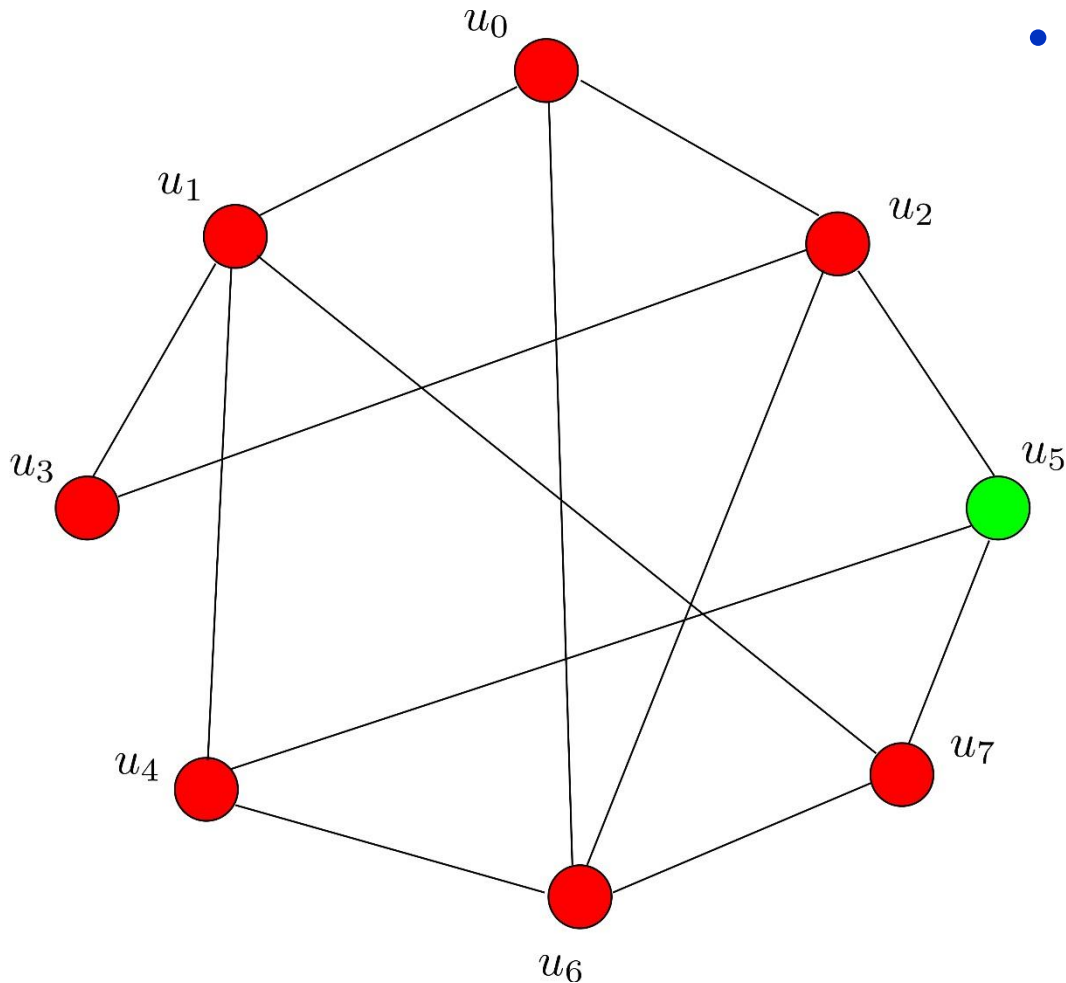
Important remark: Even slightly different assumptions can completely change the algorithmic solutions required or even make a problem impossible to solve

An Illustration



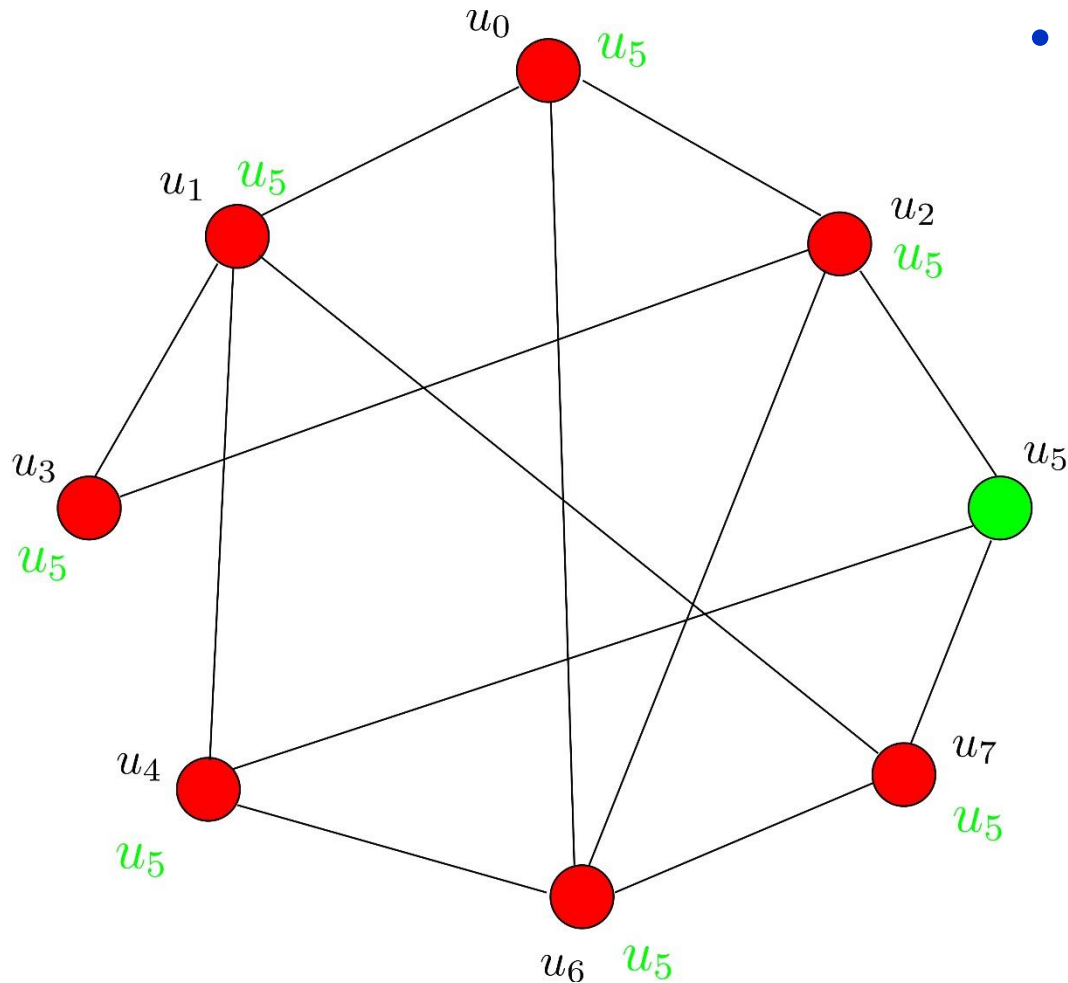
- All processors are initially non-leaders
 - no processor elected yet

An Illustration



- Eventually a **unique leader** will be **elected**
 - node u_5 in this case


An Illustration



- Possibly all nodes could eventually learn who is the elected leader

Leader Election in a Ring

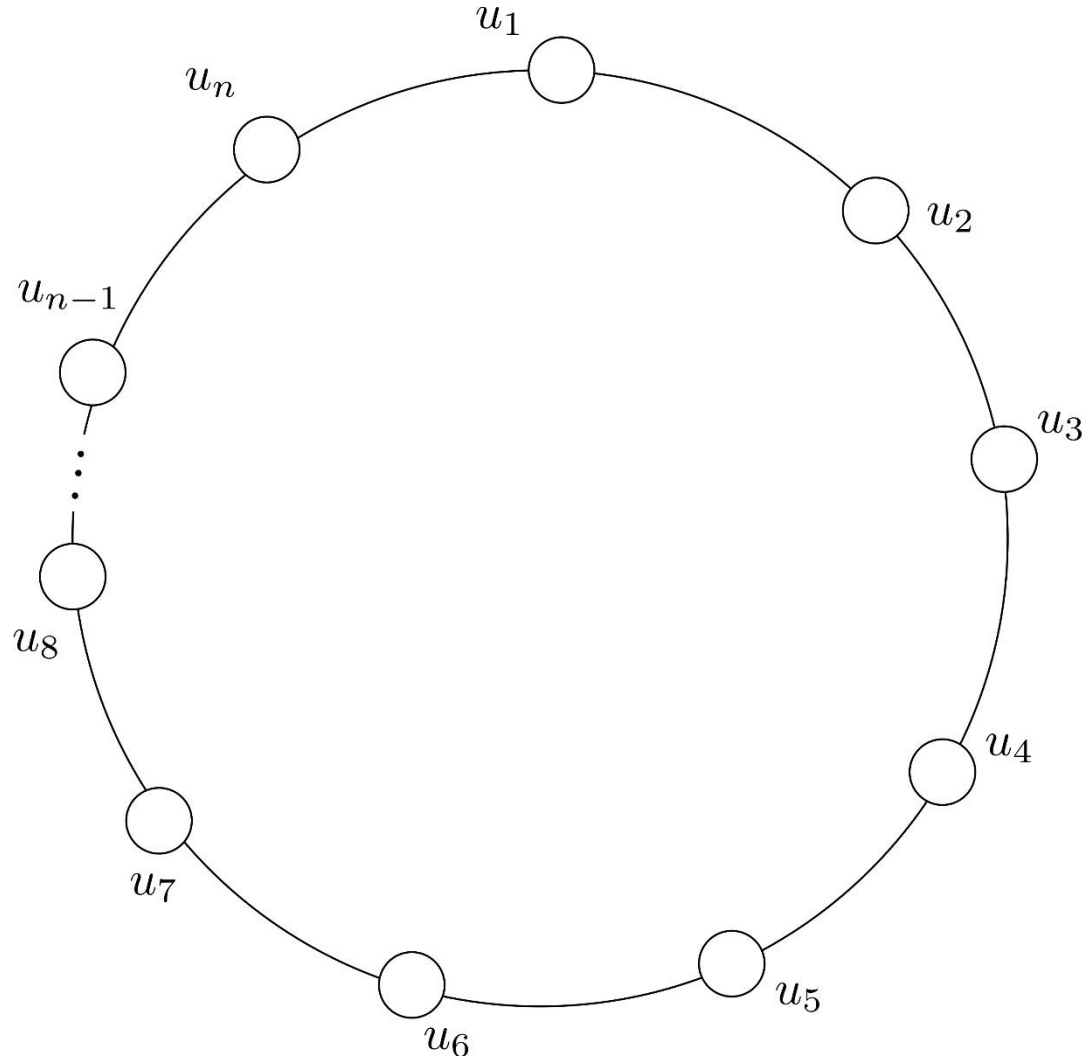
- First appeared in local area **token ring** networks
- Single token circulates the network
- The current owner of the token has sole right to initiate computation
 - e.g., send jobs to a printer in the network



- If 2 or more nodes were to attempt simultaneously to communicate they would **interfere**
- But **occasionally the token is lost**
- Processors must execute an algorithm to **regenerate the lost token**
 - **Equivalent to leader election**


The Ring Network

- Bidirectional (or undirected)




The Ring Network

- Unidirectional (or directed)



A First Minimal Setting

- Directed ring
- All processors are initially identical
 - Meaning here that they all start from exactly the same initial state




A First Minimal Setting

- Directed ring
- All processors are initially identical
 - Meaning here that they all start from exactly the same initial state

Question:

Is there an algorithm that solves leader election?



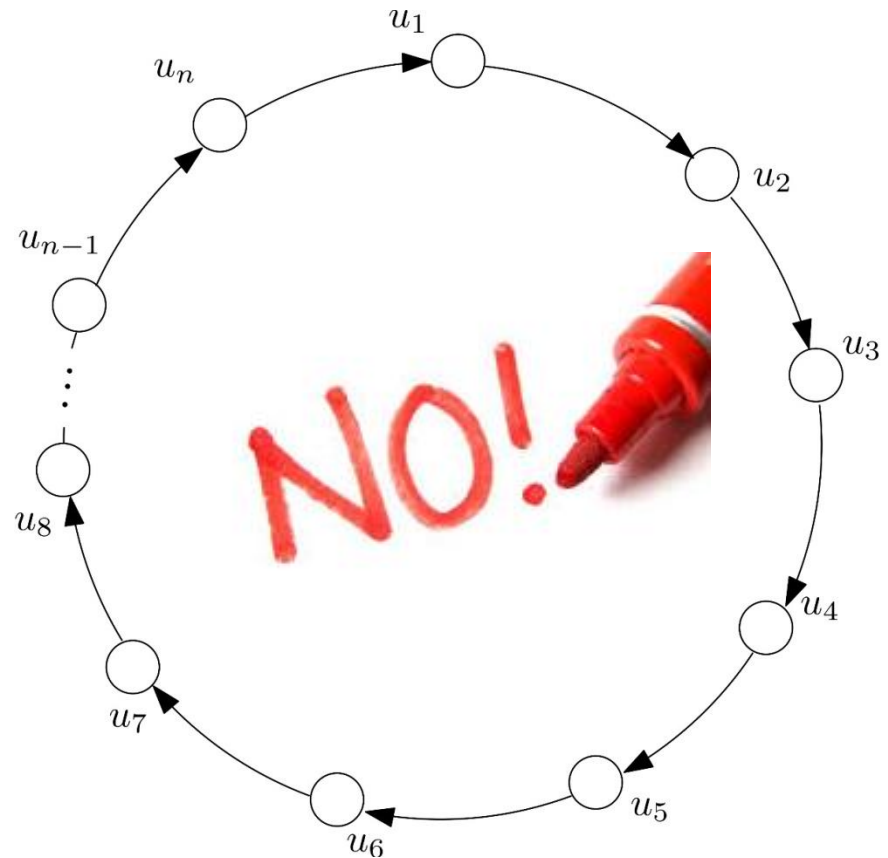
Our First Impossibility Result

Question: *Is there an algorithm that solves leader election?*

Answer: No

Very general:

No matter which algorithm you try, it will fail!



Impossibility of Leader Election with Identical Processors

Theorem. Let G be a **directed ring** of n processors. Take any algorithm A (to be executed in G), in which all processors are initially in exactly the same state. Then A **does not** solve the leader election problem.

Proof Idea.

Observation 1: To solve leader election, at some point **one processor must be in a different state than the rest**

Impossibility of Leader Election with Identical Processors

Proof Idea (continued).

Observation 2: If all processors are identical, given that

- (i) they also have identical neighbourhoods and
- (ii) they operate synchronously

we expect them all to do identical things

Summing-up: If we prove that **all processors must forever remain identical** (no matter which algorithm they execute), then we can conclude that they **cannot elect a leader**

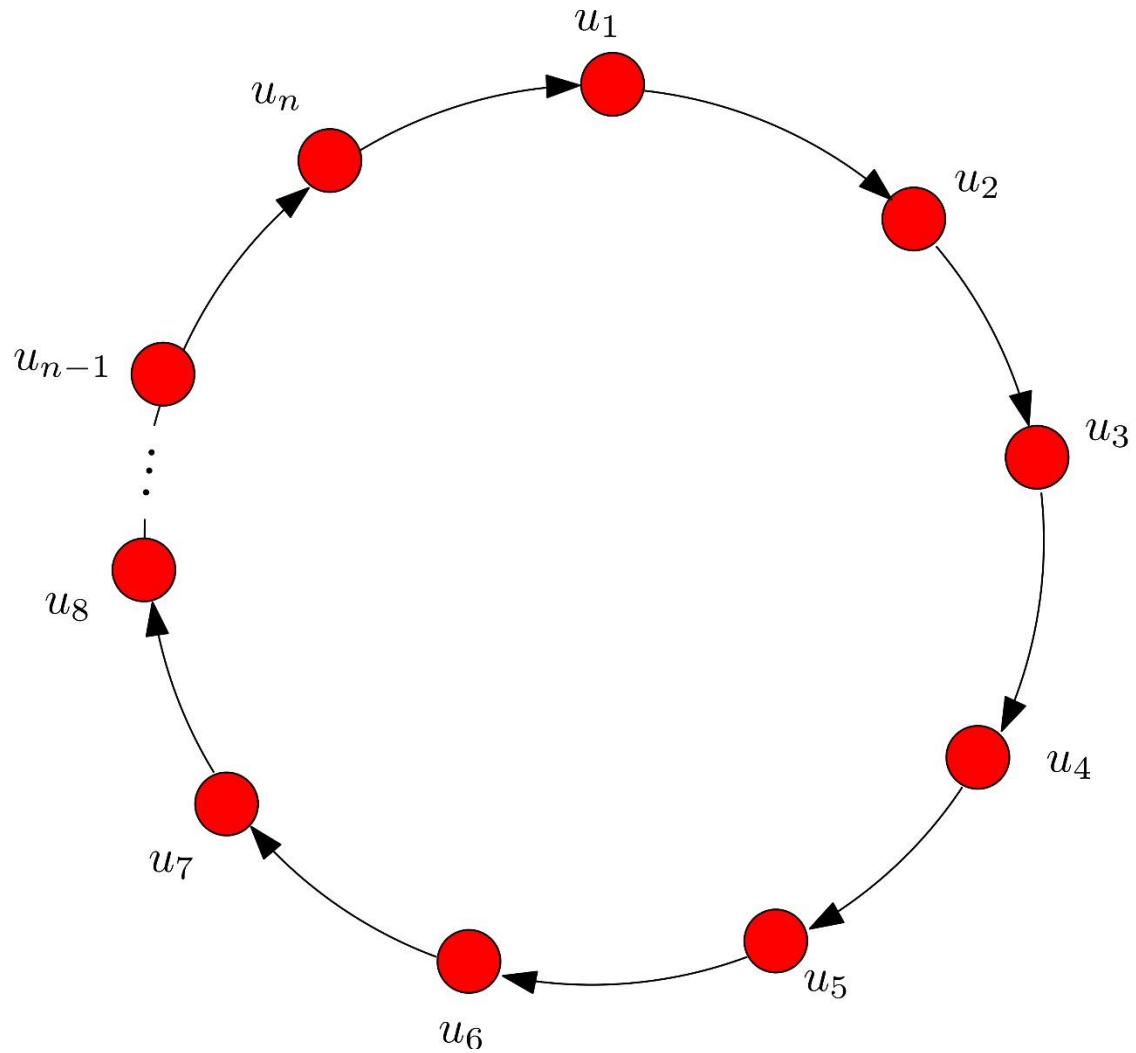
- Simply because **a leader should be non-identical to the rest**

Impossibility of Leader Election with Identical Processors


Proof. Take any algorithm A . We will show that in every round r all processors are in identical states.

- By **induction on r**
- **$r = 0$** : We know that in our setting all processors start from identical initial states
- Assume it holds for r
- **$r + 1$** : By assumption we know that it holds for r
 - That is, in round r all processors are in identical states
 - This means that they **all produce the same message** to be forwarded to their right neighbour (because the messages produced by a processor depend only on its state)
 - Thus, in $r + 1$ all processors are in identical states and receive the same incoming message
 - Therefore, they will perform the same state-update and will again obtain identical states in $r + 1$ (because the new state of a processor depends only on its previous state and the messages received)

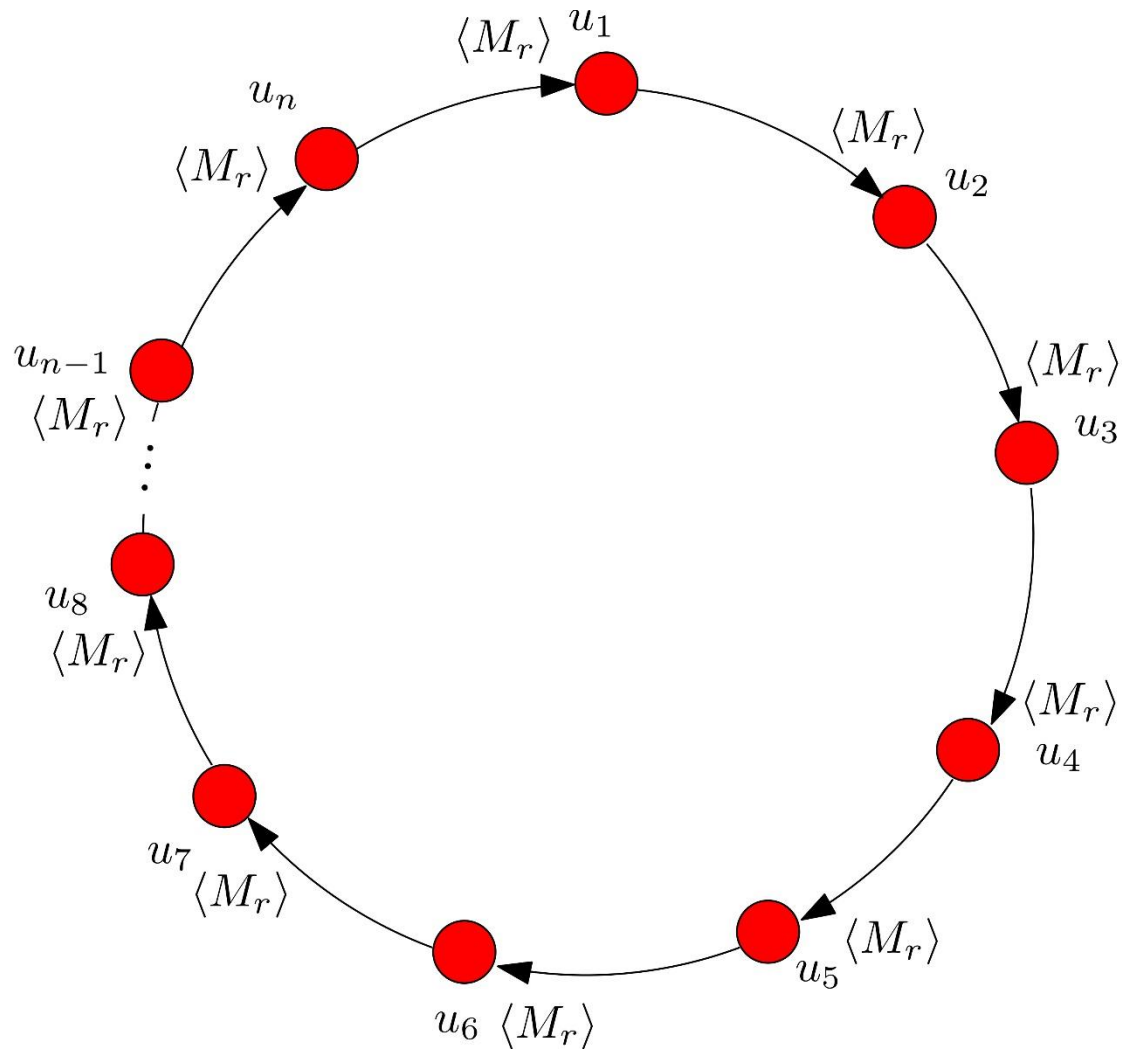
An Illustration




An Illustration



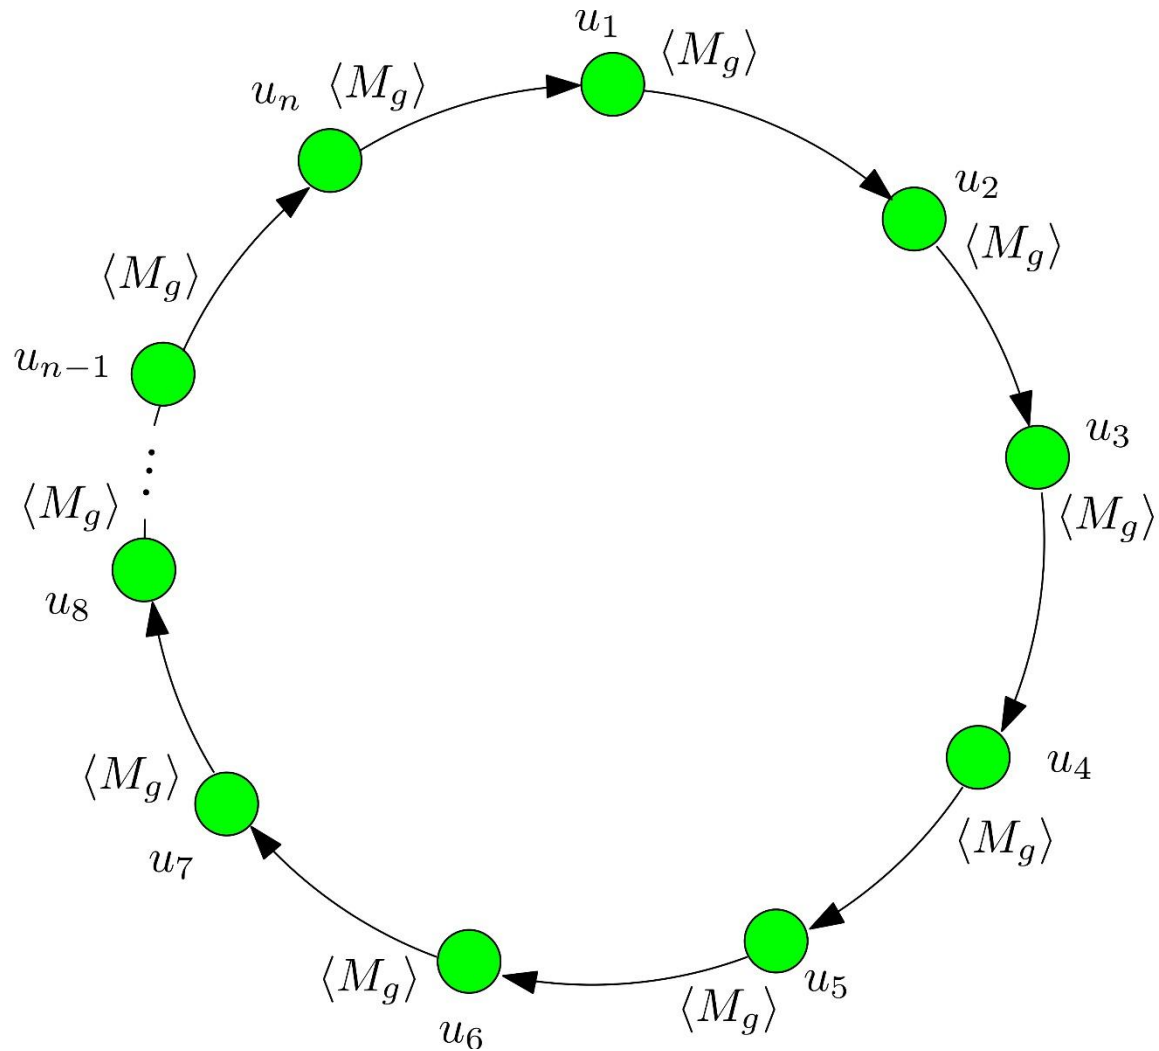
An Illustration



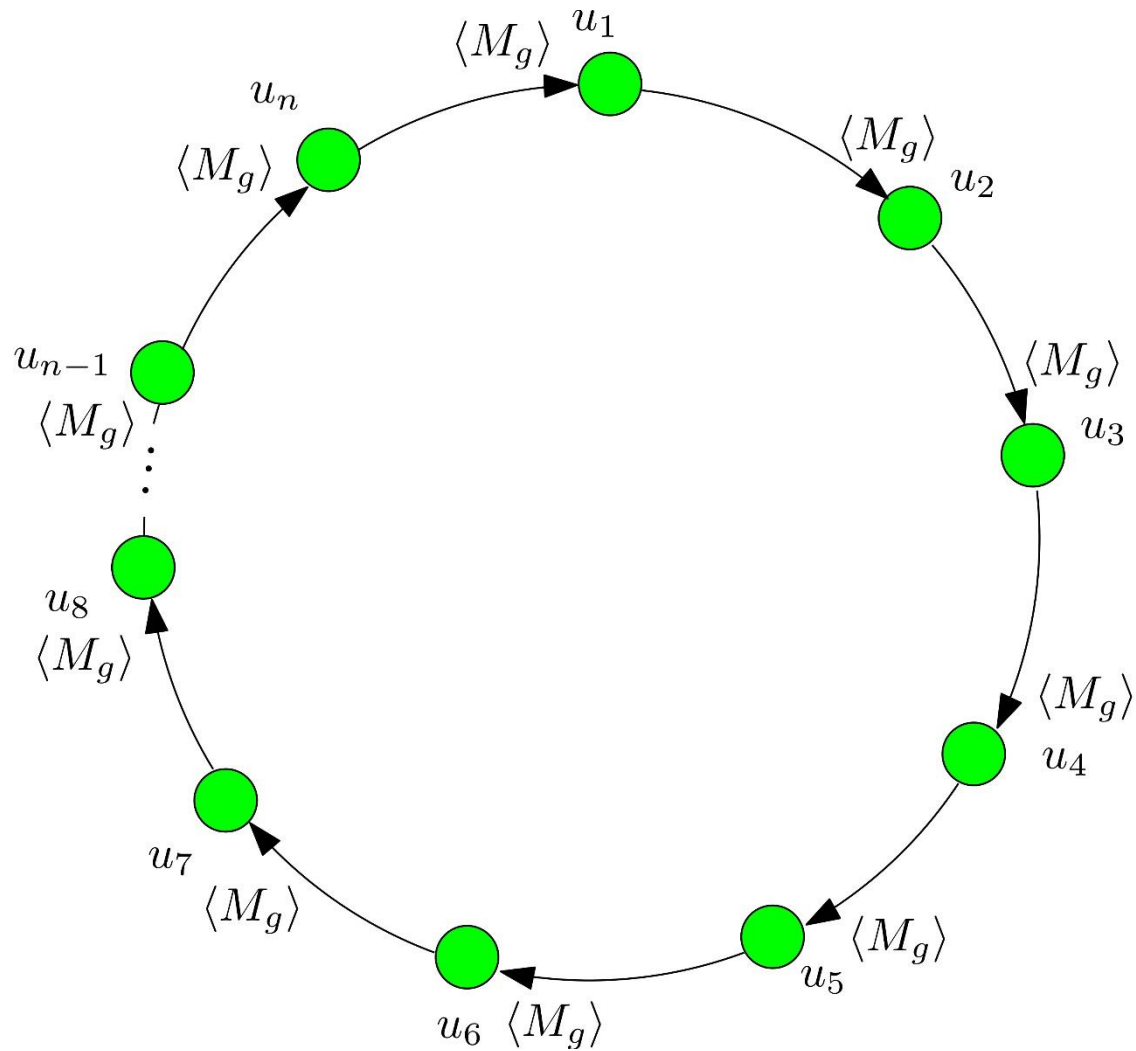
An Illustration



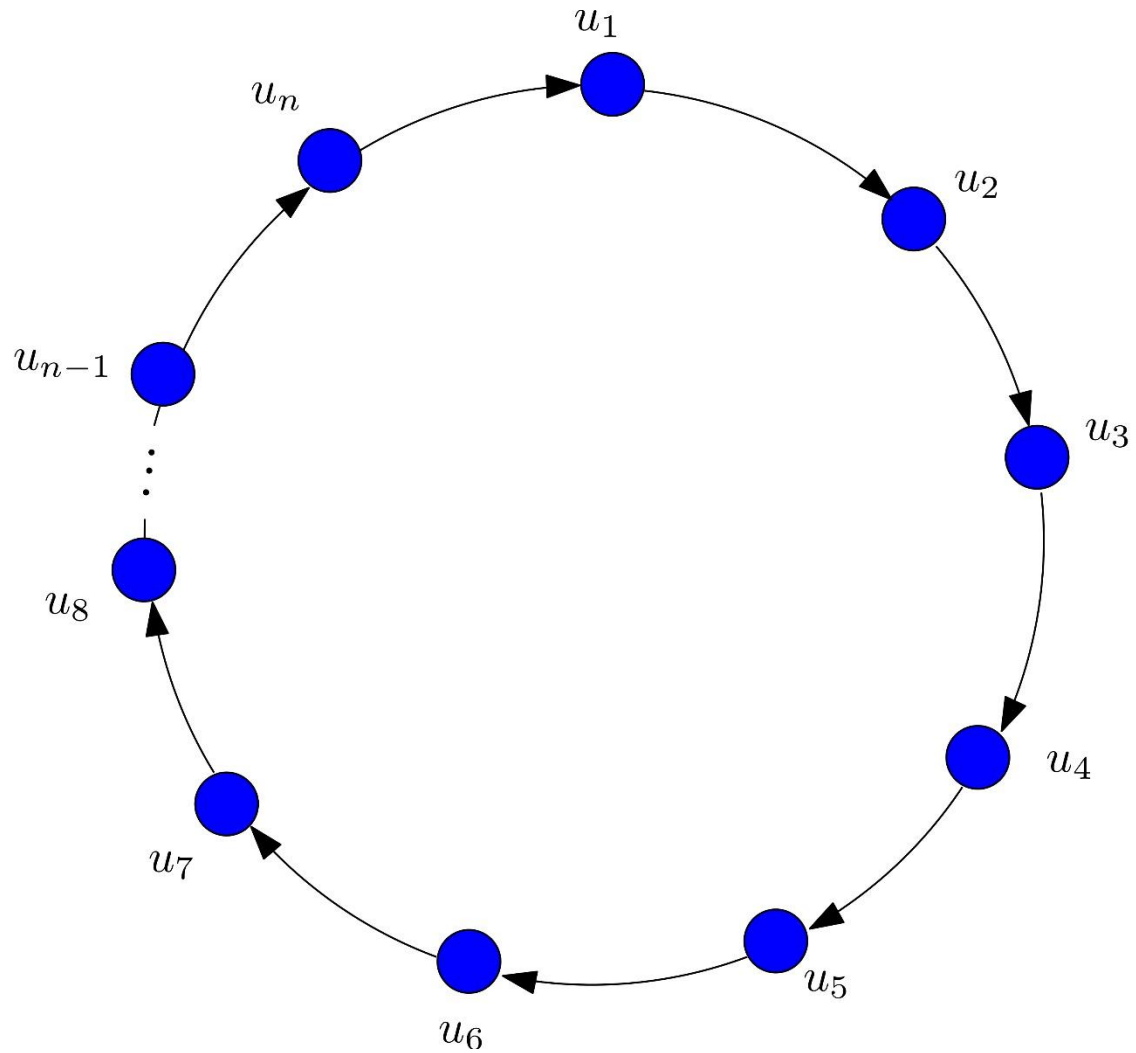
An Illustration




An Illustration



An Illustration



An Illustration



Interpretation

- Just proved that **all algorithms will fail**
- Therefore, not a drawback of a specific algorithm rather an **inherent limitation** in the setting under consideration
- Reveals the difficulties related to **symmetry** in distributed systems
 - Not always possible to **break symmetry**
 - Due to the independent/autonomous nature of processors they may very well be doing identical things
 - e.g., many of them regenerating the lost token even though only 1 is needed


Back to Optimism

- That we cannot solve a problem in a particular setting **does not mean that we cannot solve it in all settings**
 - Particularly true for distributed systems: even small modifications of the setting may greatly affect solvability/efficiency
- Immediate thought: *What if there is no such inconvenient symmetry initially?*
 - Still allow *some components of the processors states to be identical*
 - But now all processors have *distinct/unique initial states...*
 - *Any ideas?*

A Less Symmetrical Setting

Every processor has a **unique id**

- Largely available in distributed systems
- e.g., a MAC address
- Requires only $O(\log n)$ **bits** in local memory
- Not necessarily consecutive
- A processor only knows its own id initially



Leader Election with ids

- Processors have **unique ids** and **do not know n** in advance
- **LCR algorithm:** solves the problem
 - Le Lann, Chang, and Roberts [1977, 1979]
- Uses only **transmission** and **comparison** of ids
- **Simplest version:** Only the elected processor gives “output” and terminates
 - e.g., “I am the leader”
 - The other processors **never produce any output** and **do not terminate**

LCR: Informal description

- All processors **send initially their id clockwise**
- Upon receiving an **incoming id**, compare it to your own
 - If *incoming id* > *own id*, **forward** the received
 - if *incoming id* < *own id*, **discard** the received
 - if *incoming id* = *own id*, **declare yourself the leader**
- **Intuitively:**
 - The **maximum id** will manage to perform a **complete turn and return to its origin**
 - **Any other id** will at some point meet a processor with greater own id and **will be discarded before making a complete turn**

LCR: Pseudocode

Algorithm LCR

State of processor u_i :

- $myID_i$: holds the processor's unique id
- $sendID_i$: holds a unique id to be sent or *null*
- $status_i \in \{\text{"unknown"}, \text{"leader"}\}$: indicates whether u_i has been elected ("leader") or not ("unknown")

LCR: Pseudocode

Algorithm LCR

Code for processor u_i , $i \in \{1, 2, \dots, n\}$:

Initially:

u_i knows its own unique id stored in $myID_i$
 $sendID_i := myID_i$
 $status_i := \text{"unknown"}$

if round = 1 then

send $\langle sendID_i \rangle$ to unique out-neighbour

else // round > 1

upon receiving $\langle inID \rangle$ from unique in-neighbour // an id arriving from the left

if $inID > myID_i$ then // if greater than your own

$sendID_i := inID$ // forward it

send $\langle sendID_i \rangle$ to unique out-neighbour

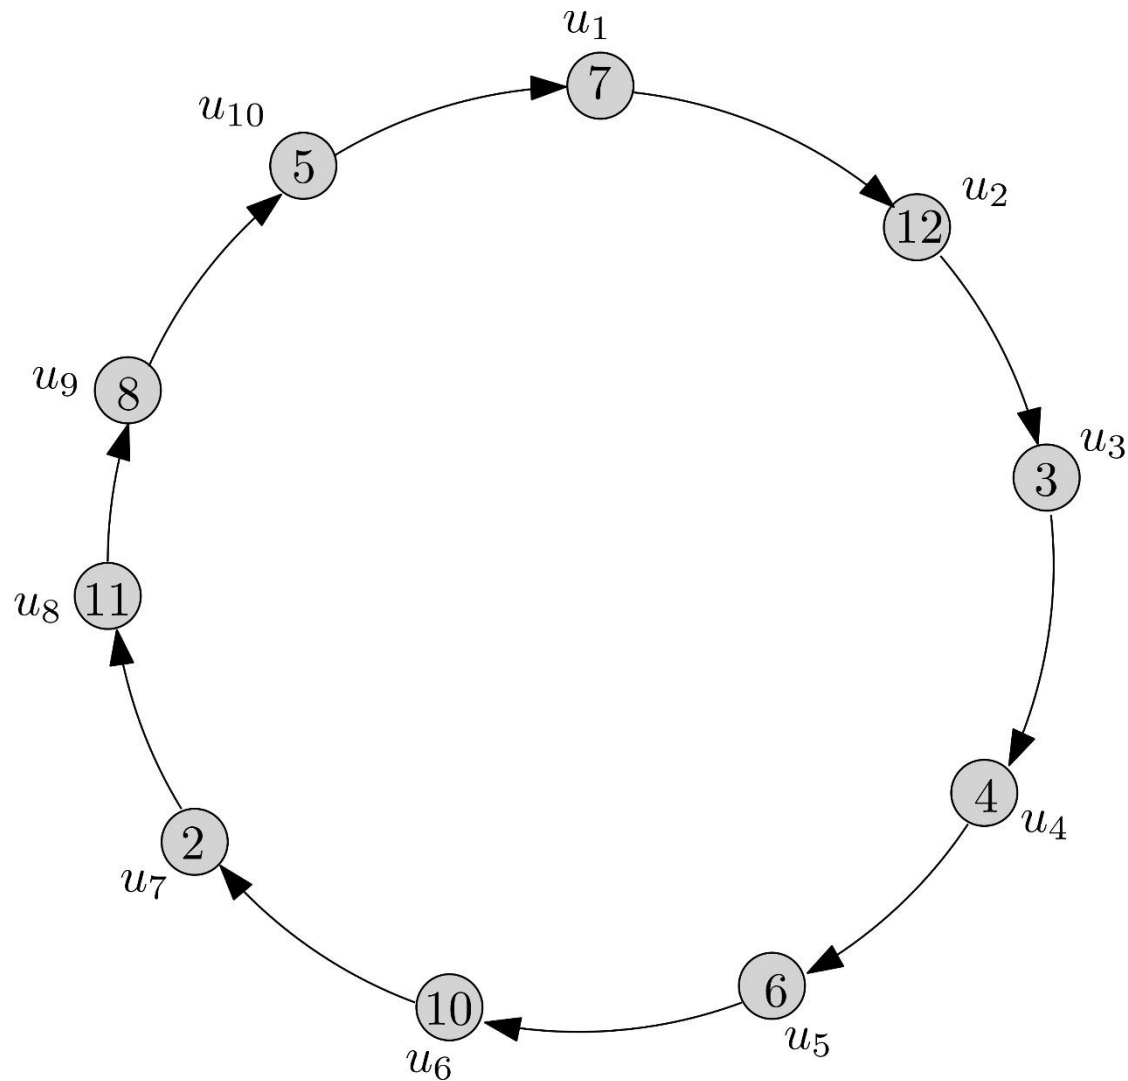
else if $inID = myID_i$ then // if equal to your own, your id managed a complete turn

$status_i := \text{"leader"}$ // therefore, elect yourself a leader


else if $inID < myID_i$ then // if smaller than own

do nothing // ignore (discard) it

Example Execution

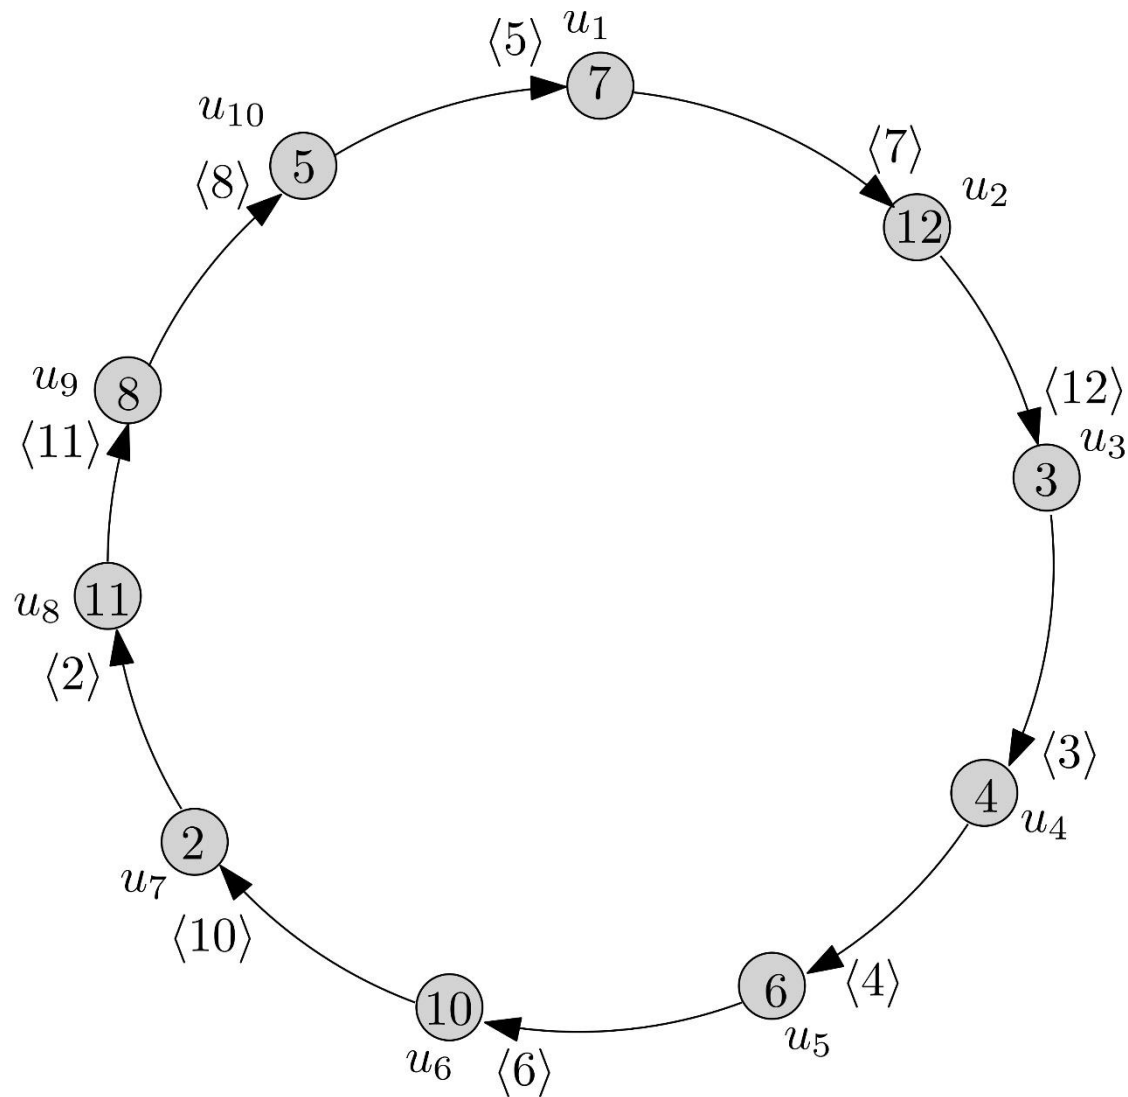


Example Execution



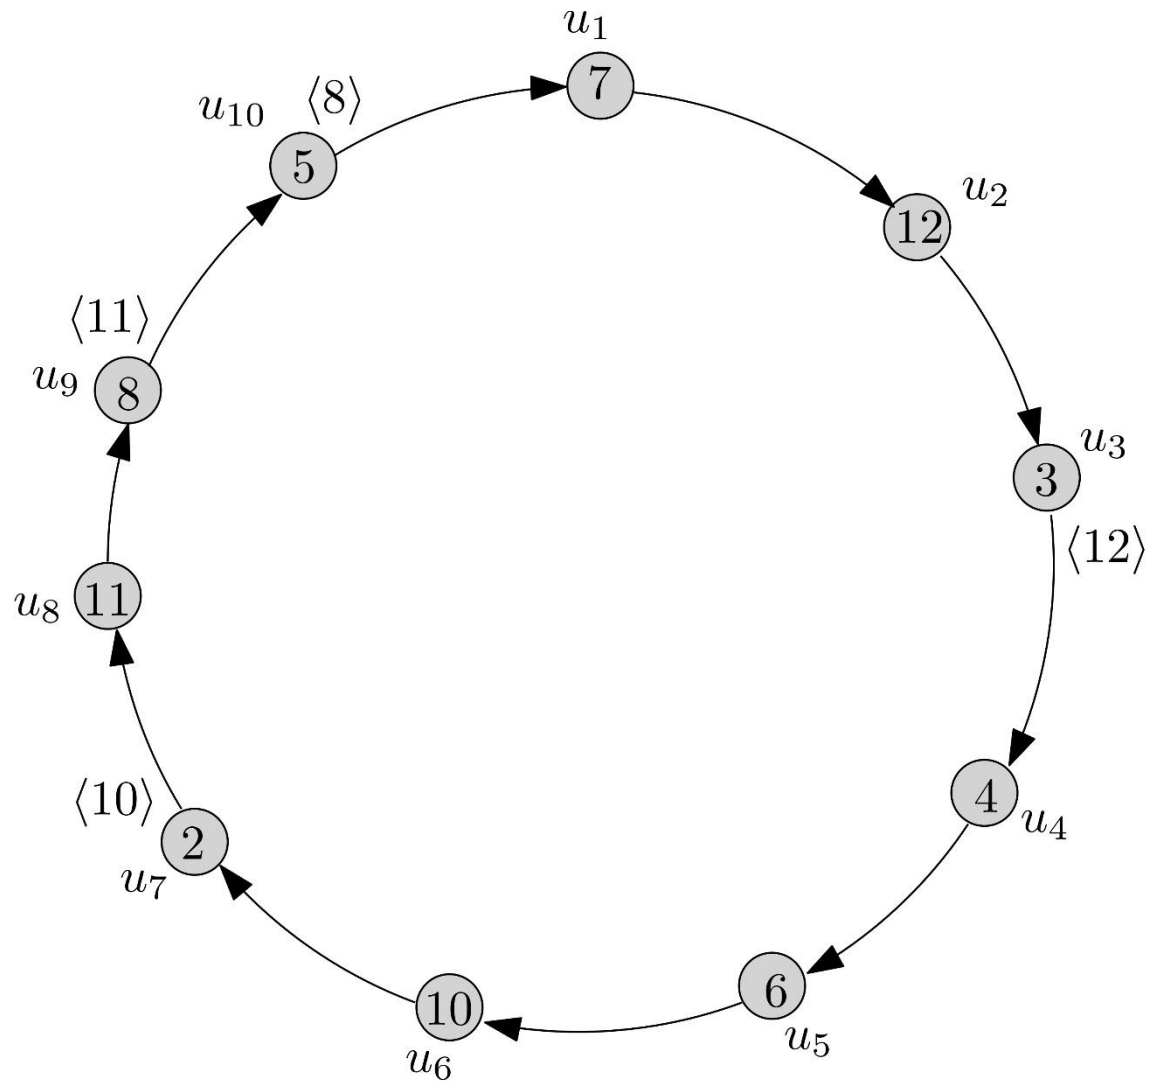
round = 1

Example Execution



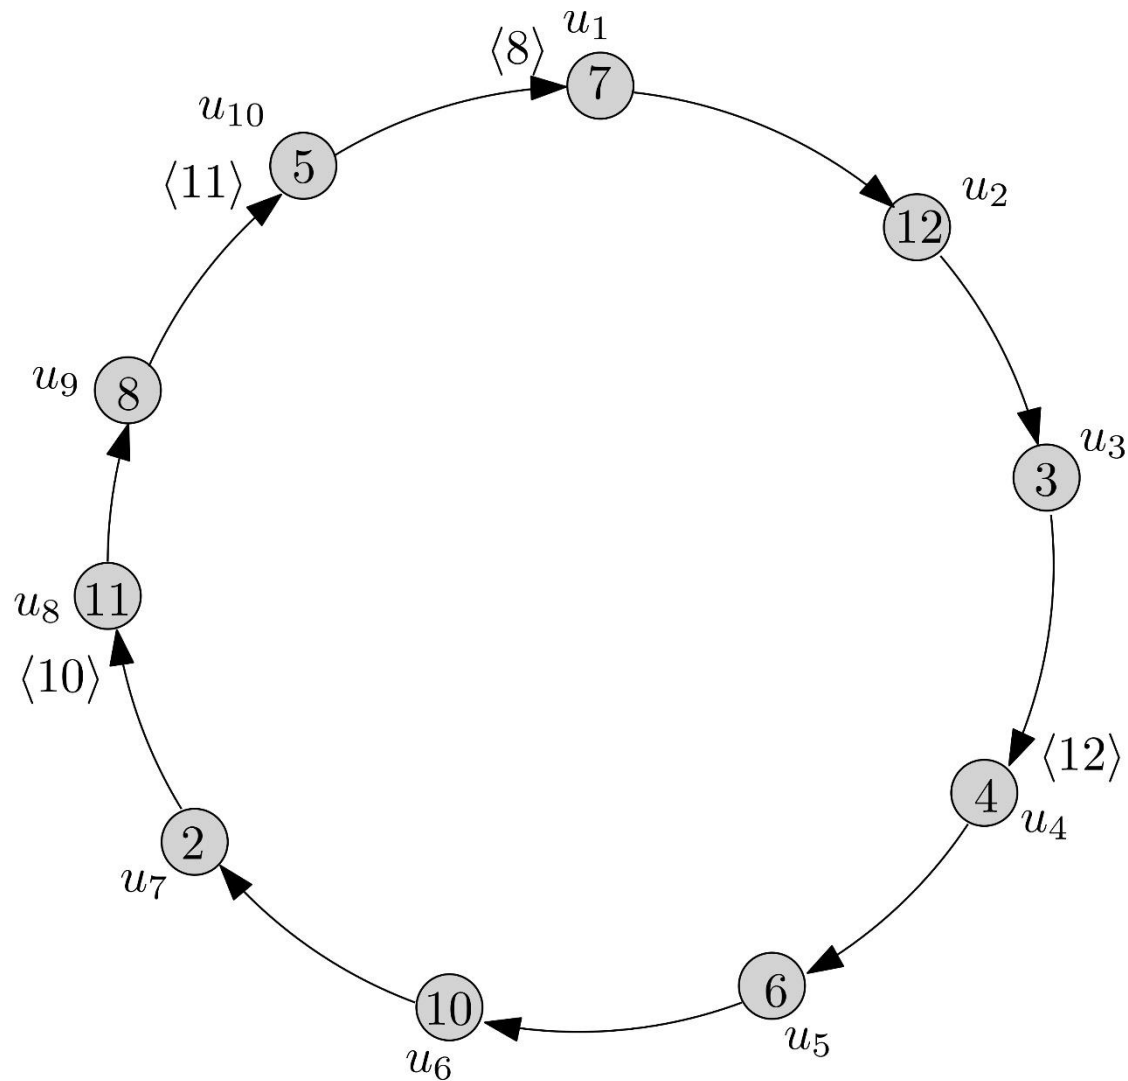
round = 1

Example Execution




round = 2

Example Execution



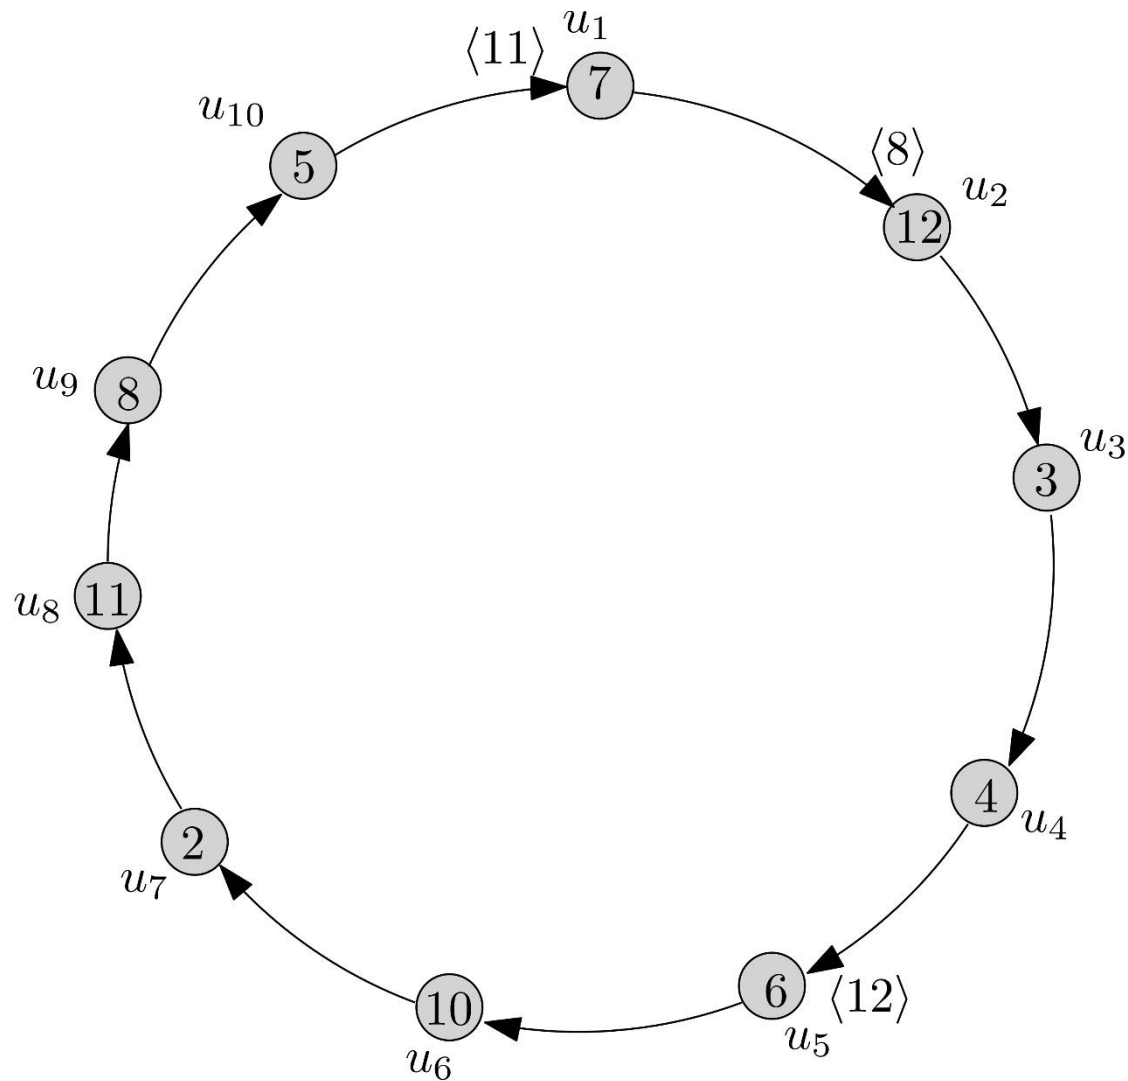
round = 2

Example Execution



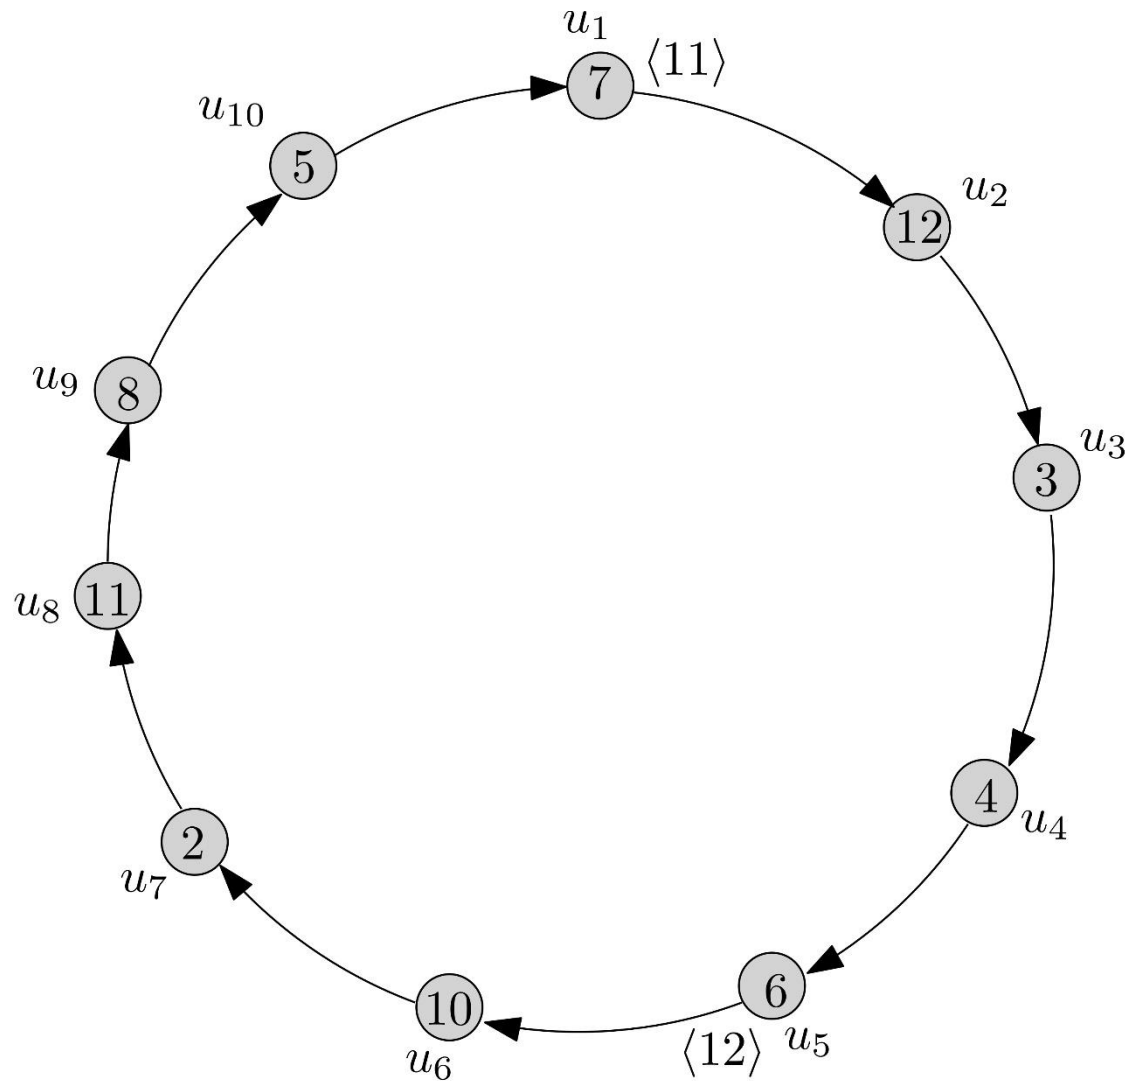
round = 3

Example Execution



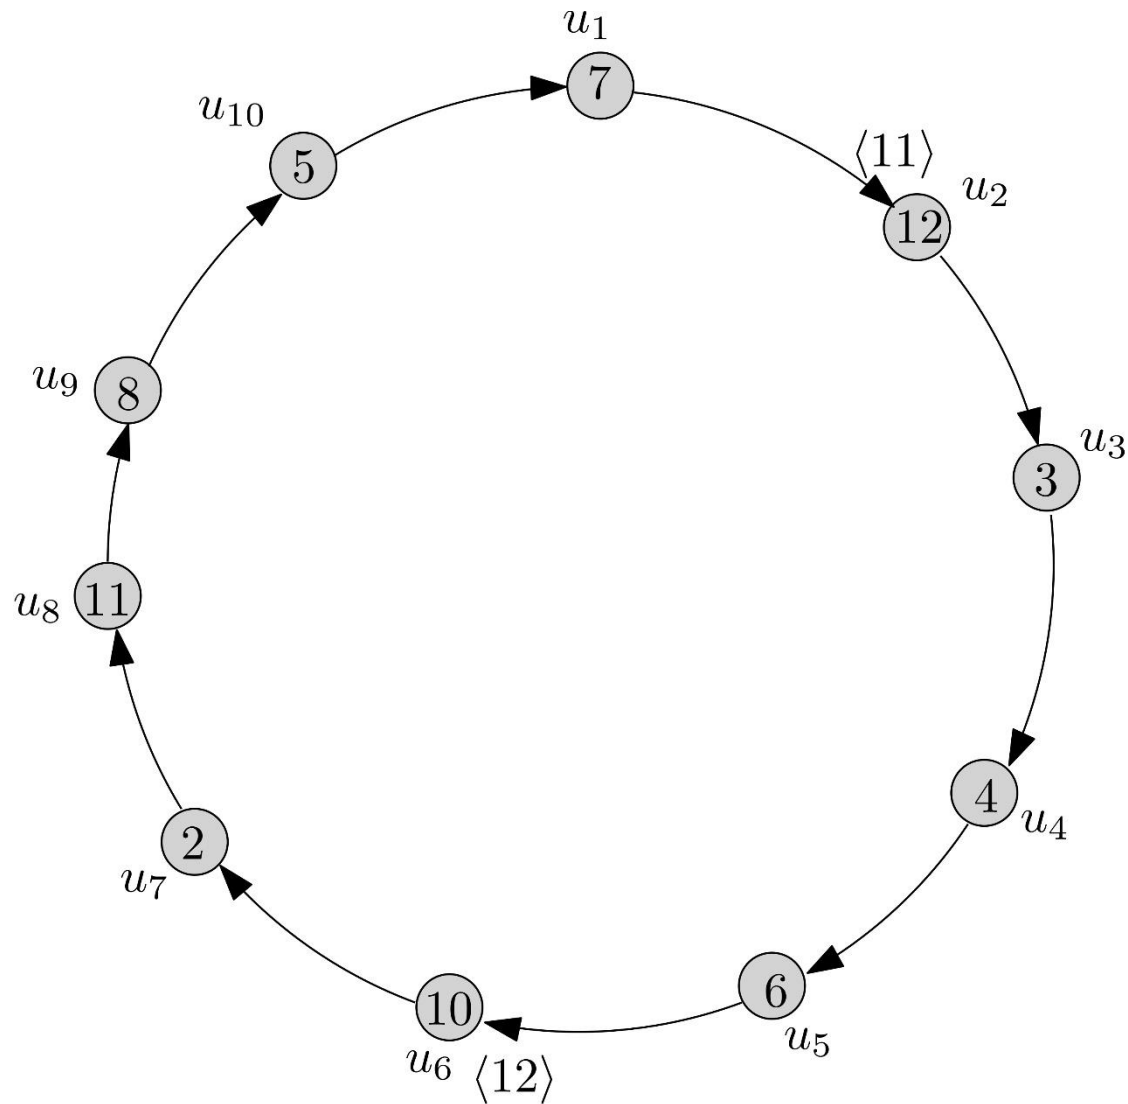
round = 3

Example Execution



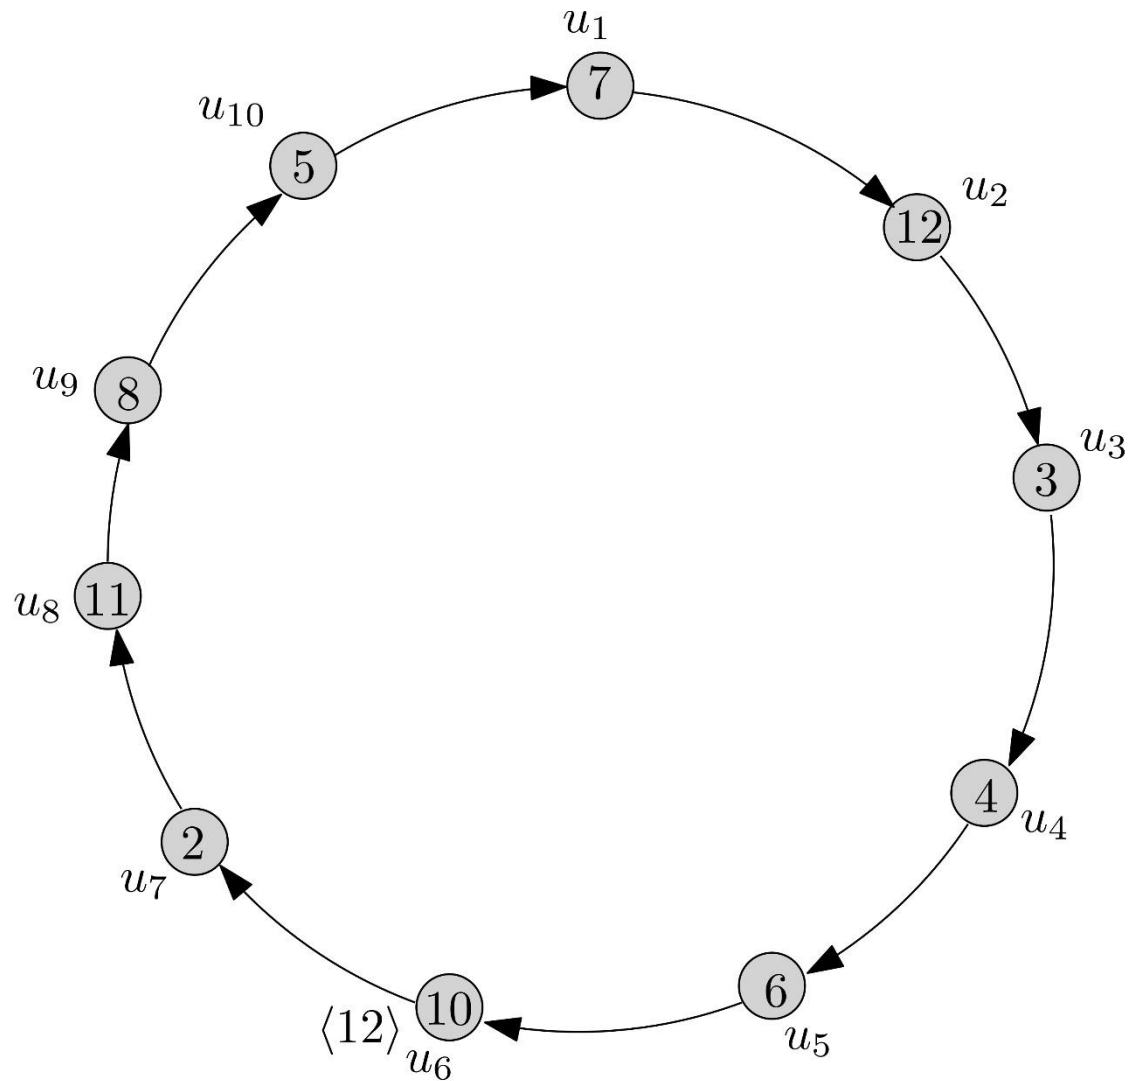
round = 4

Example Execution



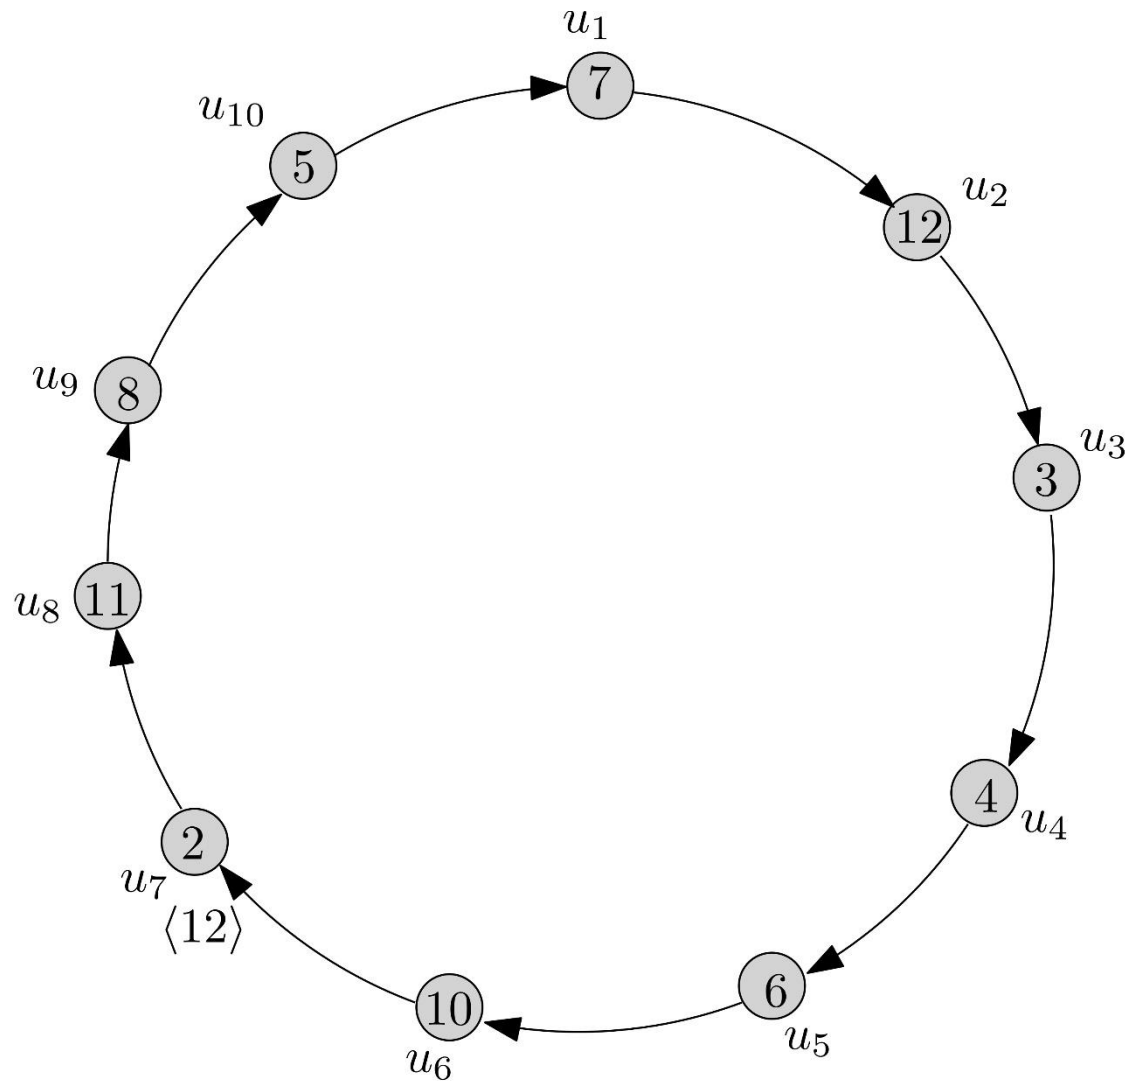
round = 4

Example Execution



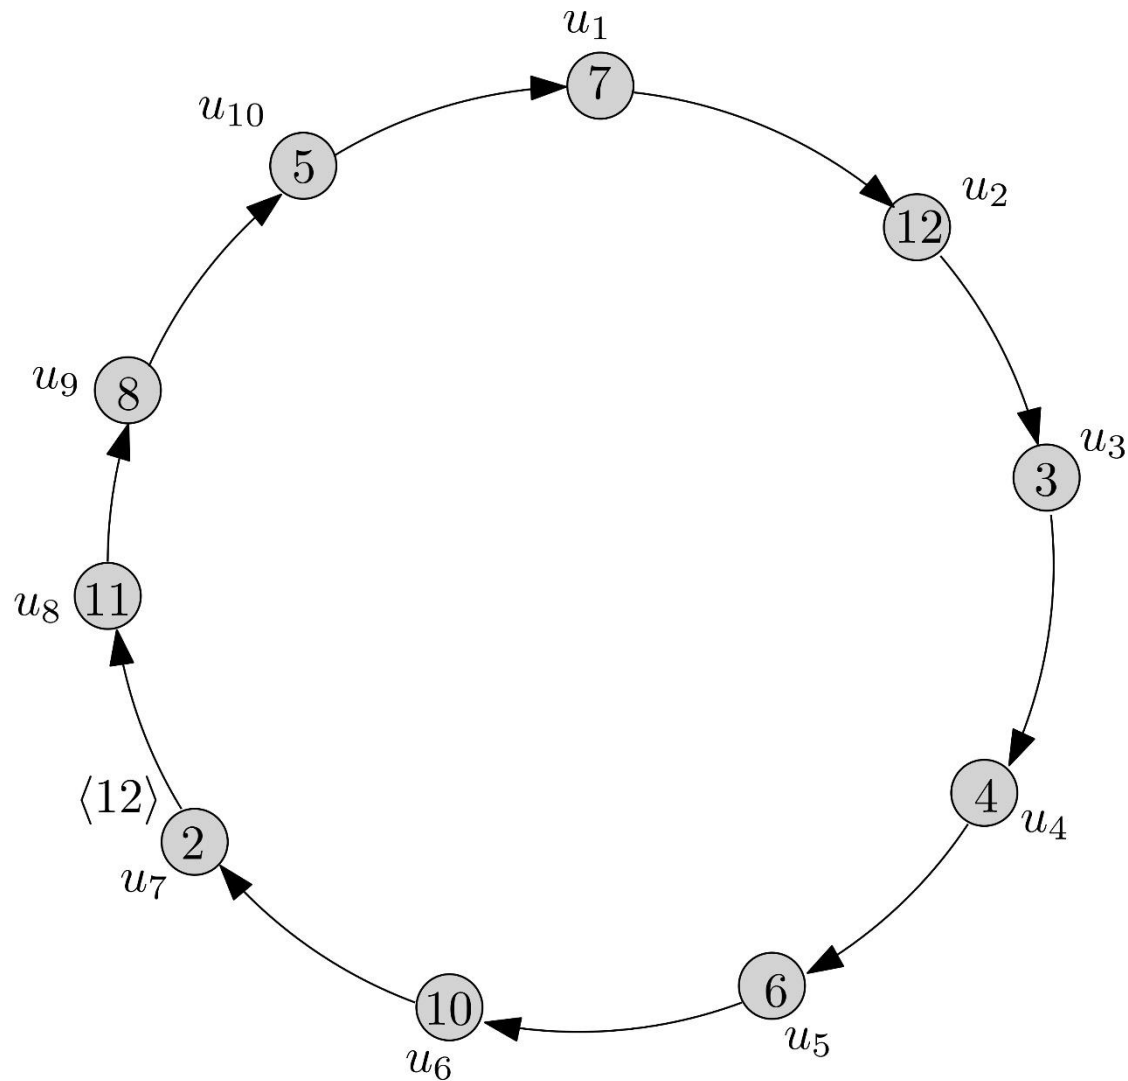
round = 5

Example Execution



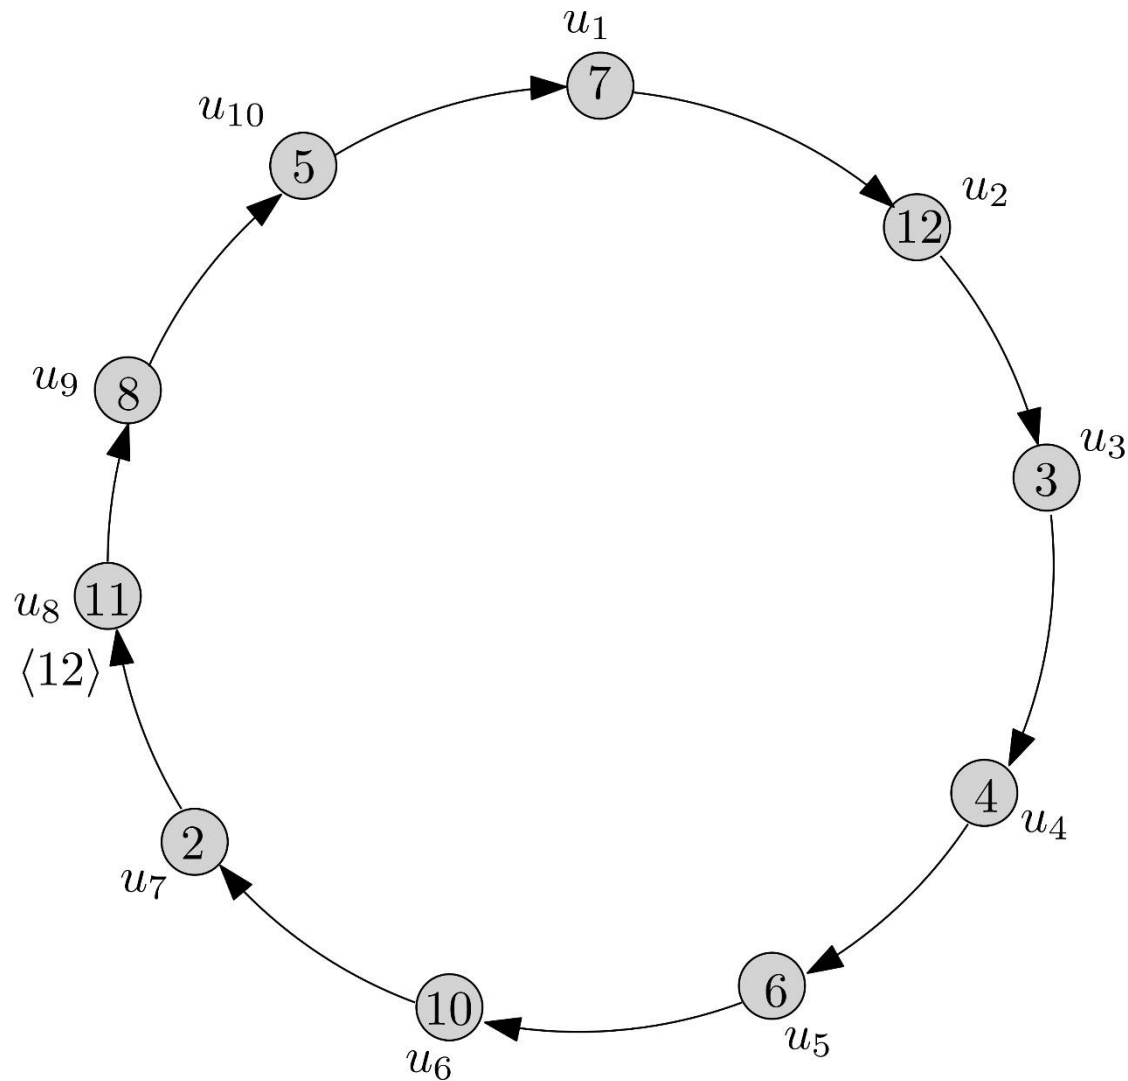
round = 5

Example Execution



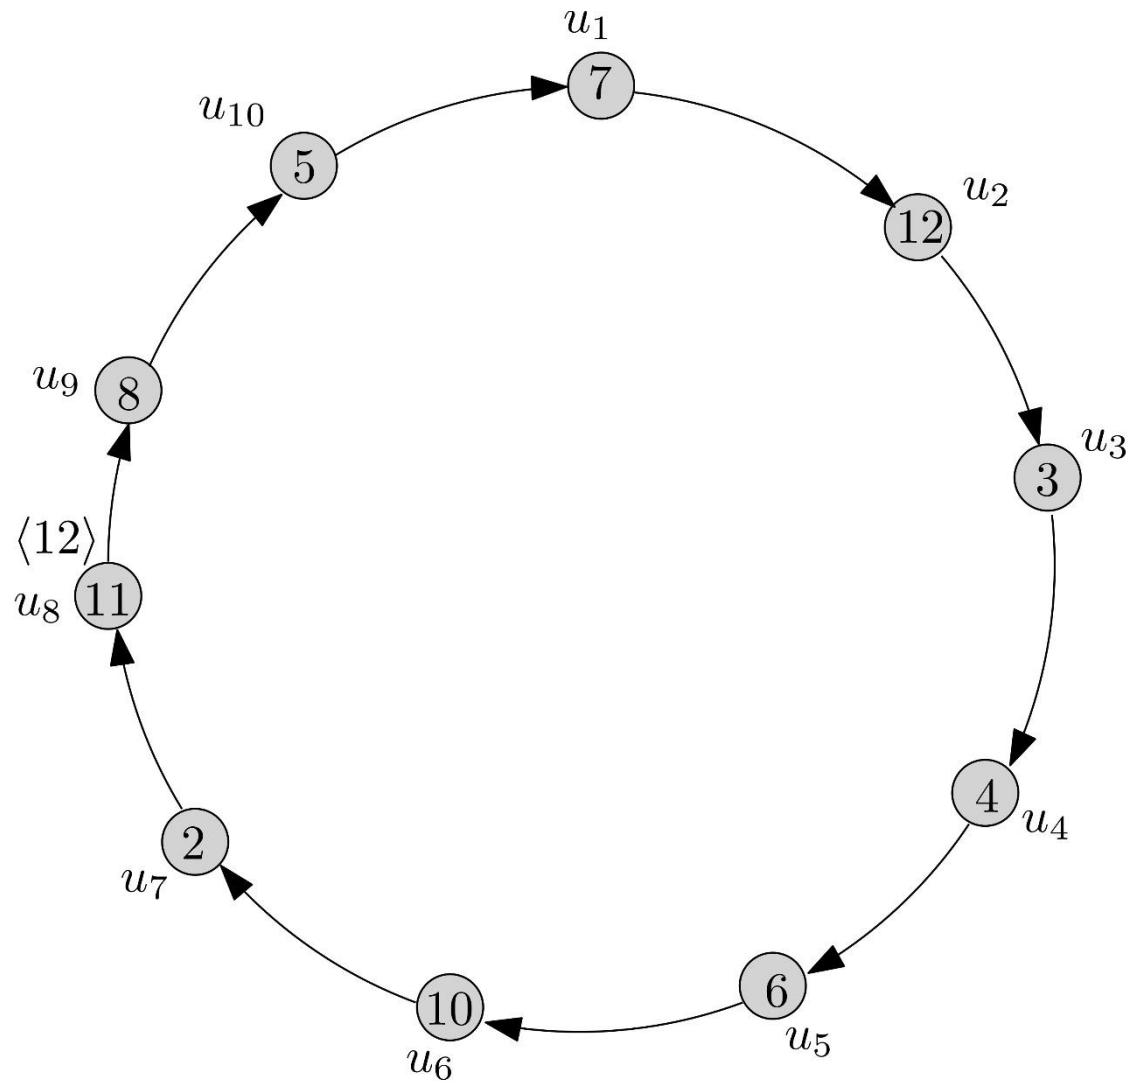
round = 6

Example Execution



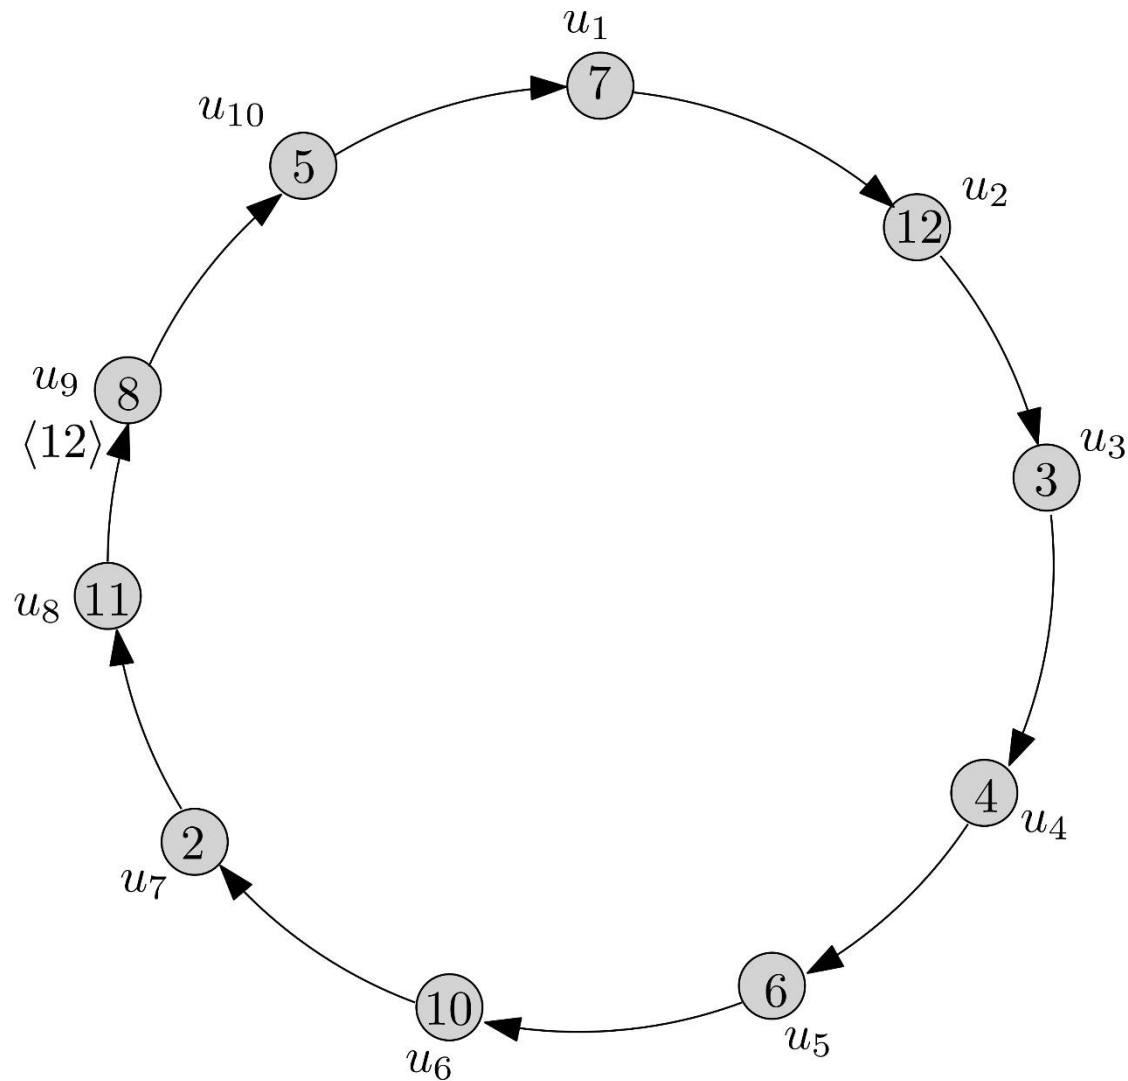
round = 6

Example Execution



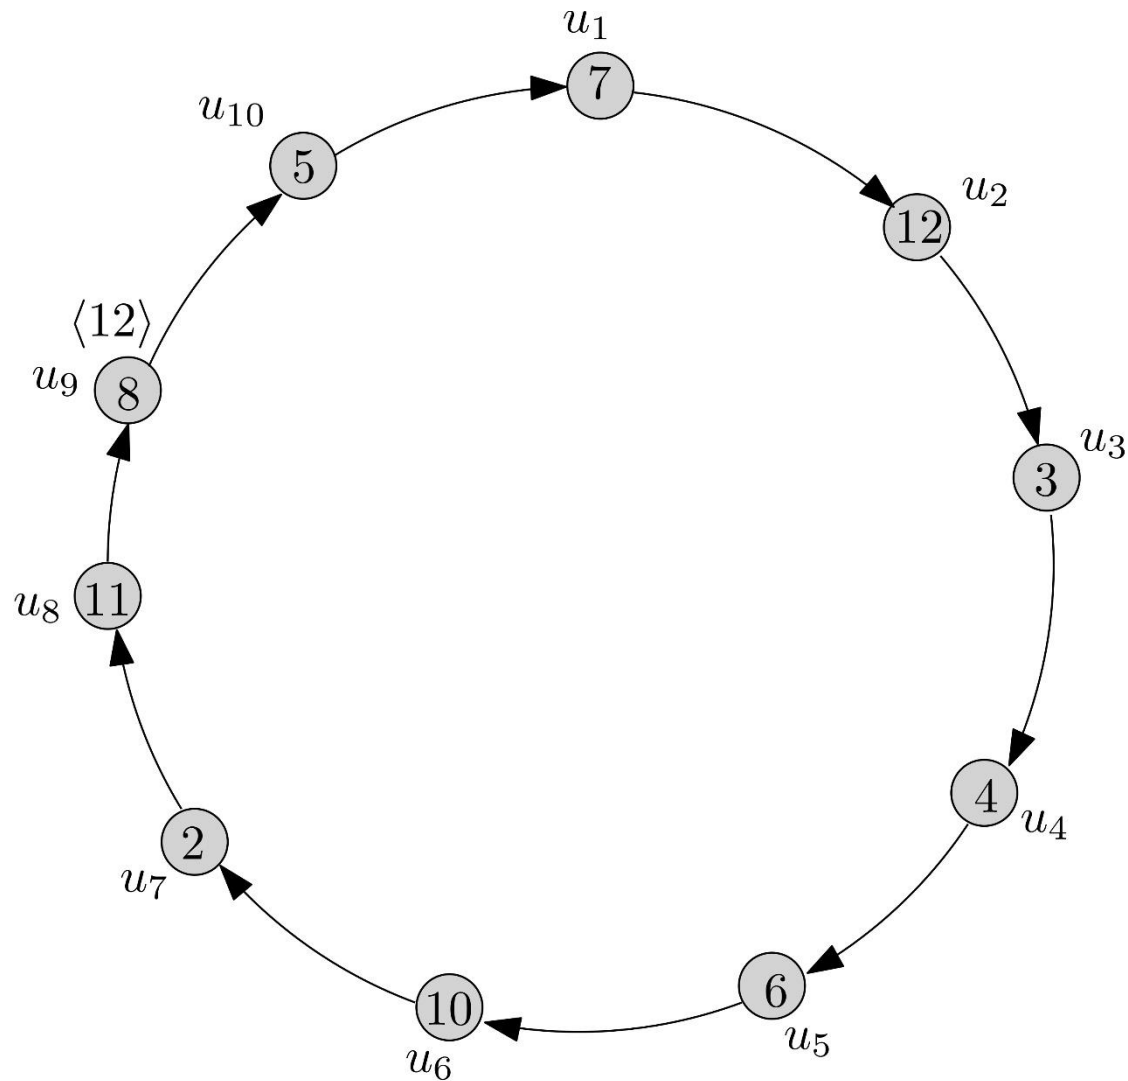
round = 7

Example Execution



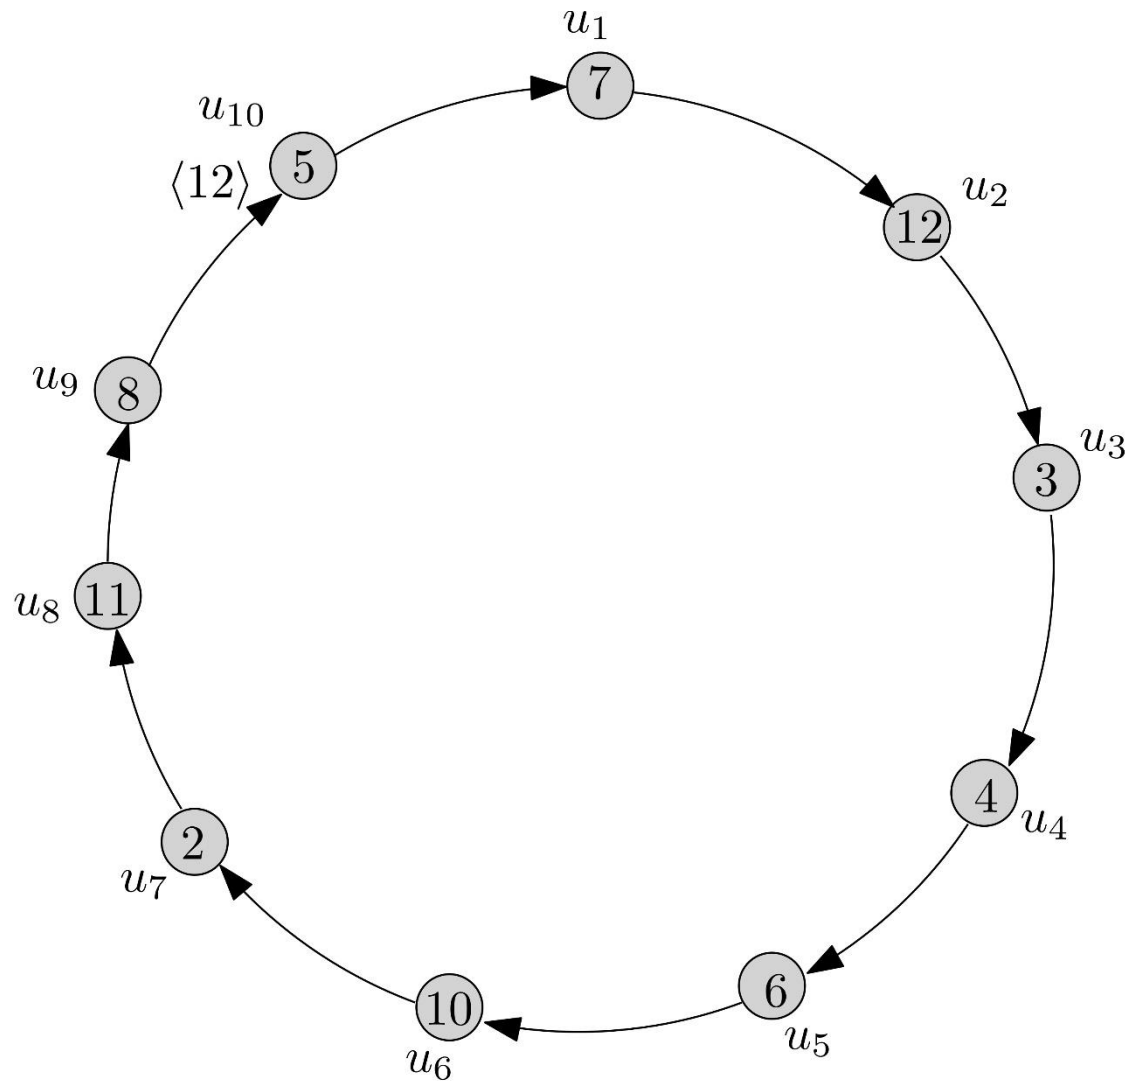
round = 7

Example Execution



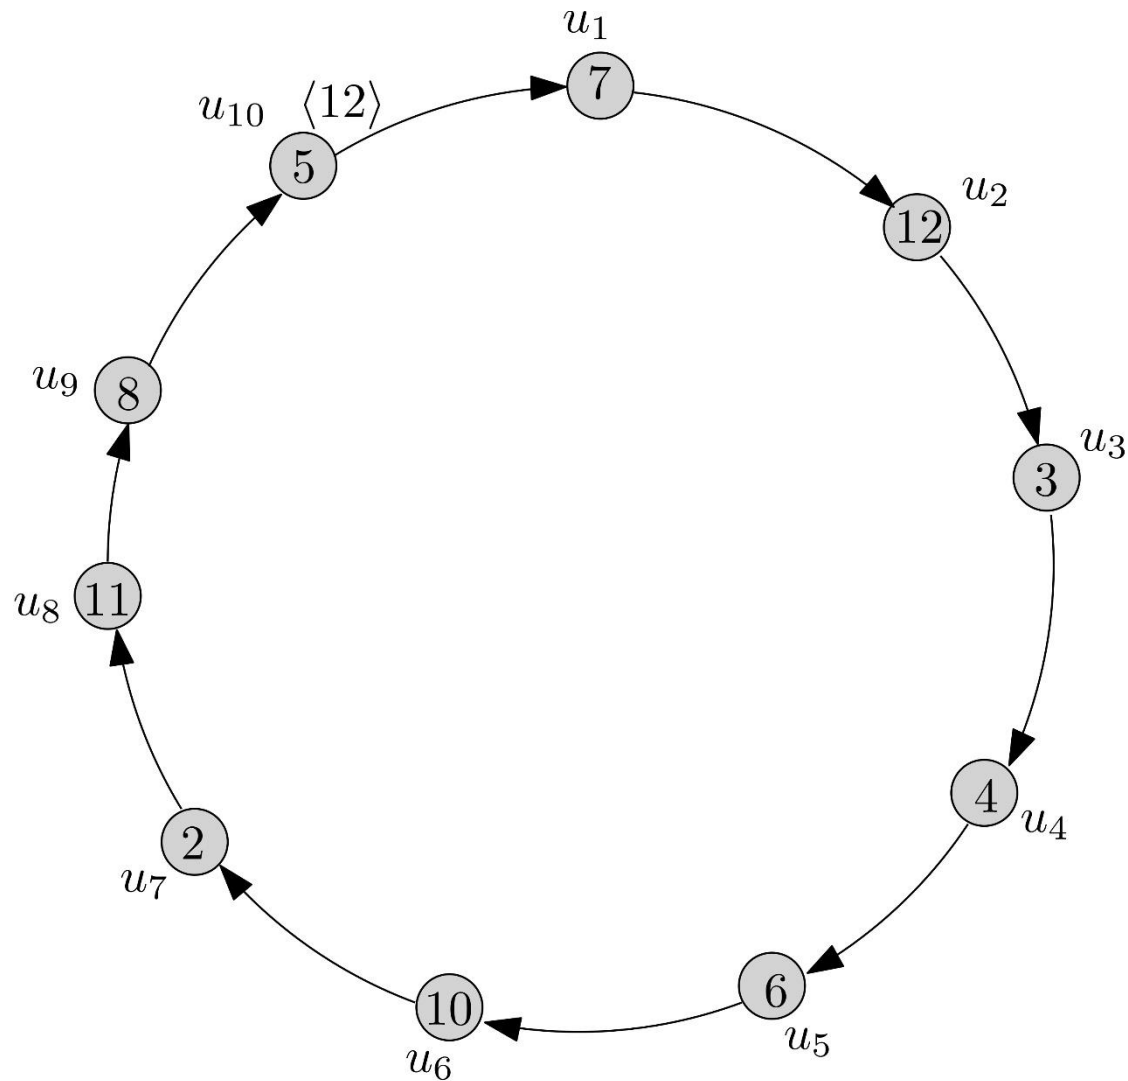
round = 8

Example Execution



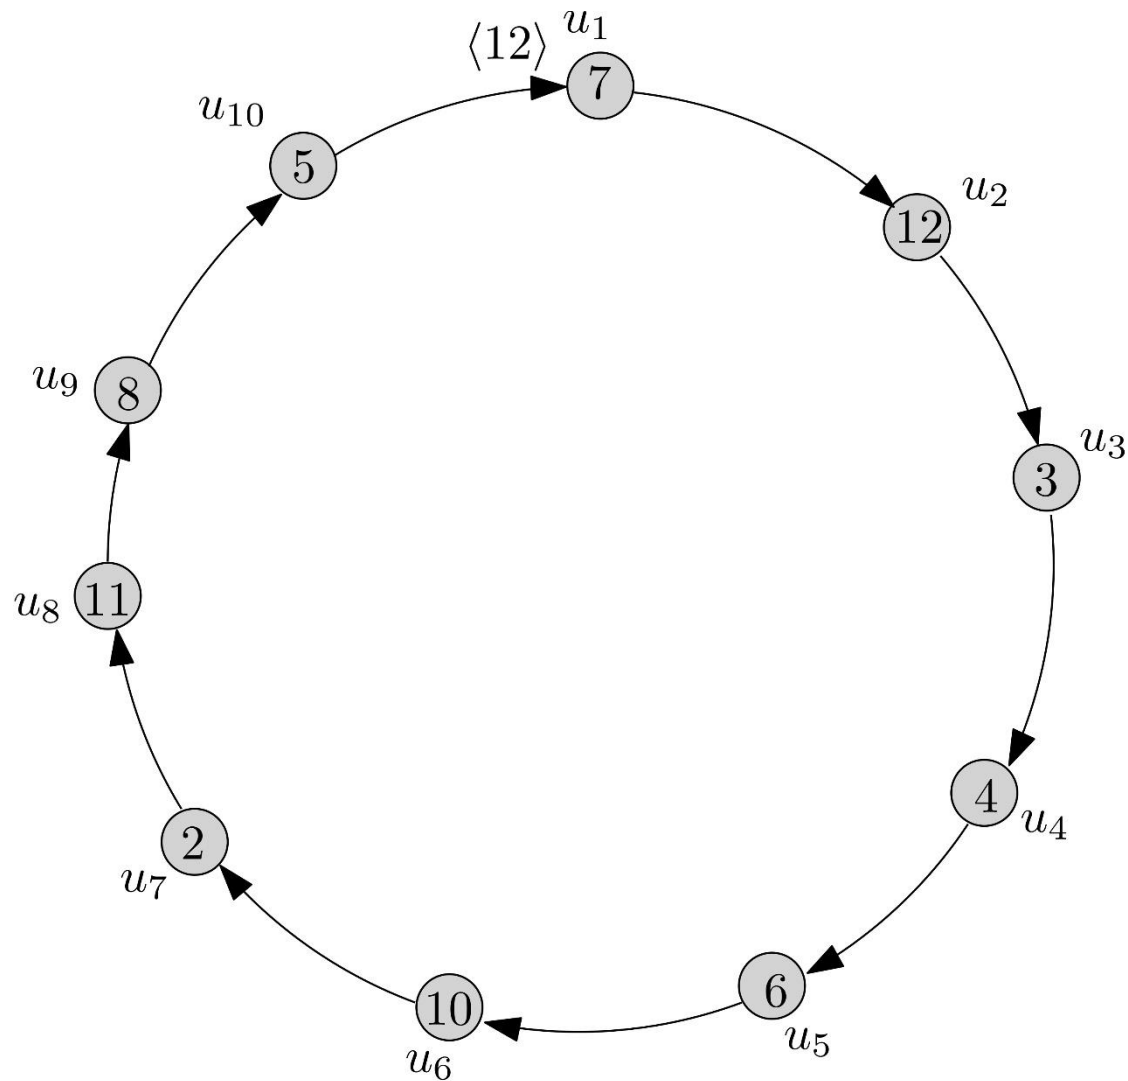
round = 8

Example Execution



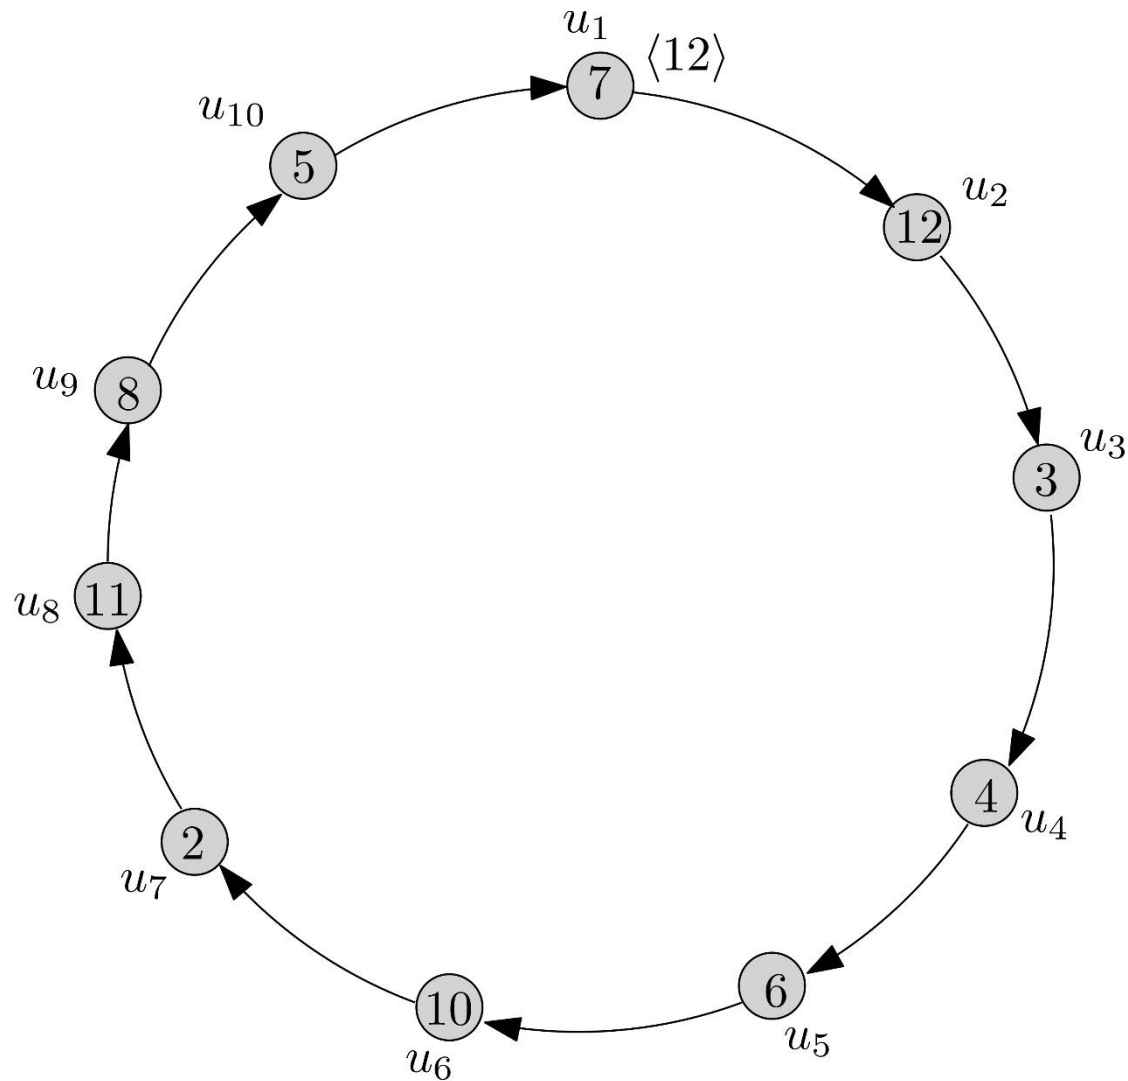
round = 9

Example Execution



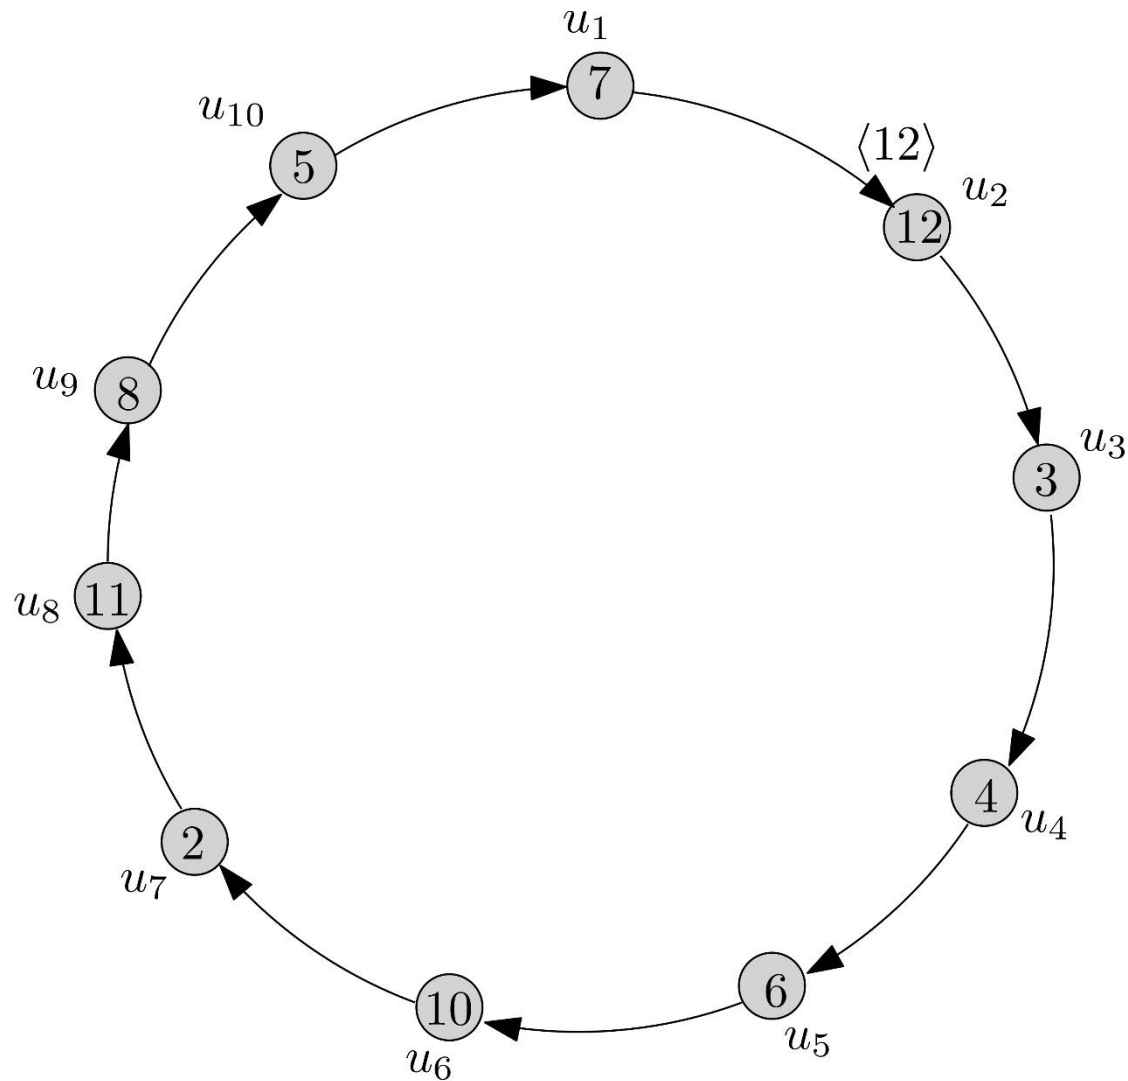
round = 9

Example Execution



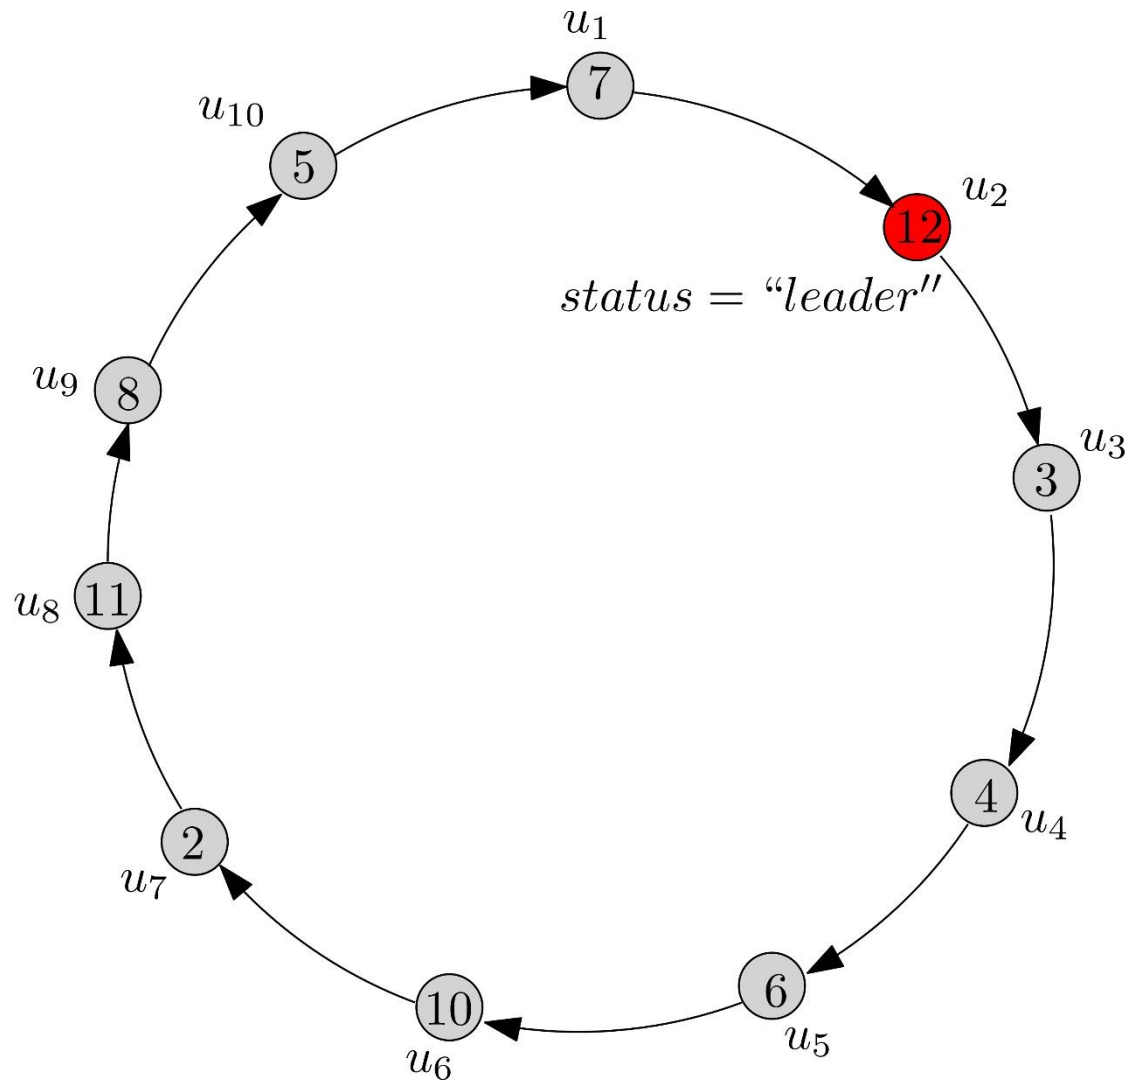
round = 10

Example Execution



round = 10

Example Execution



round = 11

Correctness and Complexity

- **Correctness:**
 - some processor is eventually elected
 - never 2 or more processors are elected
- **Time complexity:**
 - n rounds
- **Communication complexity:**
 - size of messages: encoding in bits of the maximum id
 - $O(n^2)$ messages in the worst case
 - *Can you think which is the worst case for this algorithm?*
- *How can we make all nodes terminate and know the elected leader and what will be the additional effect in performance?*

Think of these and we shall prove them in class

Summary

- Leader election is crucial for distributed systems
 - breaks symmetry
 - allows for coordination
- If all processors are initially identical then
 - impossible to elect a leader even in very simple networks
 - e.g., a ring
- Adding unique ids breaks this inconvenient initial symmetry
- The LCR algorithm elects a leader in any ring network
 - simple conceptually
 - assumes unique ids
 - n rounds
 - $O(n^2)$ messages
 - At a small additional cost can be made terminating and inform all processors of the elected one