

Distributed Systems

COMP 212

Lecture 12

Othon Michail

Communication

Interprocess Communication

- The “heart” of every distributed system.
- **Question:** how do processes on different machines exchange information?
- Answer: **with difficulty** ... ☹️
- **Message passing** harder than **shared memory**
 - But shared memory in general **not available** in DSs
- Established computer network facilities are **too primitive**, resulting in DSs that are too difficult to develop
 - Facilities not expected to change dramatically soon

Client-Server Model


Participants are divided into

- **Servers**: implementing a specific service
- **Clients**: requesting a service from a server by sending it a request and subsequent waiting for the server's reply

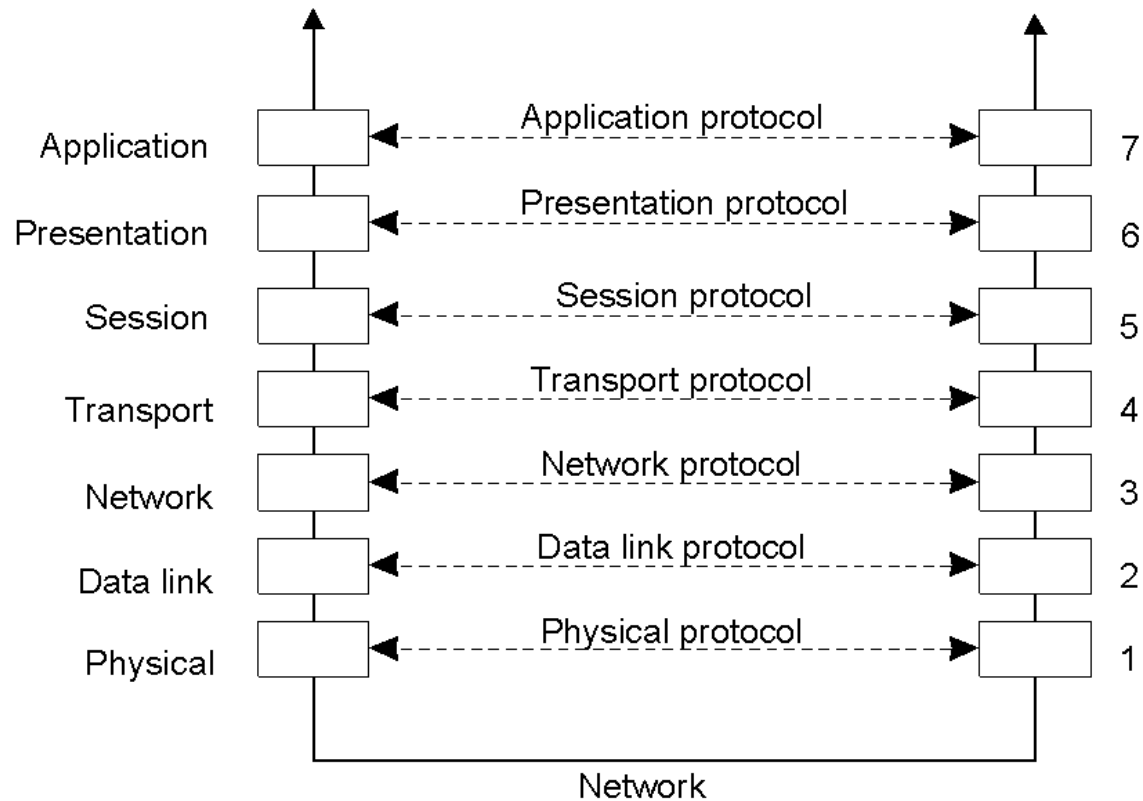
Distributed across different machines

- possibly overlapping

Follow a **request-reply** behavior



Layered Protocols



Layers, interfaces, and protocols in the OSI model

Physical Layer

- Concerned with transmitting 0s and 1s
- Specifies the mechanical, electrical and optical behaviour of the plugs and (hardware) sockets of a physical interface to the network
- Implementations of the physical layer perform functions in order to signal, modulate, and synchronize the transmission of bits over the physical network
- It may perform error detection by monitoring the quality of the electrical or optical signals received

Data Link Layer

- Basic task: error correction
 - Groups bits into frames
 - Puts a special bit pattern at the start and end of each frame
 - Also appends a checksum
 - Receiver can check correctness by recomputing the checksum and comparing to the checksum received
- Responsible for transmission of frames between nodes that are directly connected by a physical link
 - Nodes may be routers, switches, or computers
- It maps the physical connection provided by the physical layer into a layer that can transmit frames from one node to another in a relatively error-free mode
- Would be sufficient for a LAN but not for a WAN
 - Does not provide a mechanism for locating the receiver

Network Layer

- Provides a **routing** mechanism
- **Routing**: How to choose the **best path** from one machine to another
 - The shortest route is not always the best
 - What really matters is the **delay** (as with road traffic)
- **Data link layer**: node-to-node connection
- **Network layer**: routing messages along multiple nodes in the network
 - It implements connections between arbitrary nodes of the network
- **IP** (Internet Protocol)
 - Part of the **Internet protocol suite**
 - Each IP **packet** (i.e., message) is routed to its destination independent of all others
 - No internal path is selected and remembered

Transport Layer

- Provides an interface between the **underlying network** and the **application developer**
- “Turns” the network into something that developers can use
- The Application layer expects from the Transport layer to “eventually” **deliver a message without loss**
 - Provides a **reliable connection service**
- 1. Breaks message into pieces, small enough for transmission
- 2. Assigns a number to each and
- 3. Sends them all
- 4. Packets **may not arrive at the correct order**
 - Puts everything back in order before delivering to the application
- **TCP** (Transmission Control Protocol)
 - The **transport protocol of the Internet**
 - **TCP/IP** is the de facto standard for network communication

Application Layer

- Contains all applications that do not qualify as members of the other layers
 - Mainly **application-specific protocols**
 - i.e., designed to meet the communication requirements of specific applications
 - Often defining the interface to a service
- Examples:
 - **HTTP** (HyperText Transfer Protocol)
 - Handles the transfer of Web pages
 - Implemented by Web browsers and Web servers
 - **FTP** (File Transfer Protocol)
- The Session and Presentation layers are **obsolete**
 - Have been replaced by a **Middleware layer**
 - contains **application-independent protocols** (general purpose)

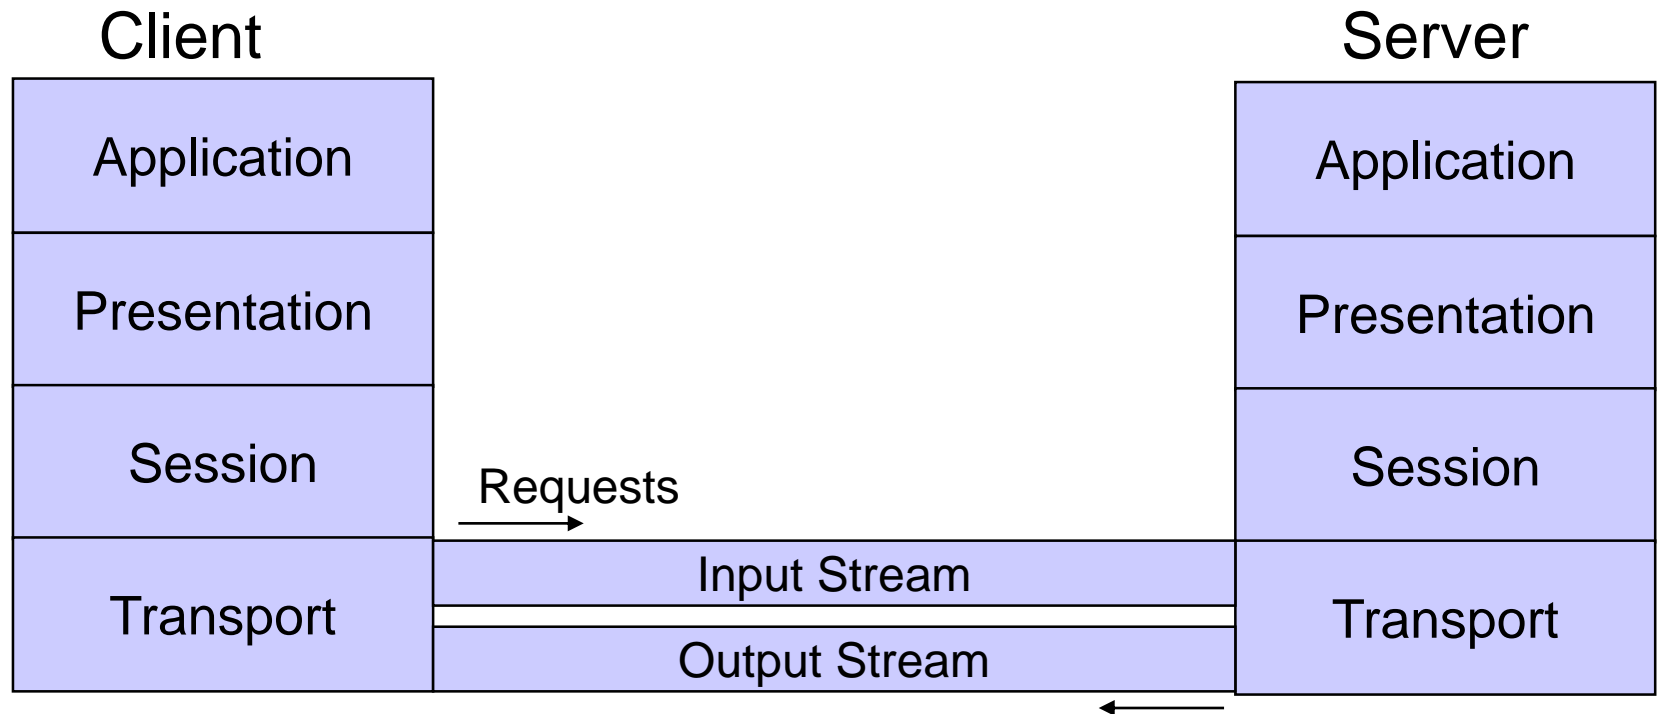
Connection (-less) Transport Layer

- **Connection-oriented** transmission protocols
 - Provide operations to:
 - Open a connection
 - Close a connection
 - Write data into such a connection
 - Read data from such a connection
- **Connection-less** transmission protocols
 - Provide the ability of sending fixed length messages between distributed hosts
 - These messages are referred to as **datagrams**


Internet Transport Protocols

- **TCP** (Transmission Control Protocol)
 - **Connection-oriented**
 - **Reliable** but **slow** (introduces overhead to guarantee delivery)
 - Applications: HTTP, POP, IMAP
- **UDP** (User Datagram Protocol)
 - **Connection-less**
 - Very **fast** but **unreliable** (does not guarantee delivery)
 - **Applications:** Domain Name System (DNS), Voice and Video streaming, VoIP
- Choice of protocol depends on whether
 - **reliability** or
 - **performance**is the major concern for a given application

TCP - Graphical Representation



UDP - Graphical Representation




- We do not study UDP in detail here

Communication at Transport Level

- Provided by all networks
- Practically, lowest-level protocol available to application developer

Sockets and Ports

- Reliable point-to-point communication provided by TCP
 - Processes bind sockets to a connection and read/write to them

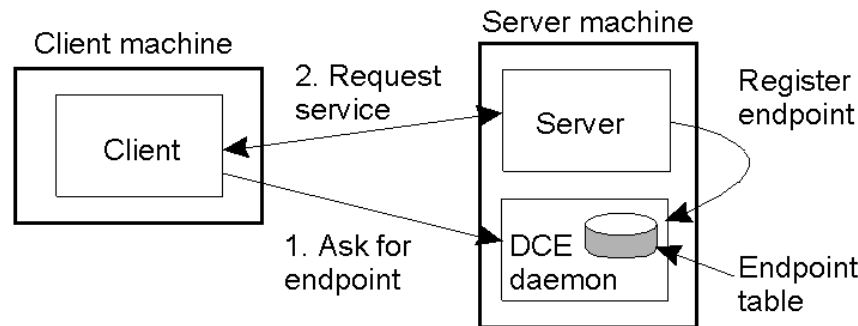


- A **socket** is one endpoint of a two-way communication link between two programs running on the network
- Is bound to a **port** number (logical construct) so that the TCP layer can identify the application that data is destined to be sent


Problem: Identifying “end-points”?

- How do clients know which end-point (or port) to contact a server at? How do they “bind” to a server?
1. Statically assigned end-points
 - e.g., TCP port 21 for telnet, TCP port 80 for HTTP, ...
 2. Dynamically assigned end-points
 - Special daemon on each machine that runs servers
 - A client first contacts the daemon
 3. “Super-server”
 - Listens to several ports, i.e., provides several services
 - Starts another process/thread, e.g., inetd on UNIX

Servers: Binding to End-Points



(a)



(b)

- a) Client-to-server binding using a daemon (DCE)
- b) Client-to-server binding using a super-server (inetd on UNIX)

Sockets in Java

Sockets in Java

- Facilities for simple network programming in Java can be found in `Java.net` package
- Main classes:
 - `Socket`
 - `ServerSocket`

The Socket Class

- Provides a number of **constructors** which enable the programmer to create a socket to a remote computer
 - Usually for programming the **client** side of the connection
- The simplest constructor has two arguments:
 - Hostname (a String: name of the computer using the DNS convention)
 - port number (an integer from 0 to $2^{16}-1 = 65535$)
- Example:

```
Socket mySock = new Socket("www.csc.liv.ac.uk", 80);
```

Reading from a Socket (1)

- The method `getInputStream` will obtain an `InputStream` object attached to the socket which can then be used to `read data`.
- **Example:**

```
String hostName = args[0]; // e.g. "www.csc.liv.ac.uk"  
int portNumber = Integer.parseInt(args[1]); // e.g. 8000
```

```
Socket myClientSocket = new Socket(hostName, portNumber);  
InputStream myInputStream = myClientSocket.getInputStream();
```

Reading from a Socket (2)

- Also we can attach a buffer to the InputStream object which can contain data until enough can be sent in an efficient way
- **Example:**

```
Socket myClientSocket = new Socket(hostName, portNumber);  
InputStream myInputStream = myClientSocket.getInputStream();  
BufferedReader input = new BufferedReader(new  
    InputStreamReader(myClientSocket));  
String lineRead = input.readLine();
```

Writing to a Socket

- The method `getOutputStream` can be used to `write` data to the socket.
- **Example:**

```
Socket myClientSocket = new Socket("localhost", 2048);  
OutputStream myOutputStream =  
    myClientSocket.getOutputStream();  
PrintWriter output = new PrintWriter(myOutputStream, true);
```


Connection is Bi-directional

```
Socket aSocket = new Socket("www.csc.liv.ac.uk", 80)

//Set up the streams associated with the socket
BufferedReader input = new BufferedReader(
    new InputStreamReader(aSocket.getInputStream()));
PrintWriter output = new PrintWriter(
    aSocket.getOutputStream(), true);

//Send message to server
output.println("Hello");

//Get reply from server
String serverReply = input.readLine();
if (serverReply.equals("Hello")) {
    ... //Process the reply
} else {
    ... //Carry out some error process
}
```

The Server Socket

- Provides a number of constructors which enable the programmer to create a socket **on a server**
- The simplest is:

`ServerSocket(int port)`

- This creates a socket on the given port (this machine)

Accepting Connections

- The most important method: **accept**
- This method **suspends** the server until a client attempts to connect to it
- The method “accept” returns the Socket object which can be used to establish a connection with the client that has connected in

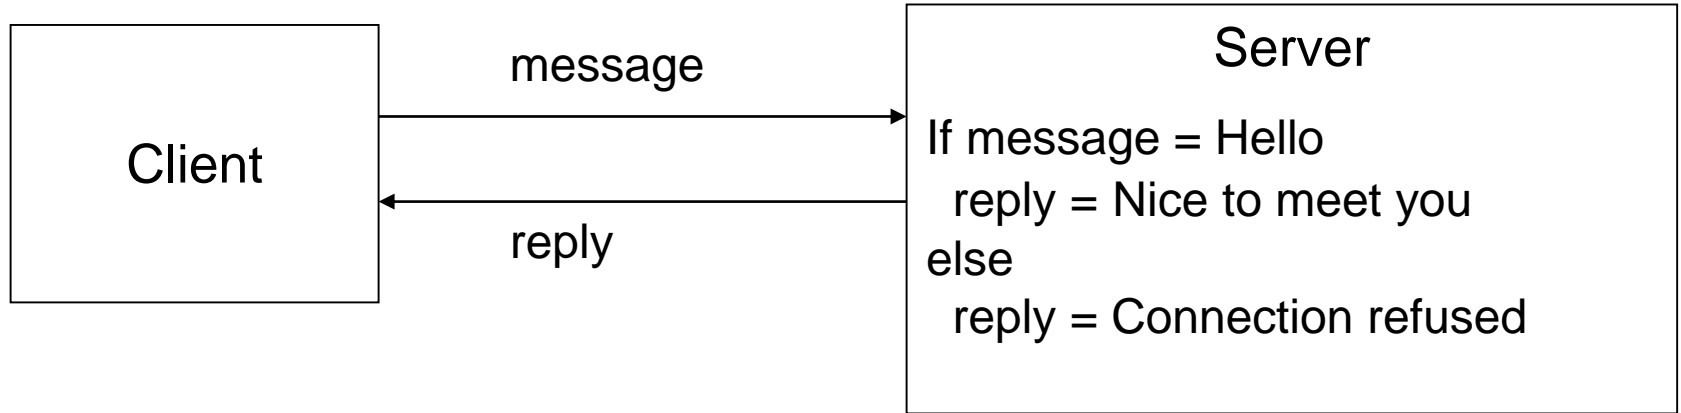
Server Socket I/O

```
ServerSocket serverSocket = new ServerSocket(2500);

// Wait for a connection
Socket clientSocket = serverSocket.accept();
// Set up the streams associated with the socket
BufferedReader in = new BufferedReader(new
    InputStreamReader(clientSocket.getInputStream()));
PrintWriter out = new
    PrintWriter(clientSocket.getOutputStream(),true);

// Read string from client
String inputLine = in.readLine();
if (inputLine.equals("Hello")) {
    out.println("Connection established");
}
```

A Trivial Communication Protocol



- The server waits for a message from a client.
- If the string is "Hello", then "Nice to meet you" is sent to the client, otherwise "Connection refused" is sent to the client.

A Complete Client (1)

```
import java.net.*;
import java.io.*;

public class HelloClient {
    public static void main(String[] args) throws IOException {

        if (args.length != 2) {
            System.err.println("Usage: java HelloClient <host name> <port number>");
            System.exit(1);
        }

        String hostName = args[0];
        int portNumber = Integer.parseInt(args[1]);

        try (
            Socket myClientSocket = new Socket(hostName, portNumber);
            PrintWriter output = new
PrintWriter(myClientSocket.getOutputStream(),true);
            BufferedReader input = new BufferedReader(new
InputStreamReader(myClientSocket.getInputStream()));
            BufferedReader stdIn = new BufferedReader(new
InputStreamReader(System.in));
        ) {
            String userInput;
            System.out.println(input.readLine()); // reads the first message from the
server and prints it
        }
    }
}
```

A Complete Client (2)

```
System.out.println("Say something to the server (and then press enter): ");
userInput = stdIn.readLine(); // reads user's input
output.println(userInput); // user's input transmitted to server
System.out.println(input.readLine()); // reads server's ack and prints it
System.out.println("-----End of communication-----");
System.out.println("\nCommunication with server " + hostName + " was
successful! Now closing...");

} catch (UnknownHostException e) {
    System.err.println("Don't know about host " + hostName);
    System.exit(1);
} catch (IOException e) {
    System.err.println("Couldn't get I/O for the connection to " + hostName);
    e.printStackTrace();
    System.exit(1);
}
}
```

Its Corresponding Server (1)

```
import java.net.*;
import java.io.*;

public class HelloServer {
    public static void main(String[] args) throws IOException {

        int portNumber;

        if (args.length < 1) {
            System.out.println("Warning: You have provided no arguments\nTrying to connect
to the default port 8000...");
            portNumber = 8000;
        } else if (args.length == 1) {
            portNumber = Integer.parseInt(args[0]);
        } else {
            System.out.println("Warning: You have provided > 1 arguments\nTrying with the
first argument to connect to a port...");
            portNumber = Integer.parseInt(args[0]);
        }

        while(true){ //in order to serve multiple clients but sequentially, one after the other
            try (
                ServerSocket myServerSocket = new ServerSocket(portNumber);
                Socket aClientSocket = myServerSocket.accept();
                PrintWriter output = new PrintWriter(aClientSocket.getOutputStream(),true);
                BufferedReader input = new BufferedReader(new
```


Its Corresponding Server (2)

```
InputStreamReader(aClientSocket.getInputStream());
    ) {

        System.out.println("Connection established with a new client with IP address:
" + aClientSocket.getInetAddress() + "\n");
        output.println("Server says: Hello Client " + aClientSocket.getInetAddress() +
". This is server " + myServerSocket.getInetAddress() +
" speaking. Our connection has been successfully established!");

        String inputLine = input.readLine();
        System.out.println("Received a new message from client " +
aClientSocket.getInetAddress());
        System.out.println("Client says: " + inputLine);
        output.println("Server says: Your message has been successfully received!
Bye...");

        System.out.println("Connection with client " + aClientSocket.getInetAddress()
+ " is now closing...\n");

    } catch (IOException e) {
        System.out.println("Exception caught when trying to listen on port "
+ portNumber + " or listening for a connection");
        System.out.println(e.getMessage());
    }
}
}
```

Summary

- **Communication** is **indispensable** for distributed systems
- **Layered approach**
- Socket-based (transport layer) communication is available in all modern computer systems
 - Hard to develop large-scale applications at this level of abstraction
 - Higher-level development methods and APIs exist