

Lecture 6 - Linear Algebra, Scientific Python, and Decision Tree

Prof. Xiaowei Huang

<https://cgi.csc.liv.ac.uk/~xiaowei/>

(Attendance Code: **805871**)

- There will be no lectures in the next week, i.e., Week 3.
 - I have to travel to France for a project meeting.
- Labs will continue without interruption.

In the last
lectures, we
discussed

Probability distribution

- Joint probability, conditional probability,
etc

Two types of probabilistic queries

- Probability query
- MAP

Safety and Reliability Problems

Topics

Linear algebra

- Scalars, vectors, matrices, tensors
- Multiplying matrices/vectors

Introduction to Scientific Python

- Numpy
- Scipy
- Matplotlib

Decision Tree (1)

A large, semi-transparent blue circle is positioned on the left side of the slide, overlapping a dark blue rectangular background.

Linear Algebra For Machine Learning

Scalar

- Single number
- Represented in lower-case italic x
 - E.g., let $x \in \mathbb{R}$ be the slope of the line
 - Defining a real-valued scalar
 - E.g., let $n \in \mathbb{N}$ be the number of units
 - Defining a natural number scalar

Vector

- An array of numbers
- Arranged in order
- Each no. identified by an index
- Vectors are shown in lower-case bold
- If each element is in R then x is in R^n
- We think of vectors as points in space
 - Each element gives coordinate along an axis

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} \Rightarrow \mathbf{x}^T = [x_1, x_2, \dots, x_n]$$

Matrix

- 2-D array of numbers
- Each element identified by two indices
- Denoted by bold typeface \mathbf{A}
- Elements indicated as $A_{m,n}$
 - E.g.,
- $A[i:]$ is *i*th row of A , $A[:,j]$ is *j*th column of A
- If A has **shape** of height m and width n with real-values then

$$\mathbf{A} = \mathbb{R}^{m \times n}$$

$$\mathbf{A} = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}$$

Tensor



Tensor is not just an array of multiple dimensions.

- Sometimes need an array with more than two axes (or dimensions)
- An array arranged on a regular grid with variable number of axes is referred to as a tensor
- Denote a tensor with bold typeface: \mathbf{A}
- Element (i,j,k) of tensor denoted by $A_{i,j,k}$

Transpose of a Matrix

- Mirror image across principal diagonal

$$A = \begin{bmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{bmatrix} \Rightarrow A^T = \begin{bmatrix} x_{11} & x_{21} & x_{31} \\ x_{12} & x_{22} & x_{32} \\ x_{13} & x_{23} & x_{33} \end{bmatrix}$$

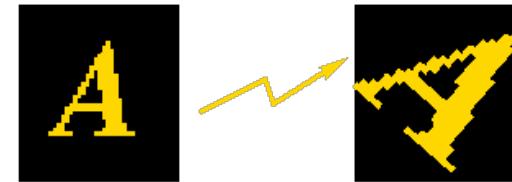
- Vectors are matrices with a single column

- Often written in-line using transpose

$$\mathbf{x} = [x_1, \dots, x_n]^T$$

- Since a scalar is a matrix with one element $a=a^T$

Linear Transformation



$$A\mathbf{x} = \mathbf{b}$$

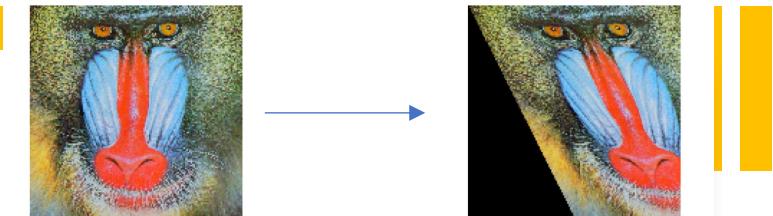
- where $A \in \mathbb{R}^{n \times n}$ and $\mathbf{b} \in \mathbb{R}^n$

$$A_{11}x_1 + A_{12}x_2 + \dots + A_{1n}x_n = b_1$$

$$A_{21}x_1 + A_{22}x_2 + \dots + A_{2n}x_n = b_2$$

...

$$A_{n1}x_1 + A_{n2}x_2 + \dots + A_{nn}x_n = b_n$$



n equations in
 n unknowns

Linear Transformation

$$A\mathbf{x} = \mathbf{b}$$

- where $A \in \mathbb{R}^{n \times n}$ and $\mathbf{b} \in \mathbb{R}^n$
- More explicitly

$$A = \begin{bmatrix} A_{1,1} & \cdots & A_{1,n} \\ \vdots & \vdots & \vdots \\ A_{n,1} & \cdots & A_{nn} \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix}$$

$n \times n \qquad n \times 1 \qquad n \times 1$

Can view A as a *linear transformation* of vector \mathbf{x} to vector \mathbf{b}

- Sometimes we wish to solve for the unknowns $x = \{x_1, \dots, x_n\}$ when A and \mathbf{b} provide constraints

Identity and Inverse Matrices

- Matrix inversion is a powerful tool to analytically solve $Ax=b$
- Needs concept of Identity matrix
- Identity matrix does not change value of vector
- when we multiply the vector by identity matrix
 - Denote identity matrix that preserves n-dimensional vectors as I_n
 - Formally $I_n \in \mathbb{R}^{n \times n}$ and $\forall \mathbf{x} \in \mathbb{R}^n, I_n \mathbf{x} = \mathbf{x}$
 - Example of I_3

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Matrix Inverse

- Inverse of square matrix A defined as $A^{-1}A=I_n$
- We can now solve $Ax=b$ as follows:

$$Ax=b$$

$$A^{-1}Ax = A^{-1}b$$

$$I_n x = A^{-1}b$$

$$x = A^{-1}b$$

- This depends on being able to find A^{-1}
- If A^{-1} exists there are several methods for finding it

Solving Simultaneous equations

- $\mathbf{Ax} = \mathbf{b}$
- Two closed-form solutions
 - Matrix inversion $\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$
 - Gaussian elimination

Norms

- Used for measuring the size of a vector
- Norms map vectors to non-negative values
- Norm of vector x is distance from origin to x
 - It is any function f that satisfies:

$$f(x) = 0 \Rightarrow x = 0$$

$$f(x+y) \leq f(x) + f(y) \quad \text{Triangle Inequality}$$

$$\forall \alpha \in \mathbb{R} \quad f(\alpha x) = |\alpha| f(x)$$

Image distance



-



=



$$\begin{bmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,32} \\ x_{2,1} & x_{22} & \dots & x_{2,32} \\ \vdots & \vdots & \ddots & \vdots \\ x_{32,1} & x_{32,2} & \dots & x_{32,32} \end{bmatrix} -$$

$$\begin{bmatrix} y_{1,1} & y_{1,2} & \dots & y_{1,32} \\ y_{2,1} & y_{22} & \dots & y_{2,32} \\ \vdots & \vdots & \ddots & \vdots \\ y_{32,1} & y_{32,2} & \dots & y_{32,32} \end{bmatrix}$$

$$= \begin{bmatrix} z_{1,1} & z_{1,2} & \dots & z_{1,32} \\ z_{2,1} & z_{22} & \dots & z_{2,32} \\ \vdots & \vdots & \ddots & \vdots \\ z_{32,1} & z_{32,2} & \dots & z_{32,32} \end{bmatrix}$$

L^P Norm

- Definition $\|x\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$

L^P Norm

- Definition $\|x\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$
- L^2 Norm
 - Called Euclidean norm, written simply as $\|x\|$
 - Squared Euclidean norm is same as $x^T x$

$$\|x\|_2 = \sqrt{\sum_i |x_i|^2}$$

$$= \sqrt{x^T x}$$

L^P Norm

- Definition $\|x\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$
- L^1 Norm
 - also called Manhattan distance

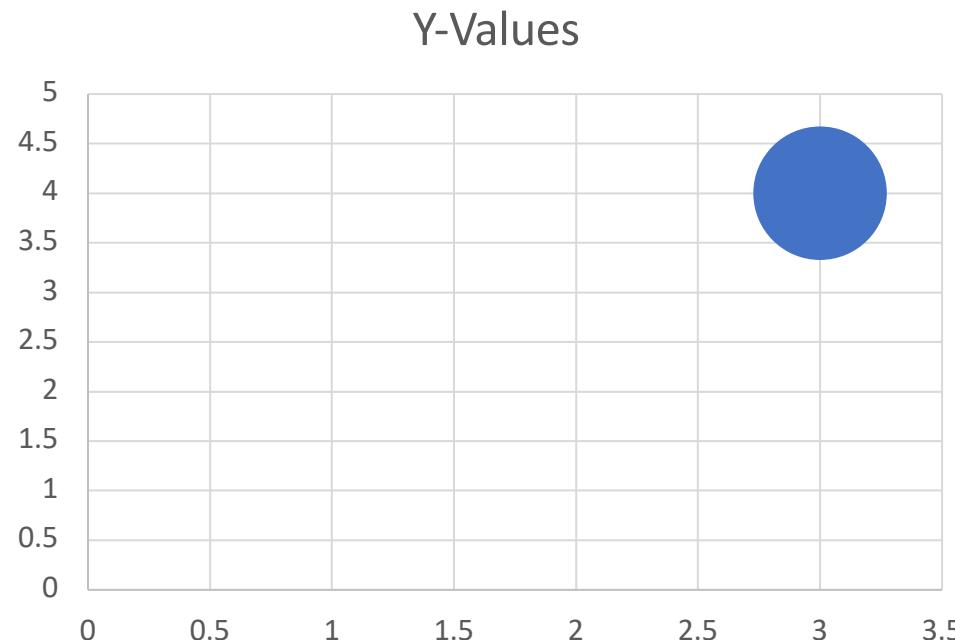
$$\|x\|_1 = \sum_i |x_i|$$

L^P Norm

- Definition $\|x\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$
- L^∞ Norm
 - also called max norm

$$\|x\|_\infty = \max_i |x_i|$$

Norms of two-dimensional Point



$$x = (3, 4)$$

$$\|x\|_1 = 3+4 = 7$$

$$\|x\|_2 = \sqrt{3^2 + 4^2} = 5$$

$$\|x\|_\infty = \max\{3, 4\} = 4$$

$$\|x\|_1 = \sum_i |x_i|$$

$$\|x\|_2 = \sqrt{\sum_i |x_i|^2}$$

$$\|x\|_\infty = \max_i |x_i|$$

Image distance

+

•

○



-



=



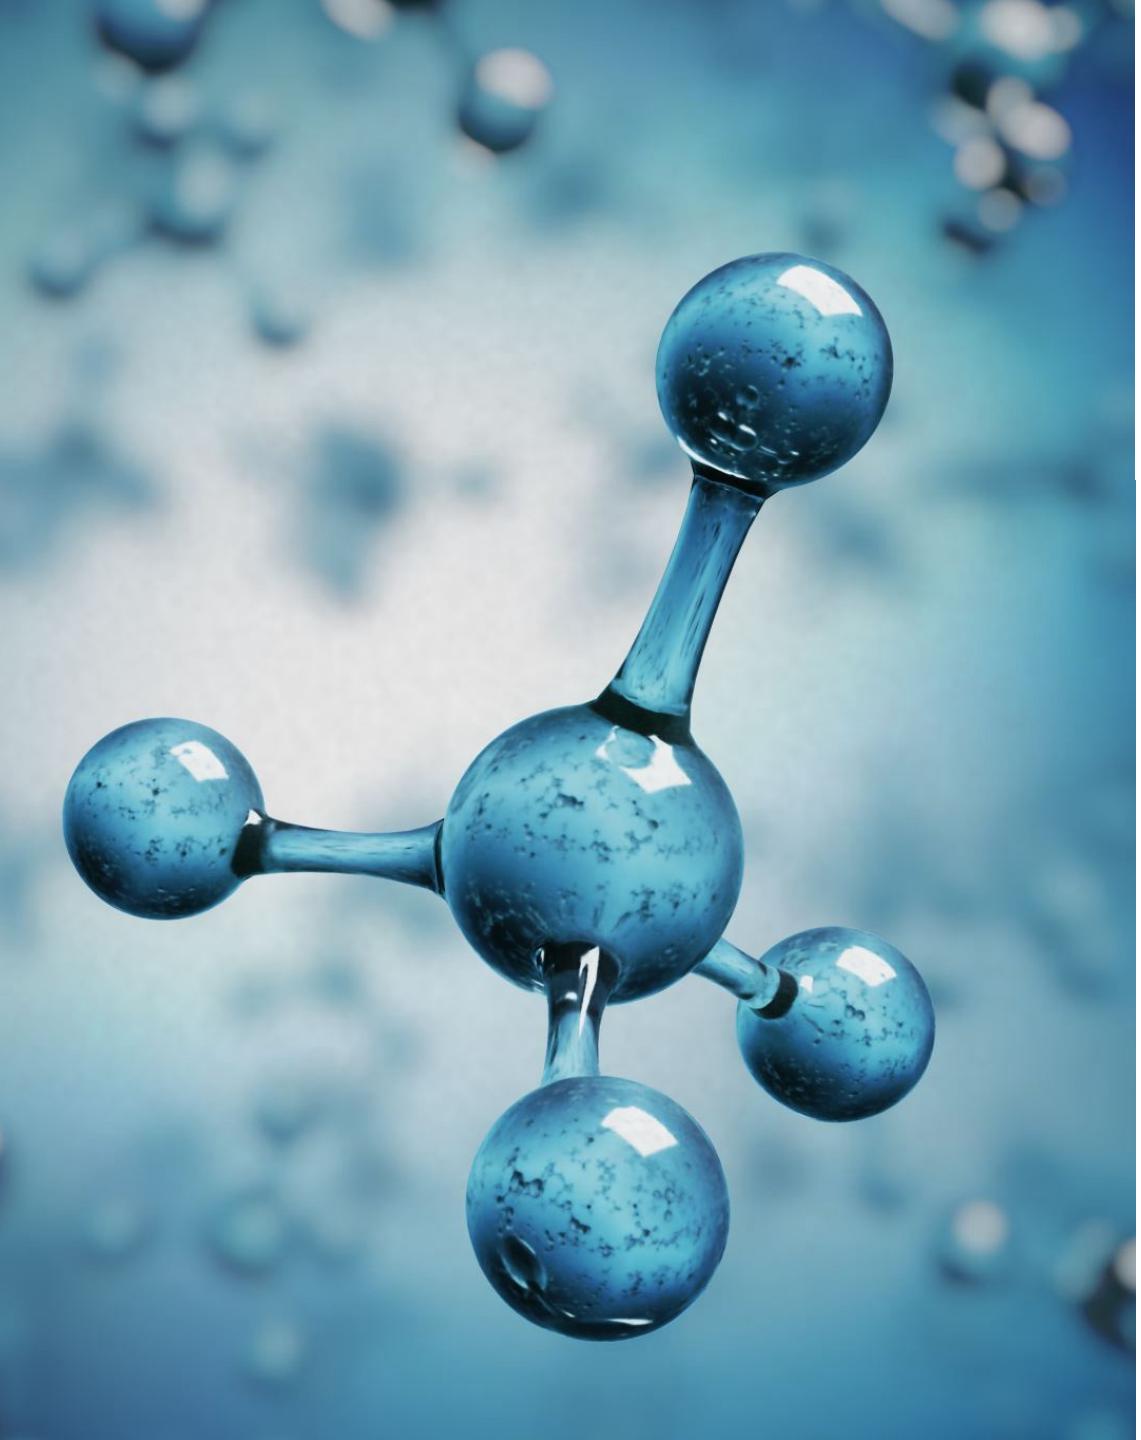
$$\begin{bmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,32} \\ x_{2,1} & x_{22} & \dots & x_{2,32} \\ \vdots & \vdots & \ddots & \vdots \\ x_{32,1} & x_{32,2} & \dots & x_{32,32} \end{bmatrix} - \begin{bmatrix} y_{1,1} & y_{1,2} & \dots & y_{1,32} \\ y_{2,1} & y_{22} & \dots & y_{2,32} \\ \vdots & \vdots & \ddots & \vdots \\ y_{32,1} & y_{32,2} & \dots & y_{32,32} \end{bmatrix} = \begin{bmatrix} z_{1,1} & z_{1,2} & \dots & z_{1,32} \\ z_{2,1} & z_{22} & \dots & z_{2,32} \\ \vdots & \vdots & \ddots & \vdots \\ z_{32,1} & z_{32,2} & \dots & z_{32,32} \end{bmatrix}$$

L^1 distance between X and Y: $\sum_{i,j} |z_{i,j}| = \sum_{i,j} |x_{i,j} - y_{i,j}|$

L^2 distance between X and Y: $\sqrt{\sum_{i,j} z_{i,j}^2} = \sqrt{\sum_{i,j} (x_{i,j} - y_{i,j})^2}$

L^∞ distance between X and Y:

$\max_{i,j} |z_{i,j}| = \max_{i,j} |x_{i,j} - y_{i,j}|$



Introduction to Scientific Python



Numpy

- Fundamental package for scientific computing with Python
- N-dimensional array object
- Linear algebra, Fourier transform, random number capabilities
- Building block for other packages (e.g. Scipy)
- Open source



import numpy as np

- Basics:

```
import numpy as np

A = np.array([[1, 2, 3], [4, 5, 6]])
print A
# [[1 2 3]
#  [4 5 6]]

Af = np.array([1, 2, 3], float)
```

- Slicing as usual

More basics

```
np.arange(0, 1, 0.2)
# array([ 0. ,  0.2,  0.4,  0.6,  0.8])

np.linspace(0, 2*np.pi, 4)
# array([ 0.0,  2.09,  4.18,  6.28])

A = np.zeros((2,3))
# array([[ 0.,  0.,  0.],
#        [ 0.,  0.,  0.]])
# np.ones, np.diag
A.shape
# (2, 3)
```

numpy.arange:
evenly spaced values within a given interval.

numpy.linspace:
evenly spaced numbers over a specified interval.

More basics

```
np.random.random((2,3))
# array([[ 0.78084261,  0.64328818,  0.55380341],
#        [ 0.24611092,  0.37011213,  0.83313416]])

a = np.random.normal(loc=1.0, scale=2.0, size=(2,2))
# array([[ 2.87799514,  0.6284259 ],
#        [ 3.10683164,  2.05324587]])

np.savetxt("a_out.txt", a)
# save to file
b = np.loadtxt("a_out.txt")
# read from file
```

- numpy.random.normal:
- Draw random samples from a normal (Gaussian) distribution
- loc: mean
- scale: standard deviation

Arrays are mutable

```
A = np.zeros((2, 2))
# array([[ 0.,  0.],
#        [ 0.,  0.]])
C = A
C[0, 0] = 1

print A
# [[ 1.  0.]
#  [ 0.  0.]]
```

Array attributes

```
a = np.arange(10).reshape((2,5))

a.ndim      # 2 dimension
a.shape     # (2, 5) shape of array
a.size      # 10 # of elements
a.T         # transpose
a.dtype     # data type
```

- numpy.reshape:
- Gives a new shape to an array without changing its data.

Basic operations

- Arithmetic operators: elementwise application

```
a = np.arange(4)
# array([0, 1, 2, 3])

b = np.array([2, 3, 2, 4])

a * b  # array([ 0,  3,  4, 12])
b - a  # array([2, 2, 0, 1])

c = [2, 3, 4, 5]
a * c  # array([ 0,  3,  8, 15])
```

- Also, we can use `+=` and `*=`.

Array broadcasting

- When operating on two arrays, numpy compares shapes. Two dimensions are compatible when
 - They are of equal size
 - One of them is 1

Array broadcasting

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 10 & 10 & 10 \\ \hline 20 & 20 & 20 \\ \hline 30 & 30 & 30 \\ \hline \end{array}
 +
 \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline \end{array}
 =
 \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 10 & 10 & 10 \\ \hline 20 & 20 & 20 \\ \hline 30 & 30 & 30 \\ \hline \end{array}
 +
 \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 10 & 10 & 10 \\ \hline 20 & 20 & 20 \\ \hline 30 & 30 & 30 \\ \hline \end{array}
 + \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 10 & 10 & 10 \\ \hline 20 & 20 & 20 \\ \hline 30 & 30 & 30 \\ \hline \end{array}
 + \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline 0 & 1 & 2 \\ \hline 0 & 1 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline 10 & 11 & 12 \\ \hline 20 & 21 & 22 \\ \hline 30 & 31 & 32 \\ \hline \end{array}$$

$$\begin{array}{c}
 \boxed{0} \\
 \boxed{10} \\
 \boxed{20} \\
 \boxed{30}
 \end{array}
 + \begin{array}{c}
 \boxed{0} \quad \boxed{1} \quad \boxed{2}
 \end{array}
 = \begin{array}{c}
 \boxed{0} \quad \boxed{0} \quad \boxed{0} \\
 \boxed{10} \quad \boxed{10} \quad \boxed{10} \\
 \boxed{20} \quad \boxed{20} \quad \boxed{20} \\
 \boxed{30} \quad \boxed{30} \quad \boxed{30}
 \end{array}
 + \begin{array}{c}
 \boxed{0} \quad \boxed{1} \quad \boxed{2} \\
 \boxed{0} \quad \boxed{1} \quad \boxed{2} \\
 \boxed{0} \quad \boxed{1} \quad \boxed{2}
 \end{array}$$

Array broadcasting with scalars

- This also allows us to add a constant to a matrix or multiply a matrix by a constant

```
A = np.ones((3,3))

print 3 * A - 1
# [[ 2.  2.  2.]
# [ 2.  2.  2.]
# [ 2.  2.  2.]]
```

Vector operations

- inner product
- outer product
- dot product (matrix multiplication)

```
# note: numpy automatically converts lists
u = [1, 2, 3]
v = [1, 1, 1]

np.inner(u, v)
# 6
np.outer(u, v)
# array([[1, 1, 1],
#        [2, 2, 2],
#        [3, 3, 3]])
np.dot(u, v)
# 6
```

Matrix operations

- First, define some matrices:

```
A = np.ones((3, 2))
# array([[ 1.,  1.],
#        [ 1.,  1.],
#        [ 1.,  1.]])
A.T
# array([[ 1.,  1.,  1.],
#        [ 1.,  1.,  1.]])  
  
B = np.ones((2, 3))
# array([[ 1.,  1.,  1.],
#        [ 1.,  1.,  1.]])
```

Matrix operations

- `np.dot(a,b)`
- If both a and b are 1-D arrays, it is inner product of vectors
- If both a and b are 2-D arrays, it is matrix multiplication

```
np.dot(A, B)
# array([[ 2.,  2.,  2.],
#        [ 2.,  2.,  2.],
#        [ 2.,  2.,  2.]])  
  
np.dot(B, A)
# array([[ 3.,  3.],
#        [ 3.,  3.]])  
  
np.dot(B.T, A.T)
# array([[ 2.,  2.,  2.],
#        [ 2.,  2.,  2.],
#        [ 2.,  2.,  2.]])  
  
np.dot(A, B.T)
# Traceback (most recent call last):
#   File "<stdin>", line 1, in <module>
# ValueError: shapes (3,2) and (3,2) not aligned: ...
# ... 2 (dim 1) != 3 (dim 0)
```

Operations along axes

```
a = np.random.random((2,3))
# array([[ 0.9190687 ,  0.36497813,  0.75644216],
#        [ 0.91938241,  0.08599547,  0.49544003]])
a.sum()
# 3.5413068994445549
a.sum(axis=0) # column sum
# array([ 1.83845111,  0.4509736 ,  1.25188219])
a.cumsum()
# array([ 0.9190687 ,  1.28404683,  2.04048899,  2.9598714 ,
#        3.04586687,  3.5413069 ])
a.cumsum(axis=1) # cumulative row sum
# array([[ 0.9190687 ,  1.28404683,  2.04048899],
#        [ 0.91938241,  1.00537788,  1.50081791]])
a.min()
# 0.0859954690403677
a.max(axis=0)
# array([ 0.91938241,  0.36497813,  0.75644216])
```

Slicing arrays

More advanced slicing

```
a = np.random.random((4,5))

a[2, :]
# third row, all columns
a[1:3]
# 2nd, 3rd row, all columns
a[:, 2:4]
# all rows, columns 3 and 4
```

- Iterating over multidimensional arrays is done with respect to the first axis: **for row in A**
- Looping over all elements: **for element in A.flat**

Iterating over arrays

- Reshape
 - using `reshape`. Total size must remain the same.
- Resize
 - using `resize`, always works: chopping or appending zeros
 - First dimension has ‘priority’, so beware of unexpected results

Reshaping

Linear algebra

import numpy.linalg	
qr	Computes the QR decomposition
cholesky	Computes the Cholesky decomposition
inv(A)	Inverse
solve(A,b)	Solves $Ax = b$ for A full rank
lstsq(A,b)	Solves $\arg \min_x \ Ax - b\ _2$
eig(A)	Eigenvalue decomposition
eig(A)	Eigenvalue decomposition for symmetric or hermitian
eigvals(A)	Computes eigenvalues.
svd(A, full)	Singular value decomposition
pinv(A)	Computes pseudo-inverse of A

Random sampling

import numpy.random	
rand(d0,d1,...,dn)	Random values in a given shape
randn(d0, d1, ...,dn)	Random standard normal
randint(lo, hi, size)	Random integers [lo, hi)
choice(a, size, repl, p)	Sample from a
shuffle(a)	Permutation (in-place)
permutation(a)	Permutation (new array)

Distributions in random

```
import numpy.random
```

The list of distributions to sample from is quite long, and includes

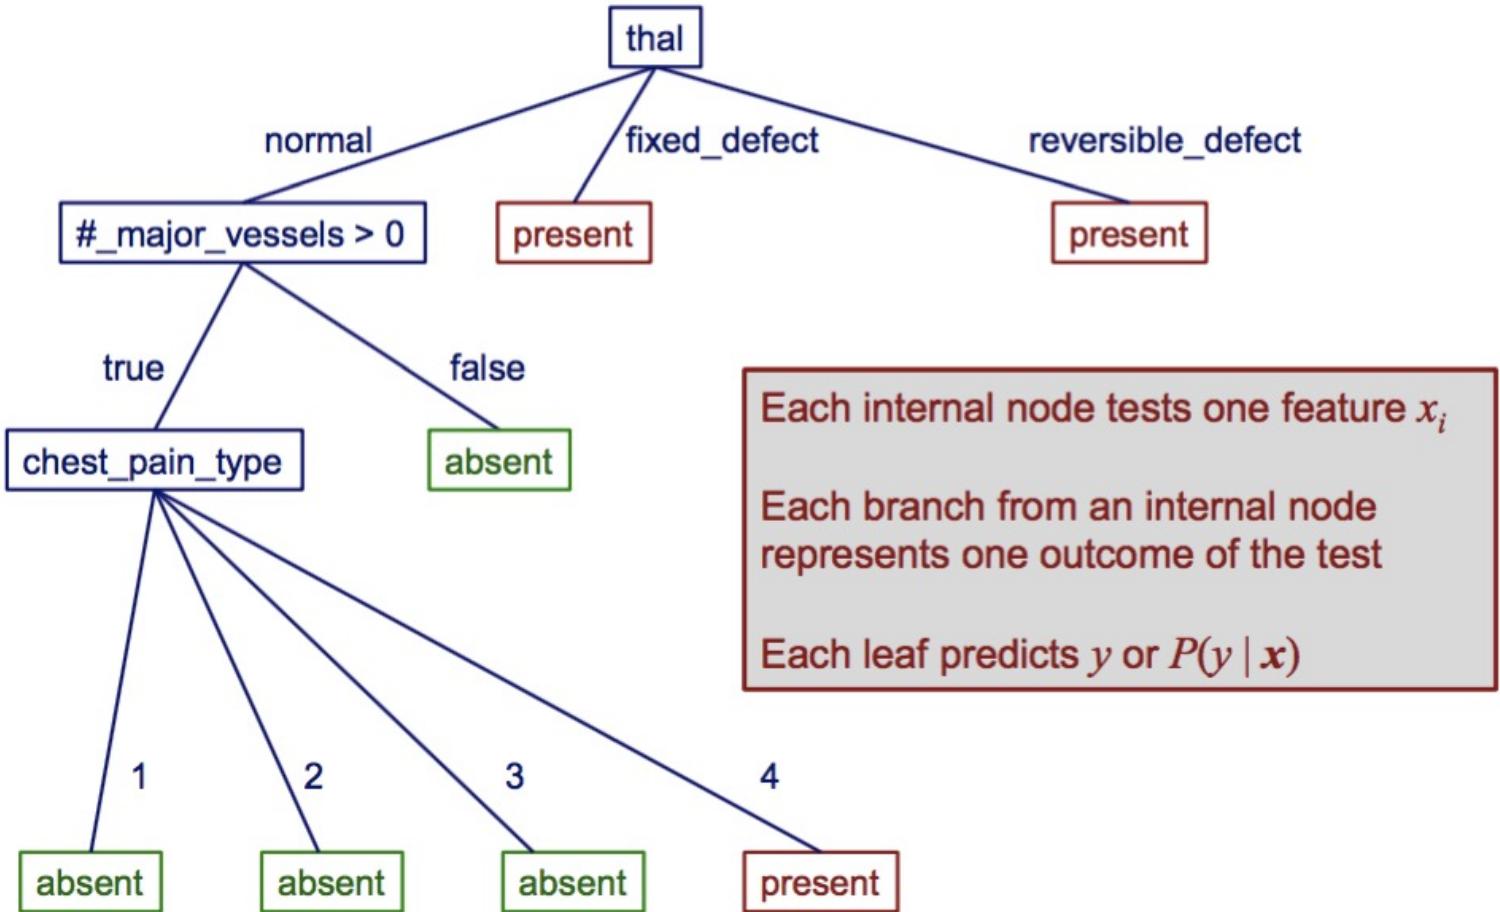
- beta
- binomial
- chisquare
- exponential
- dirichlet
- gamma
- laplace
- lognormal
- pareto
- poisson
- power

Recall: A learned decision tree

```
odor = a: e (400.0)
odor = c: p (192.0)
odor = f: p (2160.0)
odor = l: e (400.0)
odor = m: p (36.0)
odor = n
    spore-print-color = b: e (48.0)
    spore-print-color = h: e (48.0)
    spore-print-color = k: e (1296.0)
    spore-print-color = n: e (1344.0)
    spore-print-color = o: e (48.0)
    spore-print-color = r: p (72.0)
    spore-print-color = u: e (0.0)
    spore-print-color = w
        gill-size = b: e (528.0)
        gill-size = n
            gill-spacing = c: p (32.0) → if odor=almond, predict edible
            gill-spacing = d: e (0.0)
            gill-spacing = w
                population = a: e (0.0)
                population = c: p (16.0)
                population = n: e (0.0)
                population = s: e (0.0)
                population = v: e (48.0)
                population = y: e (0.0)
            spore-print-color = y: e (48.0)
odor = p: p (256.0)
odor = s: p (576.0)
odor = y: p (576.0)
```

if odor=none \wedge spore-print-color=white \wedge gill-size=narrow \wedge gill-spacing=crowded,
predict poisonous

A decision tree to predict heart disease



Decision tree exercise

- Suppose $X_1 \dots X_5$ are Boolean features, and Y is also Boolean
- How would you represent the following with decision trees?

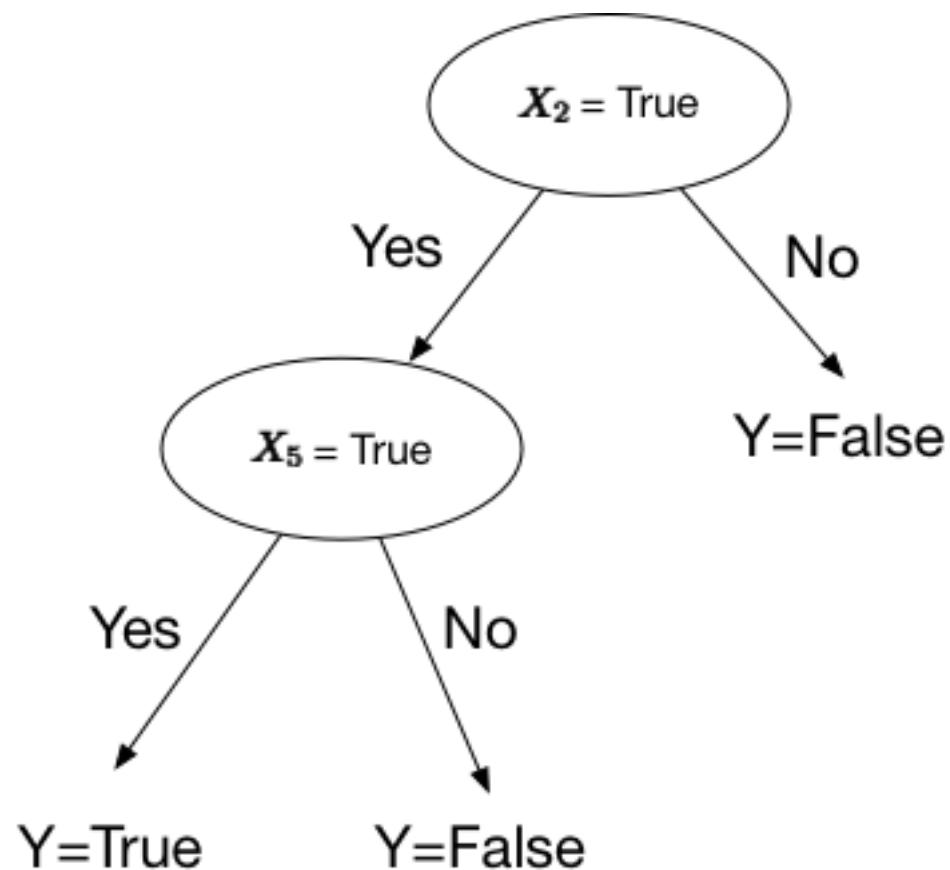
$$Y = X_2 X_5 \quad (\text{i.e., } Y = X_2 \wedge X_5)$$

$$Y = X_2 \vee X_5$$

$$Y = X_2 X_5 \vee X_3 \neg X_1$$

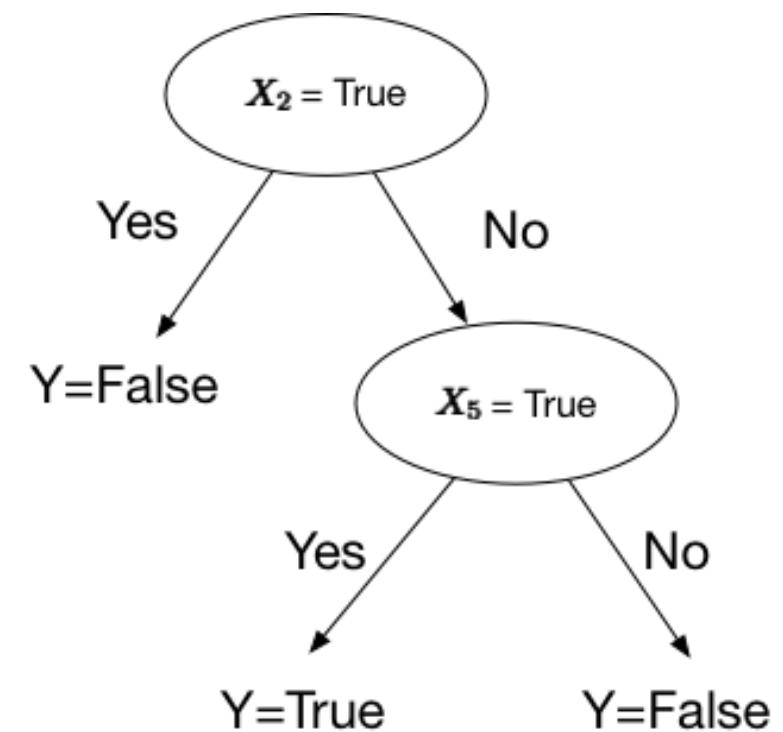
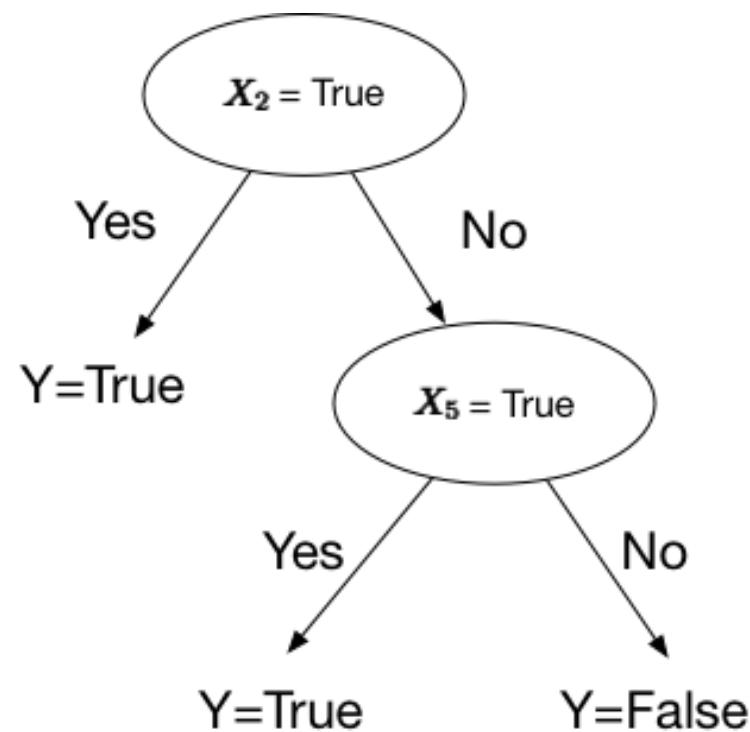
Decision tree exercise

$$Y = X_2 X_5$$



Decision tree exercise

$$Y = X_2 \vee X_5$$



Wrong!

Decision tree exercise

$$Y = X_2 X_5 \vee X_3 \neg X_1$$

