

last update: 28th February 2024

Session: 2023-24

COMP281 Lecture 10

# Principles of C and Memory Management

Phil Jimmieson



# Last Lecture

- struct, union, typedef, string

# Last Lecture

- struct, union, typedef, string

# Last Lecture

- struct, union, typedef, string

# Last Lecture

- struct, union, typedef, string

## Last Lecture

- struct, union, typedef, string

## Today's learning outcomes

- Storage Classes
- Dynamic Memory Allocation
- Stack and Heap

Storage classes

# Storage classes

Each variable has a storage class which decides the following things:

- **Scope:** where the value of the variable would be available inside a program.
- **Default initial value:** default initial value if NOT explicitly initialised.
- **Lifetime:** for how long will that variable exist.

The following storage classes are mostly used:

- **Automatic variables**
- **External variables**
- **Static variables**
- **Register variables**



# Automatic variables

Each variable has a storage class which decides the following things:

- **Scope:** local to the function block inside which they are defined.
- **Default initial value:** any random value, i.e. garbage value.
- **Lifetime:** till the end of the function/method block where the variable is defined.

```
{  
    int detail;          /* by default it's an 'auto' */  
    /* or */  
    auto int details;    /* Both are the same */  
    return 0;  
}
```

# External (Global) variables

Each variable has a storage class which decides the following things:

- **Scope:** not bound by any function; available everywhere.
- **Default initial value:** 0(zero).
- **Lifetime:** Until the program finishes executing.



Global values can be changed by any function in the program!

## Example 13

```
#include <stdio.h>
int number; /* global variable */

void fun1() {
    number = 20;
    printf("I am in function fun1. My value is %d.\n", number);
}

void fun2() {
    printf("I am in function fun2. My value is %d.\n", number);
}

int main() {
    number = 10;
    printf("I am in main function. My value is %d\n", number);
    fun1(); /* function calling, discussed in next topic */
    fun2(); /* function calling, discussed in next topic */
    return 0;
}
```

## Example 13

```
#include <stdio.h>
int number; /* global variable */

void fun1() {
    number = 20;
    printf("I am in function fun1. My value is %d.\n", number);
}

void fun2() {
    printf("I am in function fun2. My value is %d.\n", number);
}

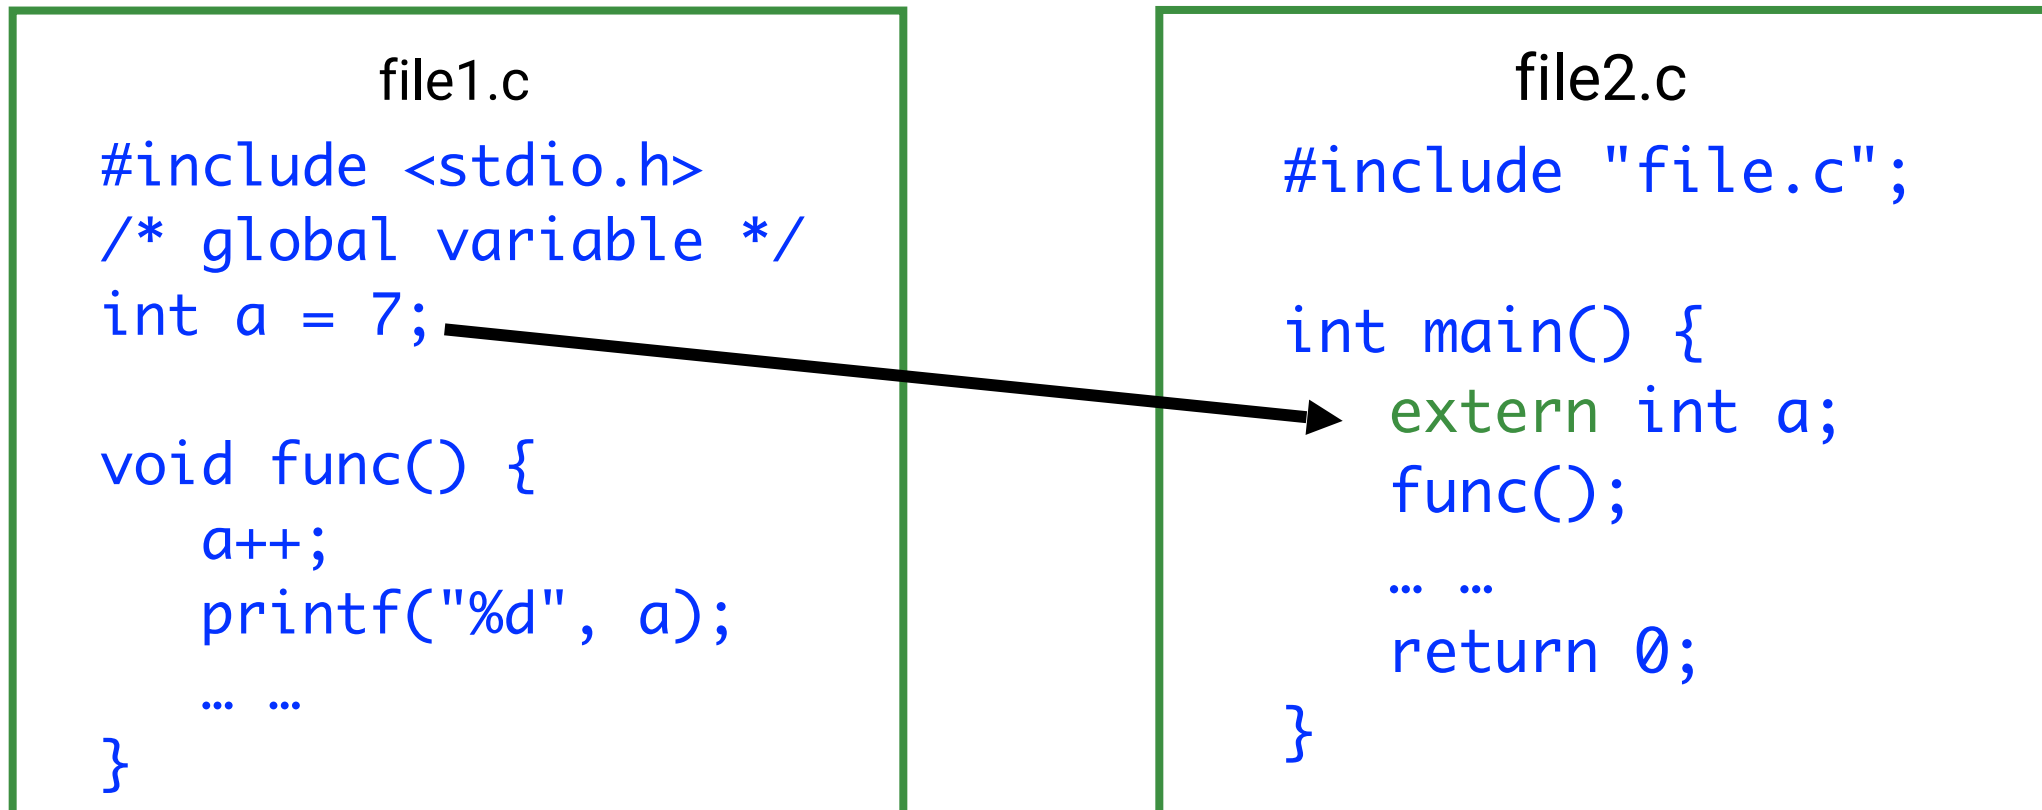
int main() {
    number = 10;
    printf("I am in main function. My value is %d\n", number);
    fun1(); /* function calling, 'number' is global variable */
    fun2(); /* function calling, 'number' is global variable */
    return 0;
}
```

**Output**

```
I am in main function. My value is 10
I am in function fun1. My value is 20.
I am in function fun2. My value is 20.
```

# Keyword extern

The `extern` keyword is used with a variable to inform the compiler that this variable is declared somewhere else. The extern declaration does not allocate storage for variables.







## Problem when extern is not used

```
int main() {  
    a = 10; /* Error: cannot find definition of variable 'a' */  
    printf("%d", a);  
    return 0;  
}
```

## Example using extern in same file

```
int main() {  
    extern int x; /* informs the compiler that it is defined  
                  somewhere else */  
    x = 10;  
    printf("%d", x);  
}  
int x; //Global variable x
```

# Static variables

Tells the compiler to persist/save the variable until the end of program

- **Scope:** local to the block in which the variable is defined.
- **Default initial value:** 0(zero).
- **Lifetime:** Until the program finishes executing.

Instead of creating and destroying a variable every time when it comes into and goes out of scope, a static variable is initialised only once and remains in existence until the end of the program.

## Example 14

```
#include <stdio.h>

void test();

int main() {
    test();
    test();
    test();
    return 0;
}

void test() {
    static int a = 0; /* a static variable */
    a = a + 1;
    printf("%d\t",a);
}
```

**Output**

1      2      3

# Register variables

Tells the compiler to store the variable in a CPU register

- **Scope:** local to the function in which it is declared.
- **Default initial value:** Any random value, i.e. garbage value.
- **Lifetime:** Until the end of function/method block, in which the variable is defined.

Syntax

```
register int number;
```



- Register variables inform the compiler to store the variable in **CPU** register instead of **memory**.
- Register variables have faster accessibility than a normal variable.
- We cannot get the address of such variables.

Which storage class should be used and when?



To improve the speed of execution of the program and to carefully use the memory space occupied by the variables:

- Use **static** storage class only when we want the value of the variable to remain same every time we call it using different function calls.
- Use **register** storage class only for those variables that are used in our program very often. CPU registers are limited and thus should be used carefully.
- Use **external (global)** storage class only for those variables that are being used by almost all the functions in the program.
- If we do not have the purpose of any of the above mentioned storage classes, then we should use the **automatic** storage class.

# Memory Management in C

C manages memory

- Statically
- Automatically
- Dynamically

## Static-duration Variables

- Allocated in the **main memory**
- Usually along with the executable code of the program
- Persist for the lifetime of the program

## Automatic-duration Variables

- Allocated on the **stack**
- Come and go as functions are called and return

# Static-duration Variables & Automatic-duration Variables

The size of the allocation must be **compile-time constant**.

Static-duration Variables  
&  
Automatic-duration Variables

The size of the allocation must be  
compile-time constant.

# Flexibility Concerns



## **Static-duration Variables**

persist for the life of the program whether needed or not

---

## **Automatic-duration Variables**

cannot persist across multiple function calls

# Lifetime Concerns

## Static-duration Variables

persist for the life of the program whether needed or not

---

## Automatic-duration Variables

cannot persist across multiple function calls

How to avoid these limitations?

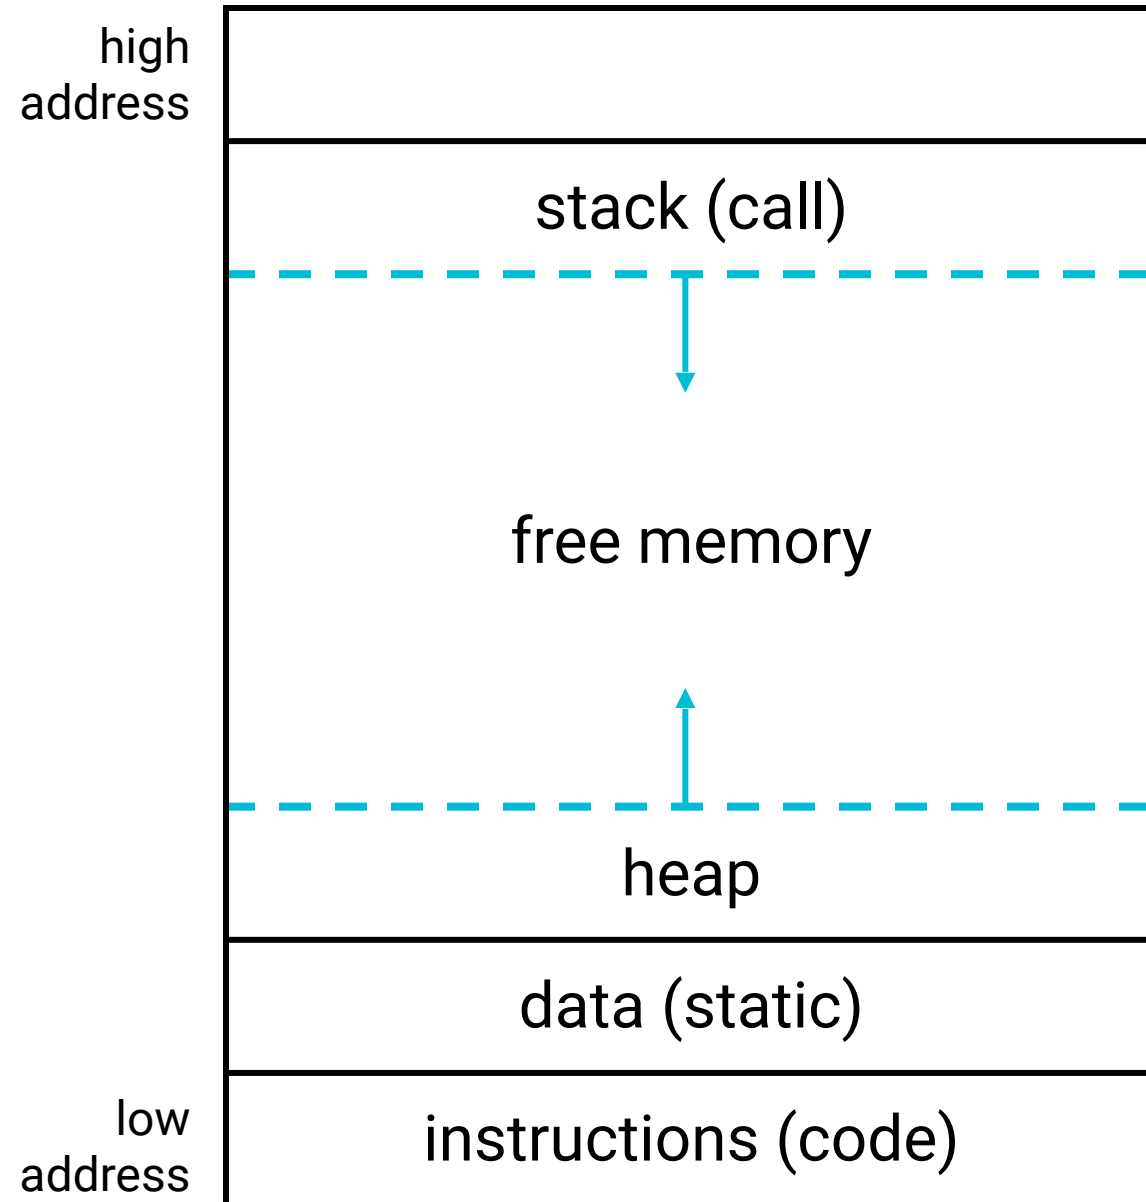
# Dynamic Memory Allocation



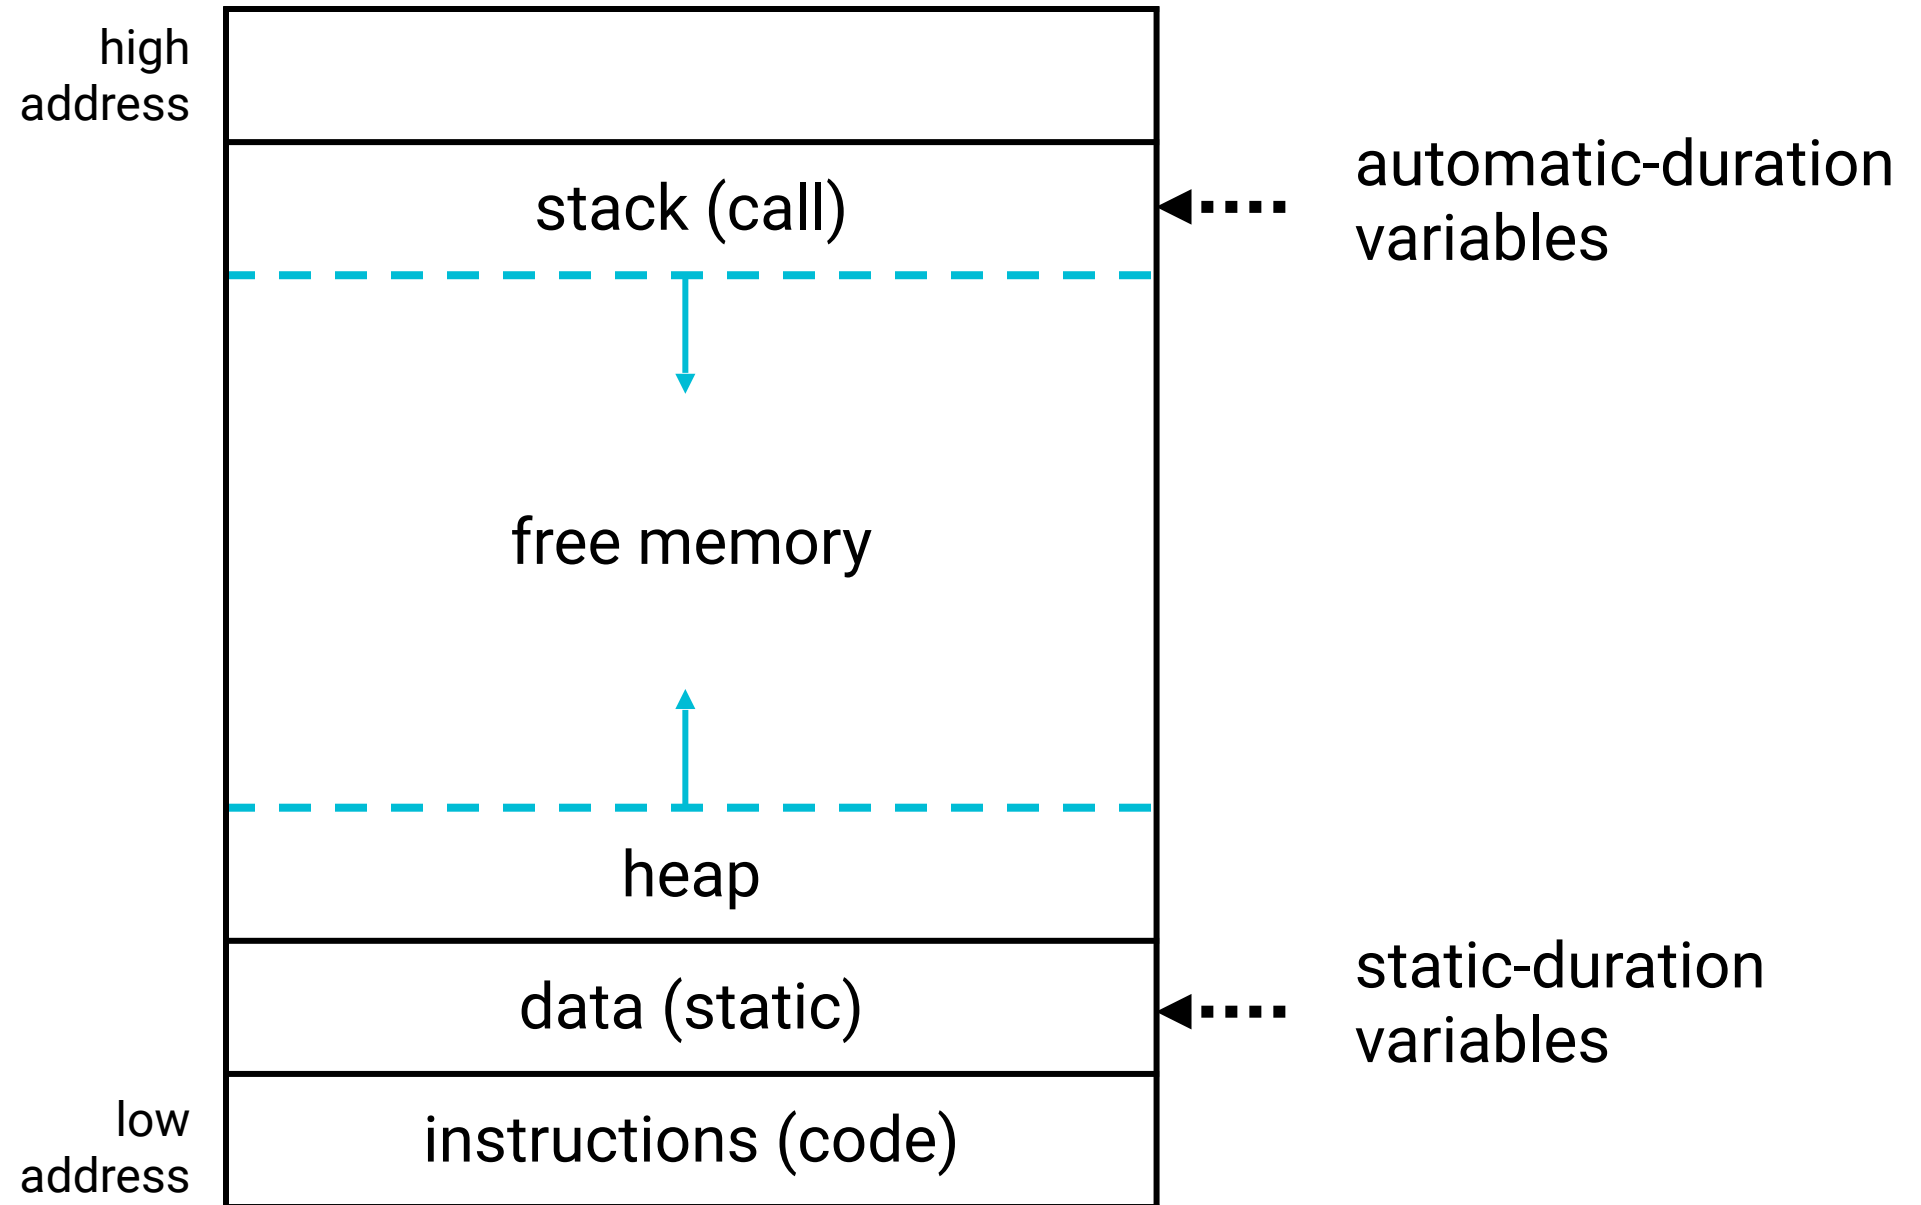
(size known at **run-time**)

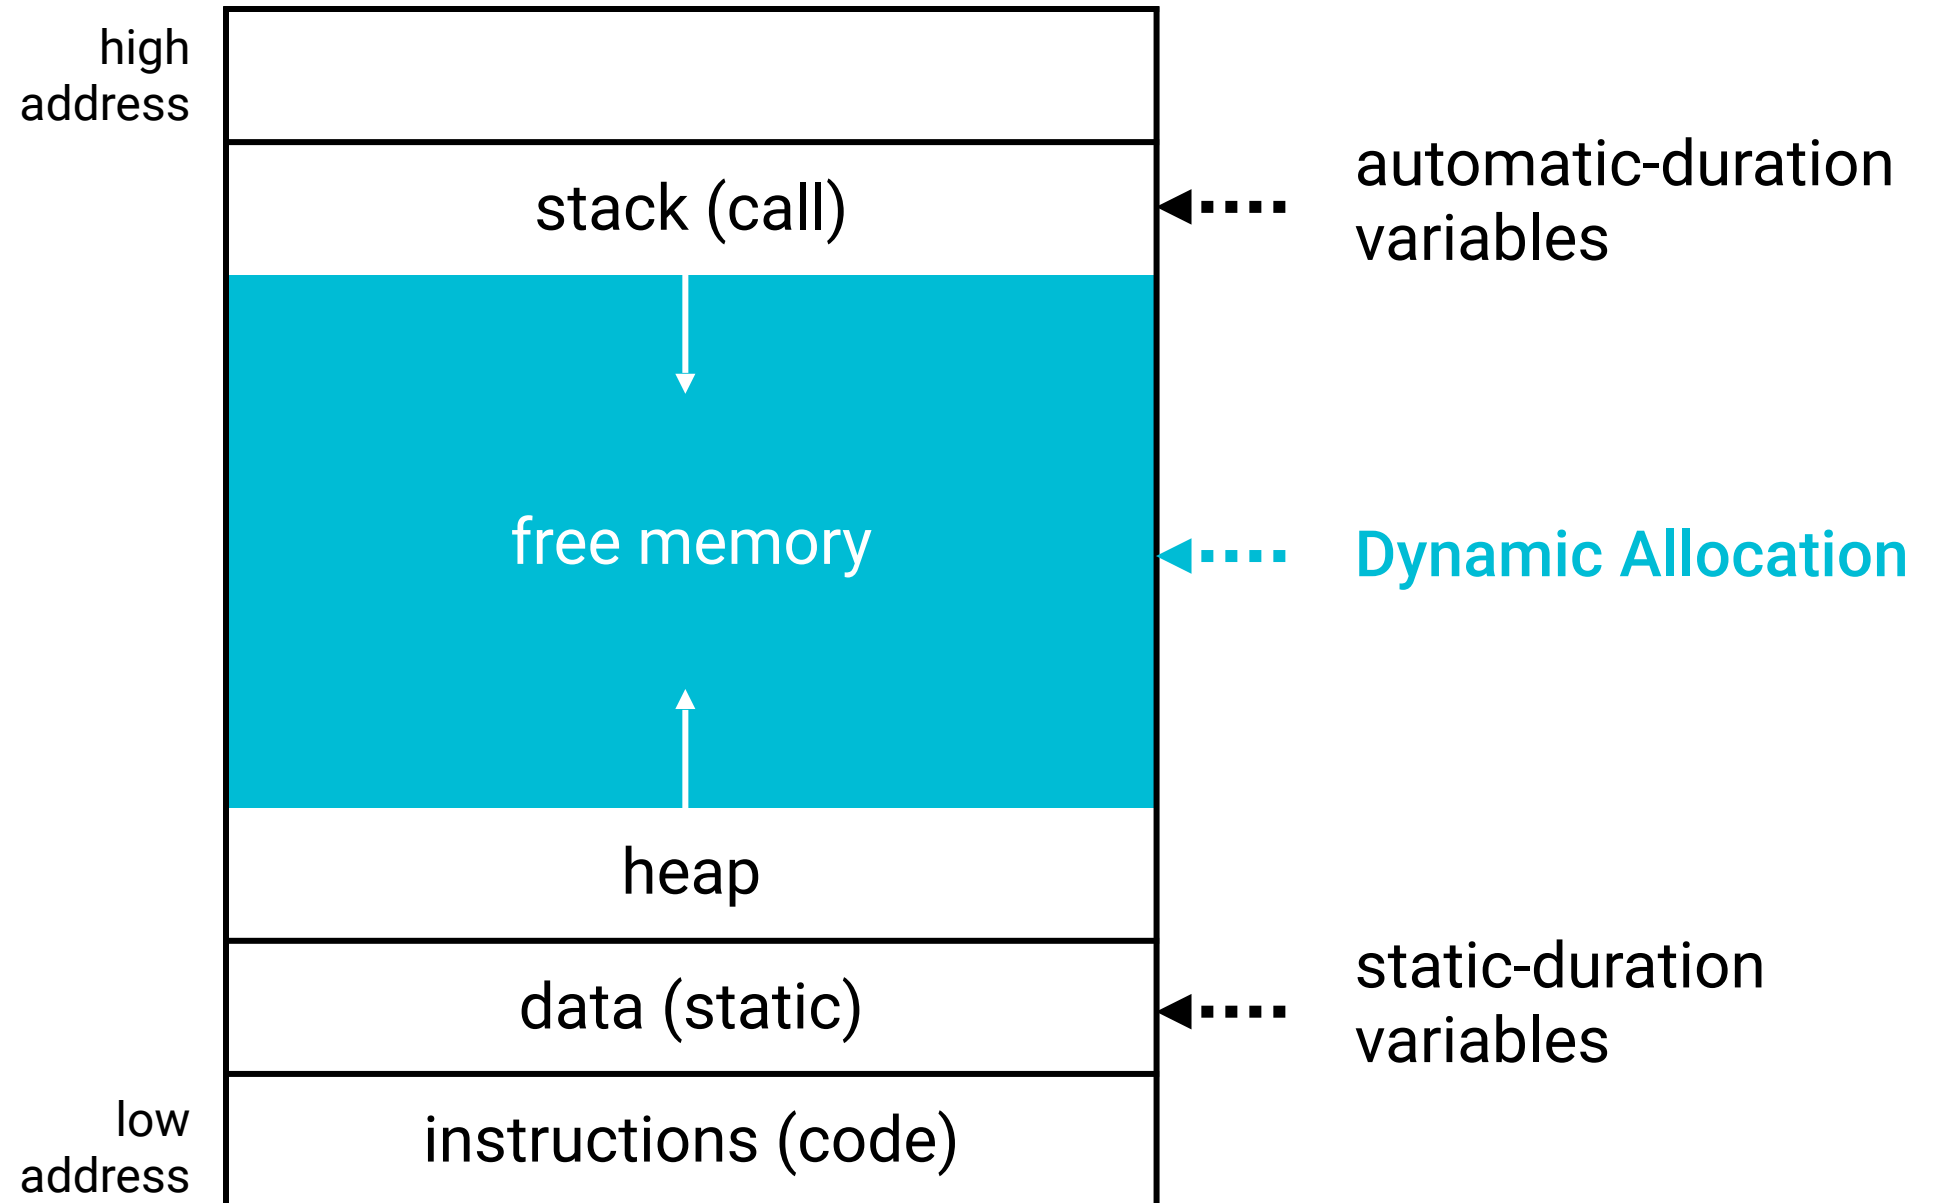
as opposed to

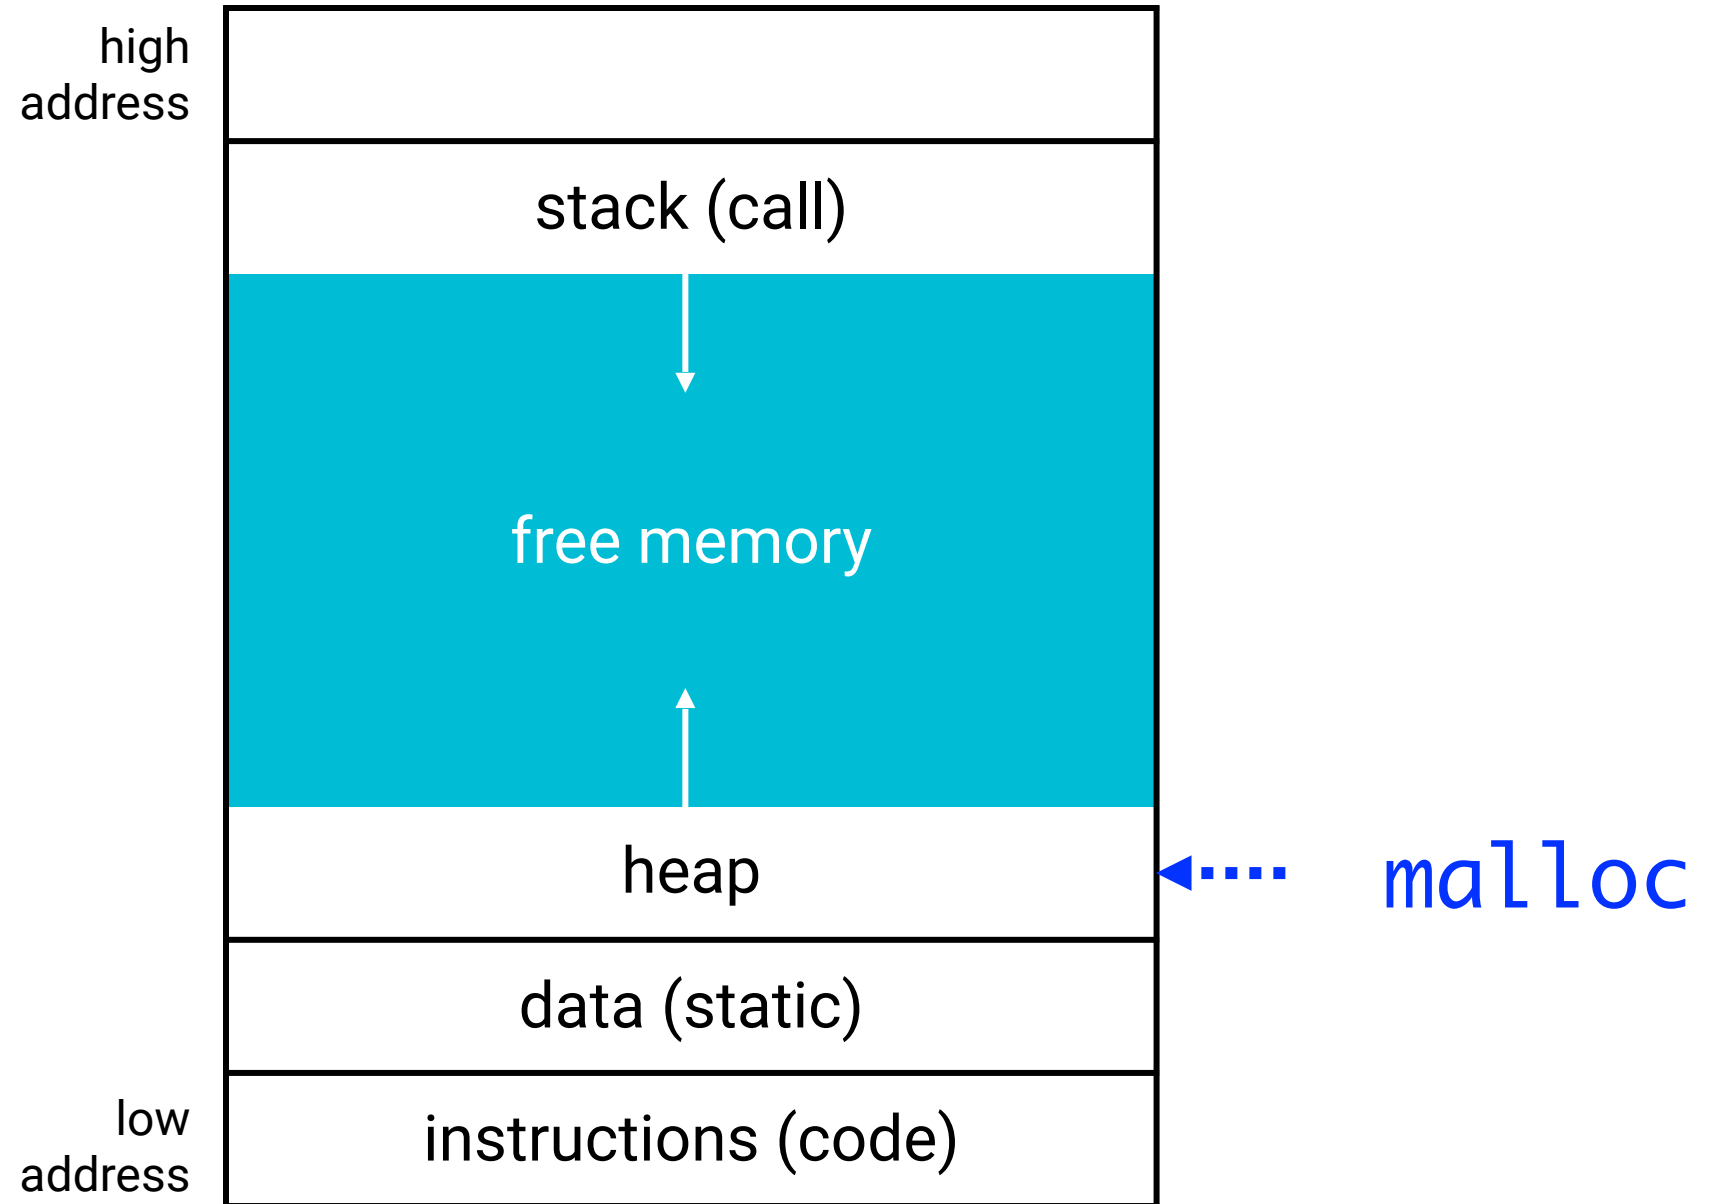
(size known at **compile-time**)











# malloc

```
void *malloc(size_t size);
```

- **malloc** stands for "memory **alloc**ator".
- **malloc** is used to allocate a block of memory on the heap.
- **malloc** takes 1 argument - the size, in bytes, of the chunk of memory to be allocated



Return a **void \*** ? what does **void \*** mean?

- It's a "pointer to anything"
- Actual type either doesn't matter or will be given later by a type cast



Return a **void \*** ? what does **void \*** mean?

- It's a "pointer to anything"
- Actual type either doesn't matter or will be given later by a **type cast**

# malloc

`malloc` is often used for dynamically allocated arrays

e.g., to dynamically allocate an array of 10 `ints`:

```
int *arr;
```

```
arr = (int *)malloc(10 * sizeof(int));
```


# malloc

`malloc` is often used for dynamically allocated arrays

e.g., to dynamically allocate an array of 10 `ints`:

```
int *arr;
```

```
arr = (int *)malloc(10 * sizeof(int));
```





# malloc

`malloc` is often used for dynamically allocated arrays

e.g., to dynamically allocate an array of 10 `ints`:

```
int *arr;
```

```
arr = (int *)malloc(10 * sizeof(int));
```

This computes the number of bytes that 10 `ints` occupy in memory, then requests that many bytes from `malloc` and assigns the result to an int pointer named `arr`

# calloc

```
void *calloc(size_t nmemb, size_t size);
```

- **calloc** is a variant of **malloc**
- **calloc** takes **2 arguments** – the number of "things" to be allocated and the size of each "thing" (in bytes).
- **calloc** returns a pointer pointing to the chunk of memory that was allocated.

# calloc

```
void *calloc(size_t nmemb, size_t size);
```

- `calloc` is a variant of `malloc`
- `calloc` takes **2 arguments** – the number of "things" to be allocated and the size of each "thing" (in bytes).
- `calloc` returns a pointer pointing to the chunk of memory that was allocated.
- `calloc` also sets each of the values in the allocated memory to zero (`malloc` doesn't).

# calloc

`calloc` is also used for dynamically allocated arrays

e.g., to dynamically allocate an array of 10 `ints`:

```
int *arr;
```

```
arr = (int *)calloc(10, sizeof(int));
```

# calloc

`calloc` is also used to dynamically allocated arrays

e.g., to dynamically allocate an array of 10 `ints`:

```
int *arr;
```

```
arr = (int *)calloc(10, sizeof(int));
```

The diagram illustrates the mapping between the code above and the `calloc` function signature below. It consists of three horizontal lines with arrows pointing down to the function signature. The first arrow points from the `(int *)` cast to the `void *` return type. The second arrow points from the `10` argument to the `size_t nmemb` parameter. The third arrow points from the `sizeof(int)` argument to the `size_t size` parameter.

```
void *calloc(size_t nmemb, size_t size);
```

# calloc

`calloc` is also used to dynamically allocated arrays

e.g., to dynamically allocate an array of 10 `ints`:

```
int *arr;
```

```
arr = (int *)calloc(10, sizeof(int));
```

vs

```
arr = (int *)malloc(10 * sizeof(int));
```



```
void *malloc(size_t size);
```

vs

```
void *calloc(size_t nmemb, size_t size);
```



	malloc	calloc
argument	1 argument - the mount of memory to allocate in bytes	2 arguments - the number of variables to allocate; and the size in bytes of a single variable.
initialisation	No	0 (zero)

## malloc/calloc return value

`malloc` and `calloc` both return **a pointer** pointing to the newly-allocated chunk of memory (**the address**).



## malloc/calloc return value

`malloc` and `calloc` both return a pointer pointing to the newly-allocated chunk of memory (the address).

**BUT!**

Not guaranteed to succeed.

Maybe there is not enough memory available 😞

What to do if they fail?

## malloc/calloc return value

If they fail, they return **NULL**

We must always check for **NULL**  
when using **malloc** or **calloc**

## malloc/calloc return value

```
int *arr = (int *)malloc(10 * sizeof(int));  
/* code that uses arr... */
```

## malloc/calloc return value



```
int *arr = (int *)malloc(10 * sizeof(int));  
/* code that uses arr... */
```



```
int *arr = (int *)malloc(10 * sizeof(int));  
if (arr == NULL) {  
    fprintf(stderr, "out of memory!\n");  
    exit(1);  
}  
/* code that uses arr... */
```

What if the program no longer needs  
the dynamic array?

# free

When the program no longer needs the dynamic array, we should call `free` to return the memory it occupies to the system.

(this is also known as "deallocating" the memory)

```
free( arr );
```

When we free `arr`, we can no longer use the array (of course...)

# free

- When we free some memory, the memory is NOT erased or destroyed, instead, the OS is informed that we no longer need it, so it may use it for e.g. another program or some other variables in our program.
- Trying to use memory after freeing it can cause a segmentation violation (when a program is trying to read or write an illegal memory location → program crash).




# free

```
void *foo(int n) {  
    int i = 7;  
    int *j;  
    j = (int *)malloc(n * sizeof(int));  
    /* j = (int *)calloc(n, sizeof(int)); */  
    return j;  
} /* memory storing the variables i and j deallocated  
here; block of memory pointed to by j is not */
```

# free

```
void bar() {  
    int *arr = foo(99); // foo creates the array  
    arr[0] = 2;  
    arr[1] = 4;  
    arr[2] = 16;  
    ...  
    free(arr); /* deallocate memory */  
}
```

# free

```
void bar() {  
    int *arr = foo(99);  
    arr[0] = 2;  
    arr[1] = 4;  
    arr[2] = 16;  
    ...  
      
    

---

    free(arr); /* deallocate memory */  
}
```

# Memory Leaks

- Failure to deallocate memory using **free** leads to build-up of **non-reusable** memory, which is no longer used by the program.
- This **wastes** memory resources and can lead to allocation failures when these resources are **exhausted**.

# Memory Leaks

```
void leaker() {  
    int *arr = (int *)malloc (10 * sizeof(int));  
    /* Now have allocated space for 10 ints;  
     * do something with it and return without  
     * calling free() */  
} /* arr memory is leaked here. */
```

After `leaker()` returns, nothing points to the memory allocated in the function -> memory leak 😞

# Memory Leaks

```
void not_leaker() {  
    int *arr = (int *)malloc (10 * sizeof(int));  
    /* Now have allocated space for 10 ints;  
     * do something with it */  
    free(arr); /* free arr's memory */  
} /* no memory leaked here. */
```

Here, we explicitly `free()` the memory allocated by `malloc()` before exiting the function -> no memory leak 😊

# Memory Leaks

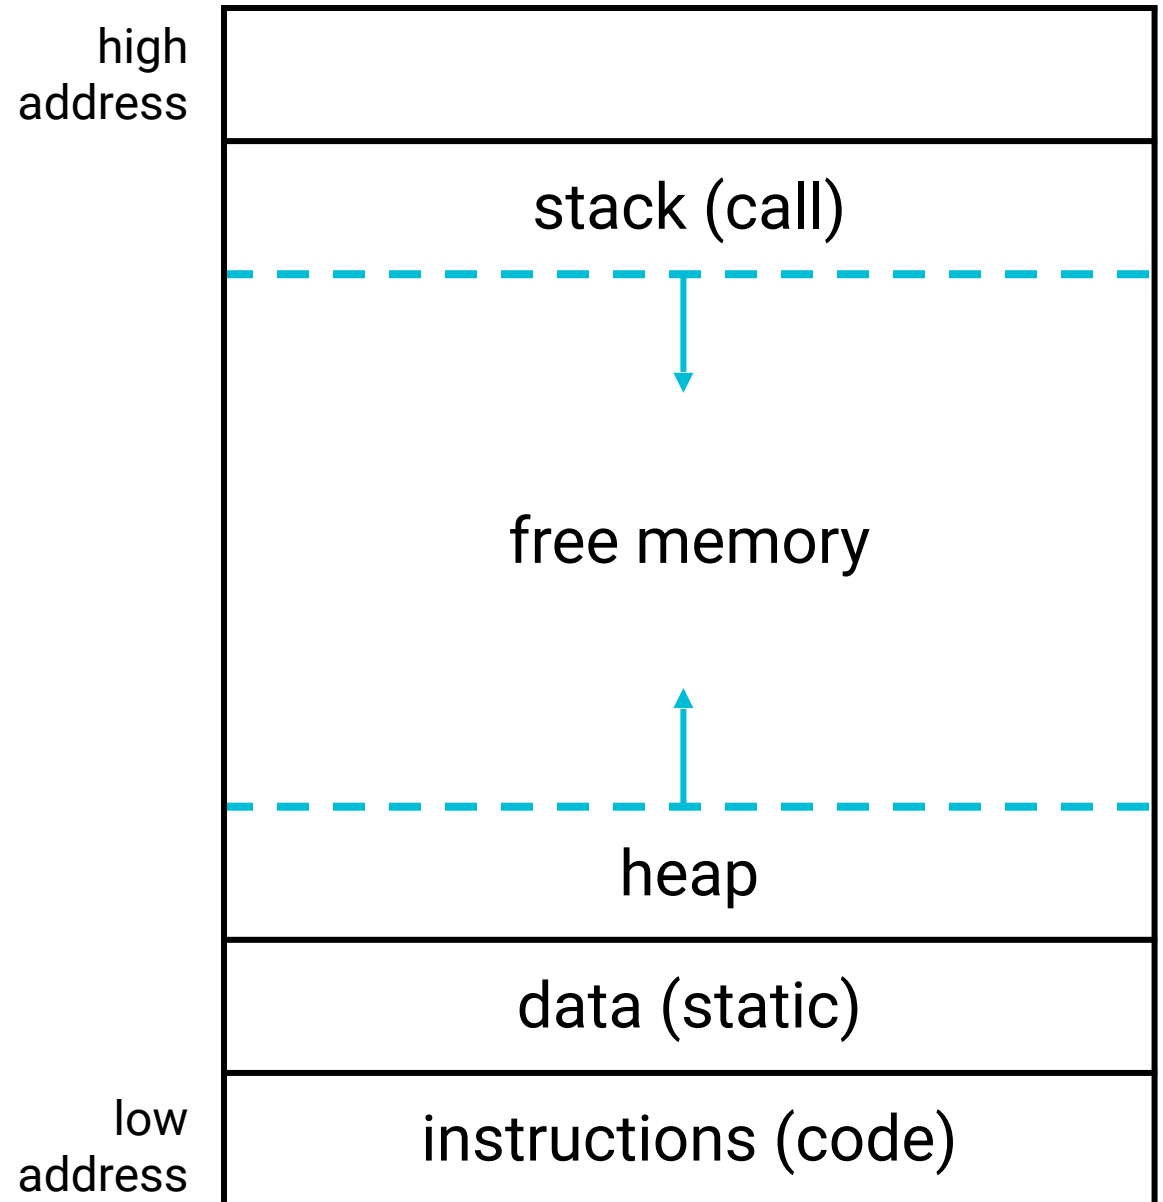
```
void not_leaker_2() {  
    int arr[10];  
    /* Now have allocated space for 10 ints;  
     * do something with it */  
} /* no memory leaked here. */
```

Here, we don't need to `free()` the memory, because it was allocated locally (on the "stack") -> no memory leak 😊

# Memory Leaks

```
void not_leaker_2() {  
    int arr[10];  
    /* Now have allocated space  
     * do something with it  
    */  
} /* no memory leaked here
```

Here, we don't need to `free()`  
allocated locally (on the "stack"





# Memory Leaks

```
void crasher() {  
    int arr[10];  
    /* Now have allocated space for 10 ints;  
     * do something with it */  
    free(arr); /* shocking! */  
}
```

Here, we `free()` the memory we don't need to free  
-> anything can happen 💣

## Example 1

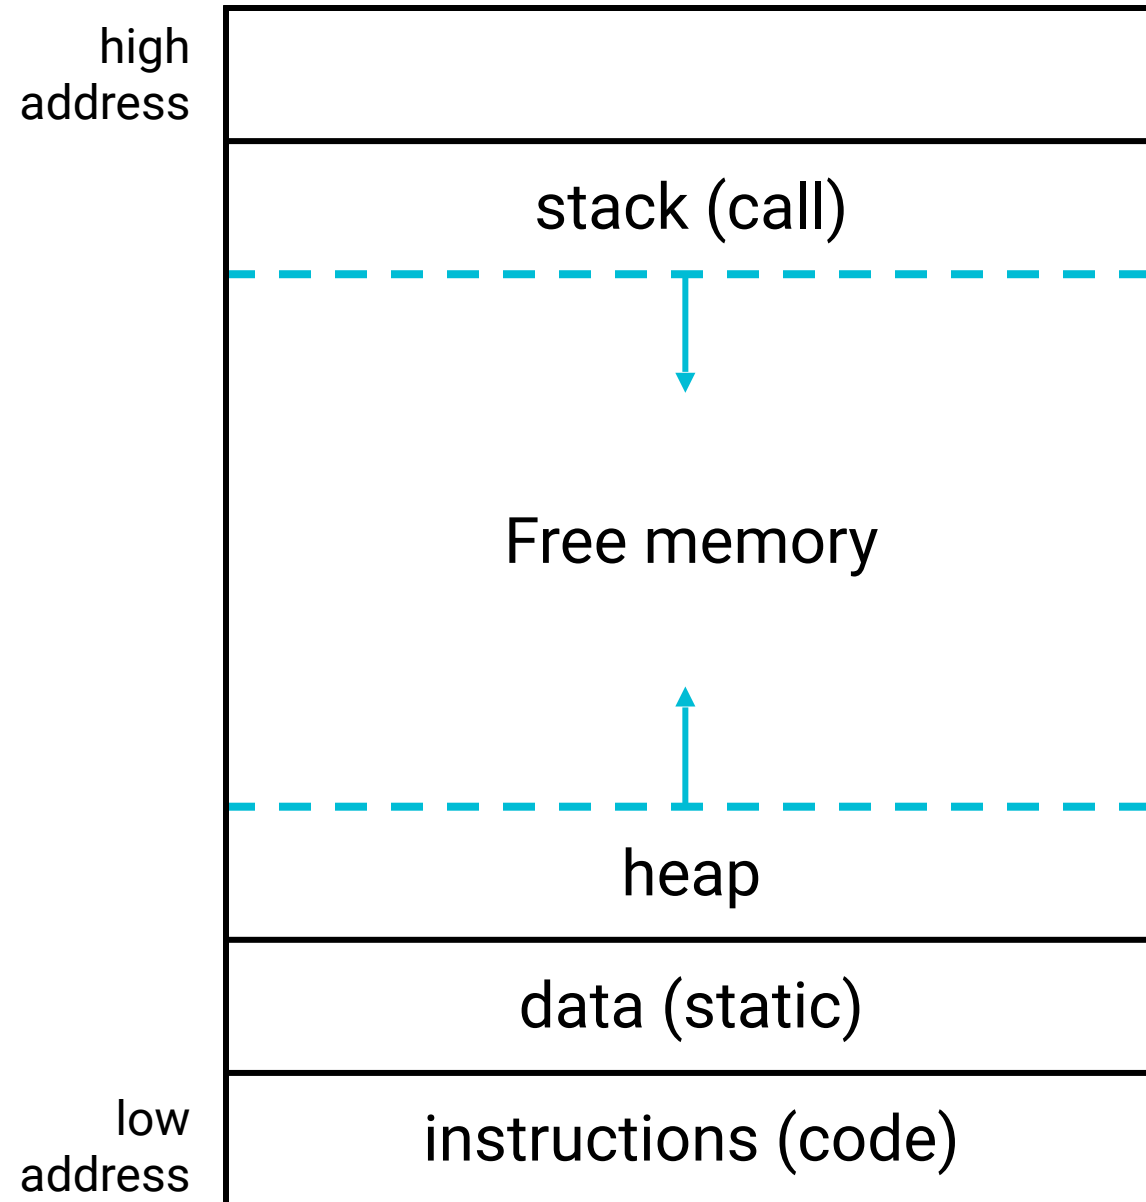
```
#include <stdio.h>
#include <stdlib.h>
int main() {
    int i, n;
    int *item;
    printf("Enter total number of items: ");
    scanf("%d", &n);
    item = (int*)calloc(n, sizeof(int));
    if (item == NULL) {
        printf("Error. Not enough space available");
        exit(1);
    }
    for (i = 0; i < n; i++) {
        scanf("%d", item+i);
    }
    for (i = 1; i < n; i++) {
        if (*item > *(item+i)) {
            *item = *(item+i);
        }
    }
    printf("Smallest item is %d\n", *item);
    free(item);
    return 0;
}
```

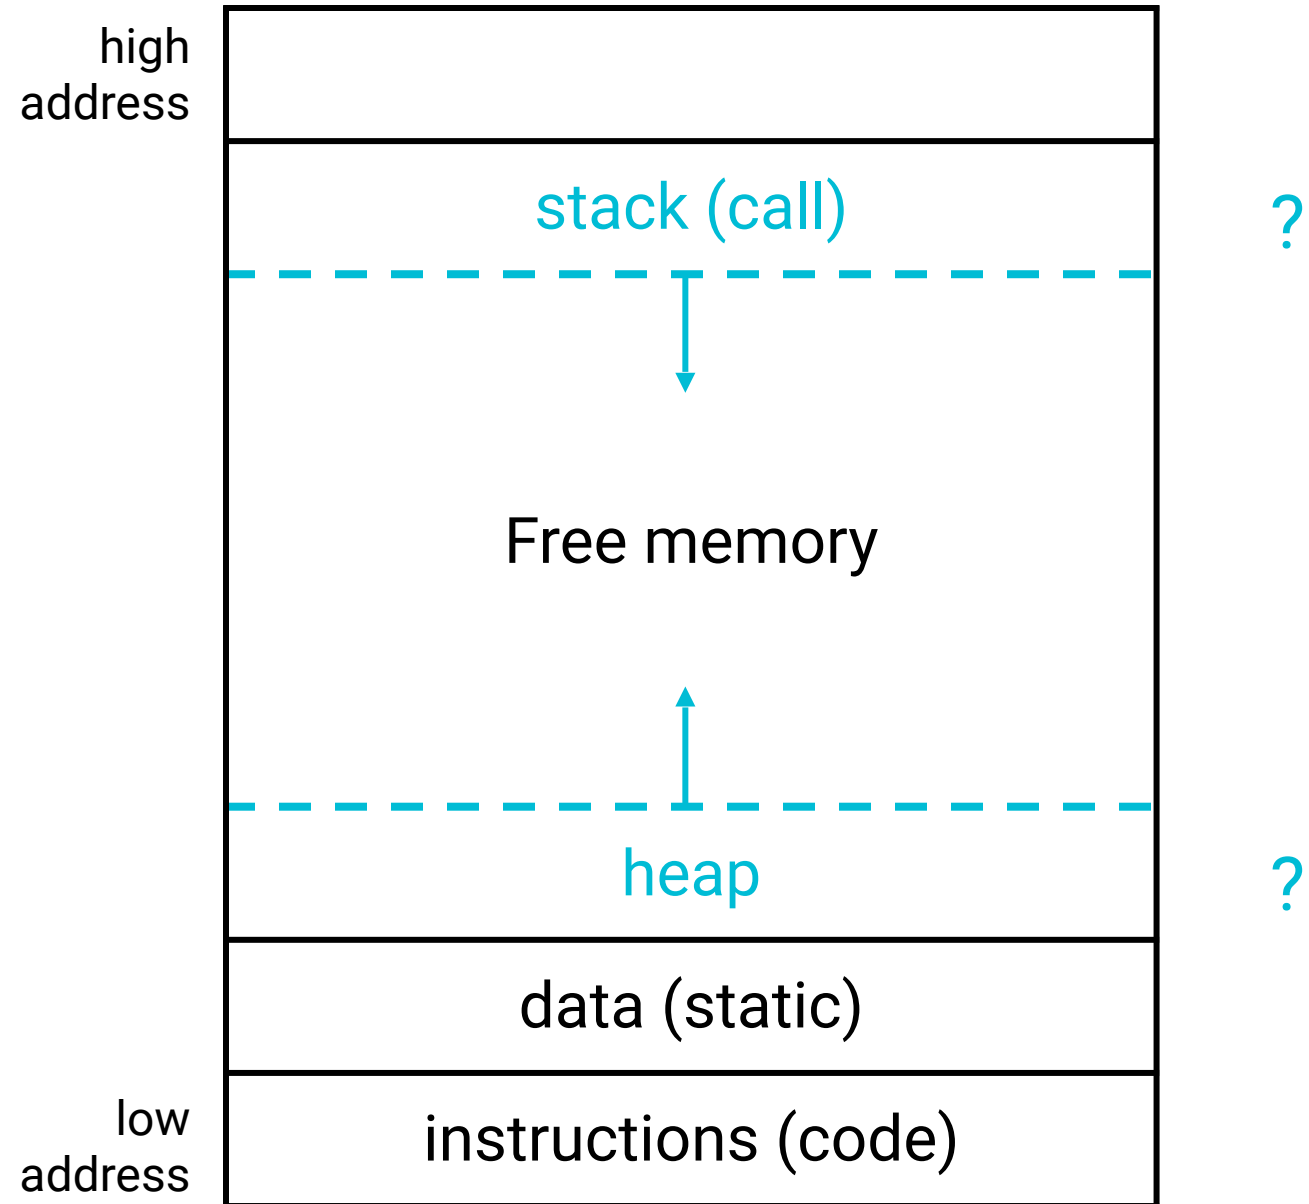
## Example 1

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    int i, n;
    int *item;
    printf("Enter total number of items: ");
    scanf("%d", &n);
    item = (int*)calloc(n, sizeof(int));
    if (item == NULL) {
        printf("Error. Not enough space available");
        exit(1);
    }
    for (i = 0; i < n; i++) {
        scanf("%d", item+i);
    }
    for (i = 1; i < n; i++) {
        if (*item > *(item+i)) {
            *item = *(item+i);
        }
    }
    printf("Smallest item is %d\n", *item);
    free(item);
    return 0;
}
```

## Output


```
Enter total number of items: 5
5 2 7 3 9
Smallest item is 2
```





The stack

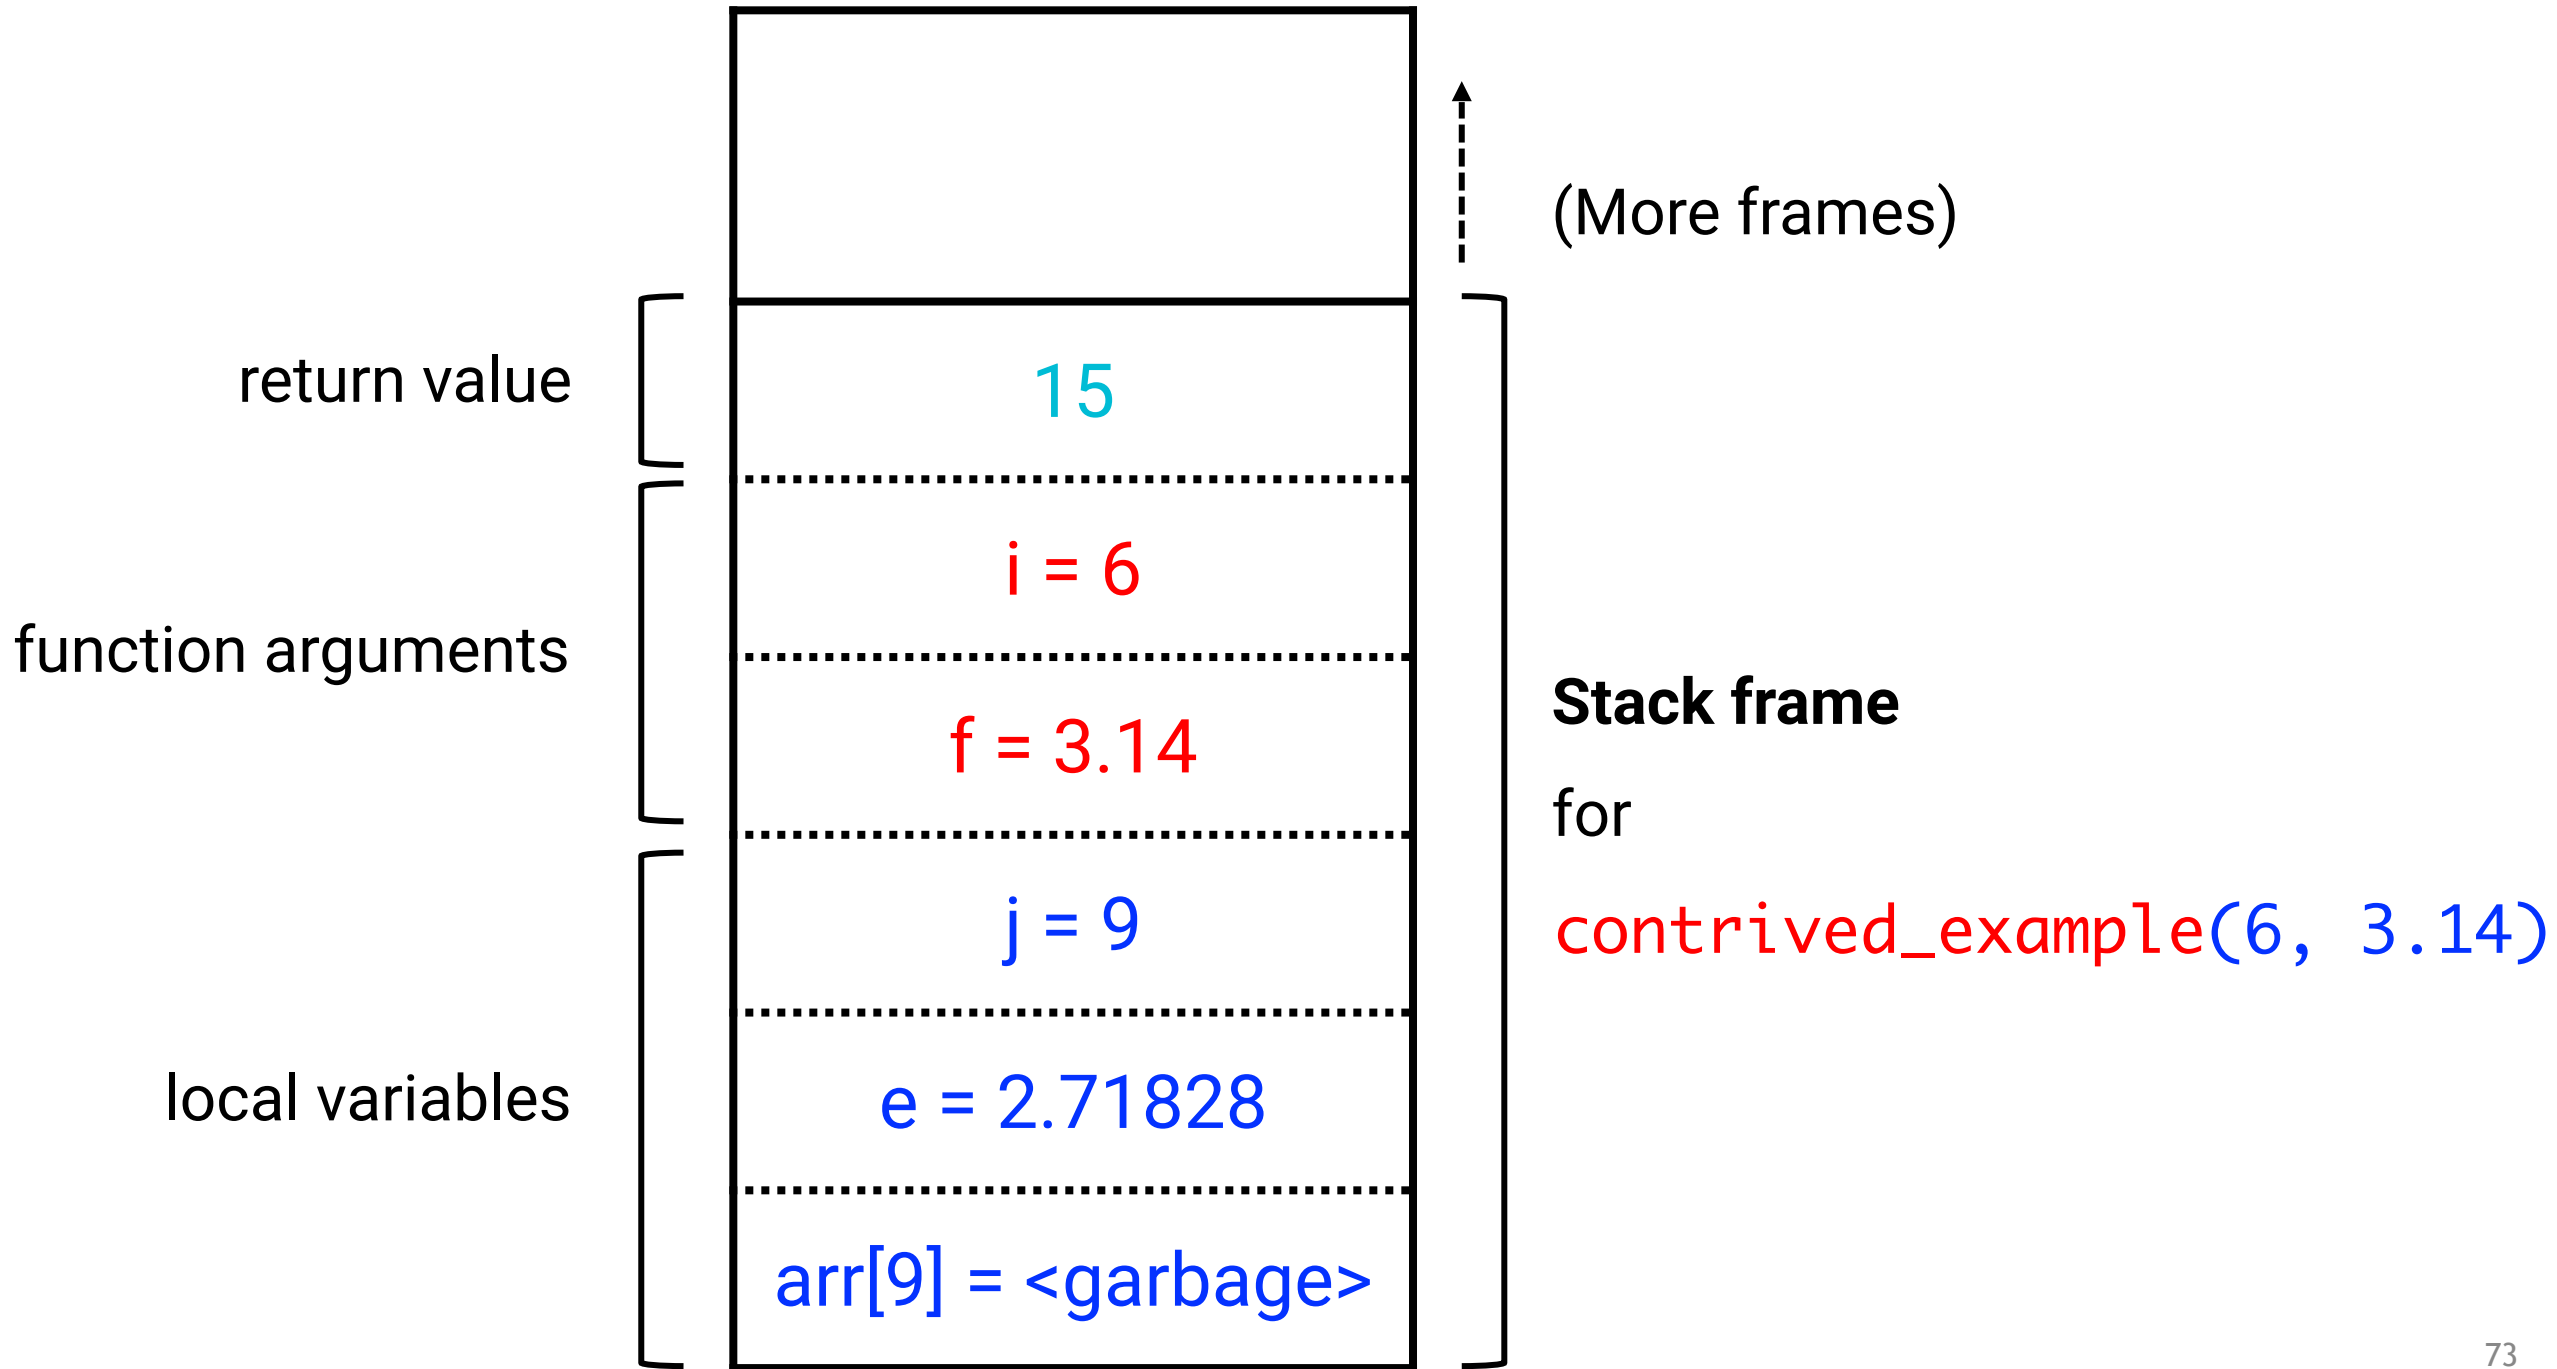
- Local variables, function arguments, return values are all stored in a **stack**.
- Each function call generates a new "**stack frame**".
- After a function returns, the stack frame disappears, along with all local variables and function arguments for that invocation.



# The Stack

```
int contrived_example(int i, float f) {  
    int j = 9;  
    double e = 2.71828;  
    int arr[9];  
    /* do something there, then return */  
    return (i + j);  
}  
  
/* somewhere in the code */  
  
int k = contrived_example(6, 3.14);
```





# The Stack

```
int factorial(int i) {  
    if (i == 0) {  
        return 1;  
    } else {  
        return i * factorial(i - 1);  
    }  
}
```

# The Stack

```
int factorial(int i) {  
    if (i == 0) {  
        return 1;  
    } else {  
        return i * factorial(i - 1);  
    }  
}
```

---

# The Stack

- What goes on the **stack** for **factorial(3)**?
- For each stack frame...
  - no local variables
  - one argument **i**
  - one return value
- Each recursive call generates a new stack frame
  - which disappears after the call is complete and the recursion unwinds.

```
int factorial(int i) {  
    if (i == 0) {  
        return 1;  
    } else {  
        return i * factorial(i - 1);  
    }  
}
```

factorial(3)

return value

?

function argument

i = 3


Stack frame

factorial(2)

return value  
function argument

factorial(3)

return value  
function argument



Stack frame

Stack frame

factorial(1)

return value

?

function argument

$i = 1$

factorial(2)

return value

?

function argument

$i = 2$

factorial(3)

return value

?

function argument

$i = 3$

Stack frame

Stack frame

Stack frame

factorial(0)

return value

?

function argument

$i = 0$

Stack frame

factorial(1)

return value

?

function argument

$i = 1$

Stack frame

factorial(2)

return value

?

function argument

$i = 2$

Stack frame

factorial(3)

return value

?

function argument

$i = 3$


Stack frame



factorial(0)

return value

1



function argument

$i = 0$

Stack frame

factorial(1)

return value

?

function argument

$i = 1$

Stack frame

factorial(2)

return value

?

function argument

$i = 2$

Stack frame

factorial(3)

return value

?

function argument


$i = 3$

Stack frame

factorial(1)

return value

1



function argument

i = 1

factorial(2)

return value

?

function argument

i = 2

factorial(3)

return value

?

function argument

i = 3

Stack frame

Stack frame


Stack frame

factorial(2)

return value  
function argument

factorial(3)


return value  
function argument



factorial(3)

return value

6



function argument

i = 3


Stack frame

`factorial(3)`

result: 6

The heap

- The *heap* is the general pool of computer memory.
- Memory is allocated on the heap using *malloc* / *calloc*
- The *heap* memory must be **explicitly** freed using *free*.
- Failure to do so -> **memory leak**



```
void foo2() {  
    int *arr;  
    /* allocate memory on the heap */  
    arr = (int *)calloc(6, sizeof(int));  
    /* do something with arr */  
} /* arr is NOT deallocated */
```

```
void foo3() {  
    int *arr;  
    /* allocate memory on the heap */  
    arr = (int *)calloc(6, sizeof(int));  
    /* do something with arr */  
    free(arr);  
}
```



foo2 and foo3

Local variables

int \*arr

Stack frame

stack

foo2 and foo3

Local variables

arr = 0x1234

Stack frame

stack



0x1234

arr[5]

arr[4]

arr[3]


arr[2]

arr[1]

arr[0]

heap

After **foo2** exits,  
**without** freeing  
memory



stack




**Memory  
Leak**

0x1234



heap


After **foo3** exits,  
**with** freeing  
memory



stack



0x1234



heap



- **Memory leaks** are one of the **worst** kinds of bugs
  - Often, no harm done at all
  - Eventually may cause a long-running program to crash
    - Out of memory
  - Very hard to track down
- Special tools (e.g. **valgrind**) exist to debug memory leaks

# Summary

# Today

- Storage Classes

# Today

- Storage Classes
  - auto, external, static, register variables



# Today

- Storage Classes
- Dynamic Memory Allocation

# Today

- Storage Classes
- Dynamic Memory Allocation
  - `void *malloc(size_t size);`
  - `void *calloc(size_t nmemb, size_t size);`
  - `void free(pointer);`

# Today

- Storage Classes
- Dynamic Memory Allocation
  - `void *malloc(size_t size);`
  - `void *calloc(size_t nmemb, size_t size);`
  - `void free(pointer);`
  - `void *realloc(pointer, int new_size);`  
increases or decreases the size of the specified block of memory. Reallocates it if needed.

# Today

- Storage Classes
- Dynamic Memory Allocation
- Stack and Heap
  - How Stack and Heap work
  - How to prevent memory leaks

# Today

- Storage Classes
- Dynamic Memory Allocation
- Stack and Heap
  - How Stack and Heap work
  - How to prevent memory leaks
  - `valgrind` -> to debug memory leaks

# Finally

- This is the last lecture for COMP281 this year, though the practical classes will run next week.
- Thank you very much for your hard work.
- I hope you have enjoyed COMP281

Finally

- This is the last lecture for COMP281 this year, though the practical classes will run next week.
- Thank you very much for your hard work.
- I hope you have enjoyed COMP281

Bye, for now...