

# Logistic regression

Procheta Sen

# Probabilistic vs “ordinary” classifier

- **“Ordinary” classifier** is a function  $f$  that assigns to an input object  $\bar{X}$  a predicted class  $c$  from a fixed set of classes  $\{c_1, c_2, \dots, c_k\}$ , i.e.

$$c = f(\bar{X}).$$

- **Probabilistic classifier** is a conditional distribution  $P(C | \bar{X})$ . For an input object  $\bar{X}$  it gives probabilities  $p_1, p_2, \dots, p_k$ , where

$$p_i = P(c_i | \bar{X})$$

and  $p_1 + p_2 + \dots + p_k = 1$ .

# Two types of models: discriminative vs generative

## Discriminative

- Assume that the conditional distribution  $P(C | X)$  (i.e. the probabilistic classifier) has specific form  $P_{\theta}(C | X)$  depending on some parameters  $\theta = (\theta_1, \dots, \theta_k)$
- Use training data set to **find** / **learn** parameters  $\theta_1, \dots, \theta_k$  such that the resulting distribution is “best possible” among all distributions of the assumed form

## Generative

- Assume that data come from specific distribution  $P_{\theta}(X, C)$  depending on some parameters  $\theta = (\theta_1, \dots, \theta_k)$
- Use training data set to **find** / **learn** parameters  $\theta_1, \dots, \theta_k$  such that the resulting distribution is “best possible” among all distributions of the assumed form
- Use  $P_{\theta}(X, C)$  to classify new objects

# Probabilistic classifiers

## Generative

- Naive Bayes 
$$P(H|E) = \frac{P(E, H)}{P(E)} = \frac{P(E|H)P(H)}{P(E)}$$

...

## Discriminative

- **Logistic regression** (today!)
- Multilayer perceptrons (neural networks)

...

# Setup

- We consider the binary classification problems with classes  $\{-1, +1\}$
- We want to build a probabilistic classifier that outputs the probability of a particular training instance  $\bar{X}$  being positive ( $y = +1$ ) or negative ( $y = -1$ )

# Logistic regression: main idea

- Define a separating hyperplane  $H$  **parameterised** by the feature weights  $\bar{W} = (w_1, \dots, w_d)$  and a bias parameter  $b$ , i.e.

$$H = \left\{ b + \sum_{i=1}^d w_i x_i = 0 \mid x_1, \dots, x_d \right\}$$

- **In perceptron**, to classify an input object  $\bar{X} = (x_1, \dots, x_d)$ , we used only the sign of

$$b + \bar{W}^T \bar{X} = b + \sum_{i=1}^d w_i x_i,$$

which tells in which of the two half-spaces created by the hyperplane the point is located.

# Logistic regression: main idea

- However, the actual value of  $b + \bar{W}^T \bar{X}$  conveys extra useful information: it is proportional to the distance from point  $\bar{X}$  to the hyperplane  $H$
- **Logistic regression** uses
  - sign of  $b + \bar{W}^T \bar{X}$  to **classify object**  $\bar{X}$  and
  - $|b + \bar{W}^T \bar{X}|$  to **quantify our confidence** in this classification: the larger the value, the further  $\bar{X}$  from the separating hyperplane  $H$

# Logistic regression: main idea

$b + \sum w_i x_i$  is

- 1) positive
- 2) proportional to the distance


Positive class

Negative class

Hyperplane defined by  $b + \sum w_i x_i = 0$

$b + \sum w_i x_i$  is

- 1) negative
- 2) proportional to the distance






# Logistic regression: main idea

- To interpret the confidence score  $b + \bar{W}^T \bar{X} \in [-\infty, \infty]$  as probability we would like to transform it to a value in the interval  $[0, 1]$  so that  $-\infty$  maps to 0 and  $\infty$  maps to 1.
- A function that does this is the **logistic sigmoid function**.

# Logistic sigmoid function



# Logistic sigmoid function properties

- $\sigma(x) \in [0,1]$  for any  $x \in [-\infty, \infty]$
- $1 - \sigma(x) = \sigma(-x)$
- $\frac{\partial \sigma}{\partial x} = \sigma(x) \cdot (1 - \sigma(x))$

# Logistic regression: discriminative classifier

## Discriminative classifier

- Assume that the conditional distribution  $P(C | X)$  (i.e. the probabilistic classifier) has specific form  $P_{\theta}(C | X)$  depending on some parameters  $\theta = (\theta_1, \dots, \theta_k)$
- Use training data set to **find** / **learn** parameters  $\theta_1, \dots, \theta_k$  such that the resulting distribution is “best possible” among all distributions of the assumed form

# Logistic regression: model assumption

- For an object  $\bar{X} = (x_1, \dots, x_d)$ , the probability that  $\bar{X}$  belongs to the positive class is modelled as

$$P(y = +1 \mid \bar{X}) = \sigma(a) = \frac{1}{1 + e^{-a}},$$

where  $a = b + \bar{W}^T \bar{X}$ . Hence the probability that  $\bar{X}$  belongs to the negative class is

$$P(y = -1 \mid \bar{X}) = 1 - P(y = +1 \mid \bar{X}) = 1 - \sigma(a) = \sigma(-a) = \frac{1}{1 + e^a}$$

It is convenient to write  $P(y = t \mid \bar{X}) = \sigma(t \cdot a) = \frac{1}{1 + e^{-t \cdot a}}$ , where  $t \in \{-1, +1\}$

# Logistic regression: choosing/fitting parameters

- Let  $\mathcal{D} = \{(\bar{X}_1, y_1), (\bar{X}_2, y_2), \dots, (\bar{X}_n, y_n)\}$  be the training data set
- Using the **maximum likelihood estimation method** we would like to find parameters  $b, w_1, w_2, \dots, w_d$  that maximise the likelihood function

$$\ell(b, w_1, w_2, \dots, w_d, \mathcal{D}) = \prod_{i=1}^n \sigma(y_i(b + \bar{W}^T \bar{X}_i))$$

- This is equivalent to minimising  $-\ell$  or minimising the negative log-likelihood function

$$-\ell \ell = -\log \ell = -\sum_{i=1}^n \log \sigma(y_i(b + \bar{W}^T \bar{X}_i))$$

# Logistic regression: choosing/fitting parameters

$$-\ell\ell = -\log \ell = -\sum_{i=1}^n \log \sigma(y_i(b + \bar{W}^T \bar{X}_i))$$

$$\nabla_{b, w_1, \dots, w_d}(-\ell\ell) = -\nabla_{b, w_1, \dots, w_d} \ell\ell$$

Denote  $a_i = b + \bar{W}^T \bar{X}_i$

We need to compute  $\frac{\partial \ell\ell}{\partial b}$  and  $\frac{\partial \ell\ell}{\partial w_k}$  for every  $k = 1, \dots, d$ .

# Logistic regression: choosing/fitting parameters

$$\ell \ell = \log \ell = \sum_{i=1}^n \log \sigma \left( y_i (b + \overline{W}^T \overline{X}_i) \right) = \sum_{i=1}^n \log \sigma \left( y_i \cdot a_i \right)$$

$$\frac{\partial \ell \ell}{\partial b} =$$

Logistic sigmoid function properties

1)  $\frac{\partial \sigma}{\partial x} = \sigma(x) \cdot (1 - \sigma(x))$

2)  $1 - \sigma(x) = \sigma(-x)$



# Logistic regression: choosing/fitting parameters

$$\ell \ell = \log \ell = \sum_{i=1}^n \log \sigma \left( y_i (b + \overline{W}^T \overline{X}_i) \right) = \sum_{i=1}^n \log \sigma \left( y_i \cdot a_i \right)$$

$$\frac{\partial \ell \ell}{\partial b} = \sum_{i=1}^n y_i \cdot \sigma \left( - y_i (b + \overline{W}^T \overline{X}_i) \right) = \sum_{i=1}^n y_i \cdot \sigma \left( - y_i \cdot a_i \right)$$

Logistic sigmoid function properties

1)  $\frac{\partial \sigma}{\partial x} = \sigma(x) \cdot (1 - \sigma(x))$

2)  $1 - \sigma(x) = \sigma(-x)$

# Logistic regression: choosing/fitting parameters

$$\ell \ell = \log \ell = \sum_{i=1}^n \log \sigma \left( y_i (b + \overline{W}^T \overline{X}_i) \right) = \sum_{i=1}^n \log \sigma \left( y_i \cdot a_i \right)$$

$$\frac{\partial \ell \ell}{\partial w_k} =$$

Logistic sigmoid function properties

$$1) \frac{\partial \sigma}{\partial x} = \sigma(x) \cdot (1 - \sigma(x))$$

$$2) 1 - \sigma(x) = \sigma(-x)$$

# Logistic regression: choosing/fitting parameters

$$\ell \ell = \log \ell = \sum_{i=1}^n \log \sigma \left( y_i (b + \overline{W}^T \overline{X}_i) \right) = \sum_{i=1}^n \log \sigma \left( y_i \cdot a_i \right)$$

$$\frac{\partial \ell \ell}{\partial w_k} = \sum_{i=1}^n y_i \cdot \sigma \left( - y_i (b + \overline{W}^T \overline{X}_i) \right) \cdot x_k^{(i)} = \sum_{i=1}^n y_i \cdot \sigma \left( - y_i \cdot a_i \right) \cdot x_k^{(i)},$$

where  $\overline{X}_i = (x_1^{(i)}, x_2^{(i)}, \dots, x_d^{(i)})$

# Logistic regression: choosing/fitting parameters

$$\ell \ell = \log \ell = \sum_{i=1}^n \log \sigma(y_i(b + \bar{W}^T \bar{X}_i)) = \sum_{i=1}^n \log \sigma(y_i \cdot a_i)$$

Interpretation of  $\sum_{i=1}^n y_i \cdot \sigma(-y_i \cdot a_i)$

- If  $y_i = +1$ , then  $\sigma(-y_i \cdot a_i) = \sigma(-a_i) = 1 - \sigma(a_i) = P(y = -1 \mid \bar{X}_i)$
- If  $y_i = -1$ , then  $\sigma(-y_i \cdot a_i) = \sigma(a_i) = P(y = +1 \mid \bar{X}_i)$
- Hence  $\sigma(-y_i \cdot a_i)$  is the probability of misclassifying the training object  $X_i$

$$\sum_{i=1}^n y_i \cdot \sigma(-y_i \cdot a_i) = \sum_{\bar{X}_i \in \mathcal{D}_+} P(y = -1 \mid \bar{X}_i) - \sum_{\bar{X}_i \in \mathcal{D}_-} P(y = +1 \mid \bar{X}_i)$$

# Logistic regression: update rule

**Gradient Descent** method for finding local minimum of  $f(\bar{Z}) = f(z_1, \dots, z_d)$

1. Pick an initial point  $\bar{Z}_0$
2. Iterate according to

$$\bar{Z}_{i+1} = \bar{Z}_i - \gamma_i \cdot ((\nabla_{\bar{Z}} f)(\bar{Z}_i))$$

where  $\gamma_1, \gamma_2, \dots$ , are step-sizes.

# Logistic regression: update rule

$$\text{minimize } -\ell\ell = -\log \ell = -\sum_{i=1}^n \log \sigma(y_i(b + \bar{W}^T \bar{X}_i))$$

$$\frac{\partial \ell\ell}{\partial b} = \sum_{i=1}^n y_i \cdot \sigma(-y_i \cdot a_i)$$

$$\frac{\partial \ell\ell}{\partial w_k} = \sum_{i=1}^n y_i \cdot \sigma(-y_i \cdot a_i) \cdot x_k^{(i)}, k = 1, \dots, d$$

Hence we have the following update rule

$$b \leftarrow b + \mu \cdot \sum_{i=1}^n y_i \cdot \sigma(-y_i \cdot a_i)$$

uses whole training set

$$\bar{W} \leftarrow \bar{W} + \mu \cdot \sum_{i=1}^n y_i \cdot \sigma(-y_i \cdot a_i) \bar{X}_i$$

# Online vs Batch

## Batch

- Uses the **entire training dataset** in every iteration to update the weight vector
- Popular optimisation algorithm for the batch learning of logistic regression is the Limited Memory BFGS (L-BFGS) algorithm
- Batch version is slow compared to the online version. But shows slightly improved accuracies in many cases

## Online

- Uses only **one training object** in every iteration to update the weight vector
- The Stochastic Gradient Descent algorithm (SGD)
- SGD version can require multiple iterations over the dataset before it converges (if ever)
- SGD is a technique that is frequently used with large scale machine learning tasks (even when the objective function is non-convex)

# Logistic regression online algorithm (Stochastic Gradient Descent)

**LogisticRegression**(Training data:  $\{(\bar{X}_1, y_1), \dots, (\bar{X}_n, y_n)\}$ , Learning rate  $\mu$ , **MaxIter**)

1:  $w_i = 0$  for all  $i = 1, \dots, d$ ;

2:  $b = 0$

3: **for** iter = 1 ... **MaxIter** **do**


4:   **for** i = 1 ... n **do**

5:      $a_i = b + \bar{W}^T \bar{X}_i$

6:      $w_s = w_s + \mu \cdot y_i \cdot \sigma(-y_i \cdot a_i) \cdot x_s^{(i)}$ , for all  $s = 1, \dots, d$

7:      $b = b + \mu \cdot y_i \cdot \sigma(-y_i \cdot a_i)$

8: **return**  $b, w_1, w_2, \dots, w_d$

$$\begin{aligned}\bar{W} &\leftarrow \bar{W} + \mu \cdot y_i \cdot \sigma(-y_i \cdot a_i) \bar{X}_i \\ b &\leftarrow b + \mu \cdot y_i \cdot \sigma(-y_i \cdot a_i)\end{aligned}$$




# Logistic regression prediction

**LogisticRegressionTest**( $b, w_1, w_2, \dots, w_d, \bar{X}$ )

1:  $a = b + \bar{W}^T \bar{X}$

2: if  $a > 0$  then

3:   predicted label =  $+1$  #positive class

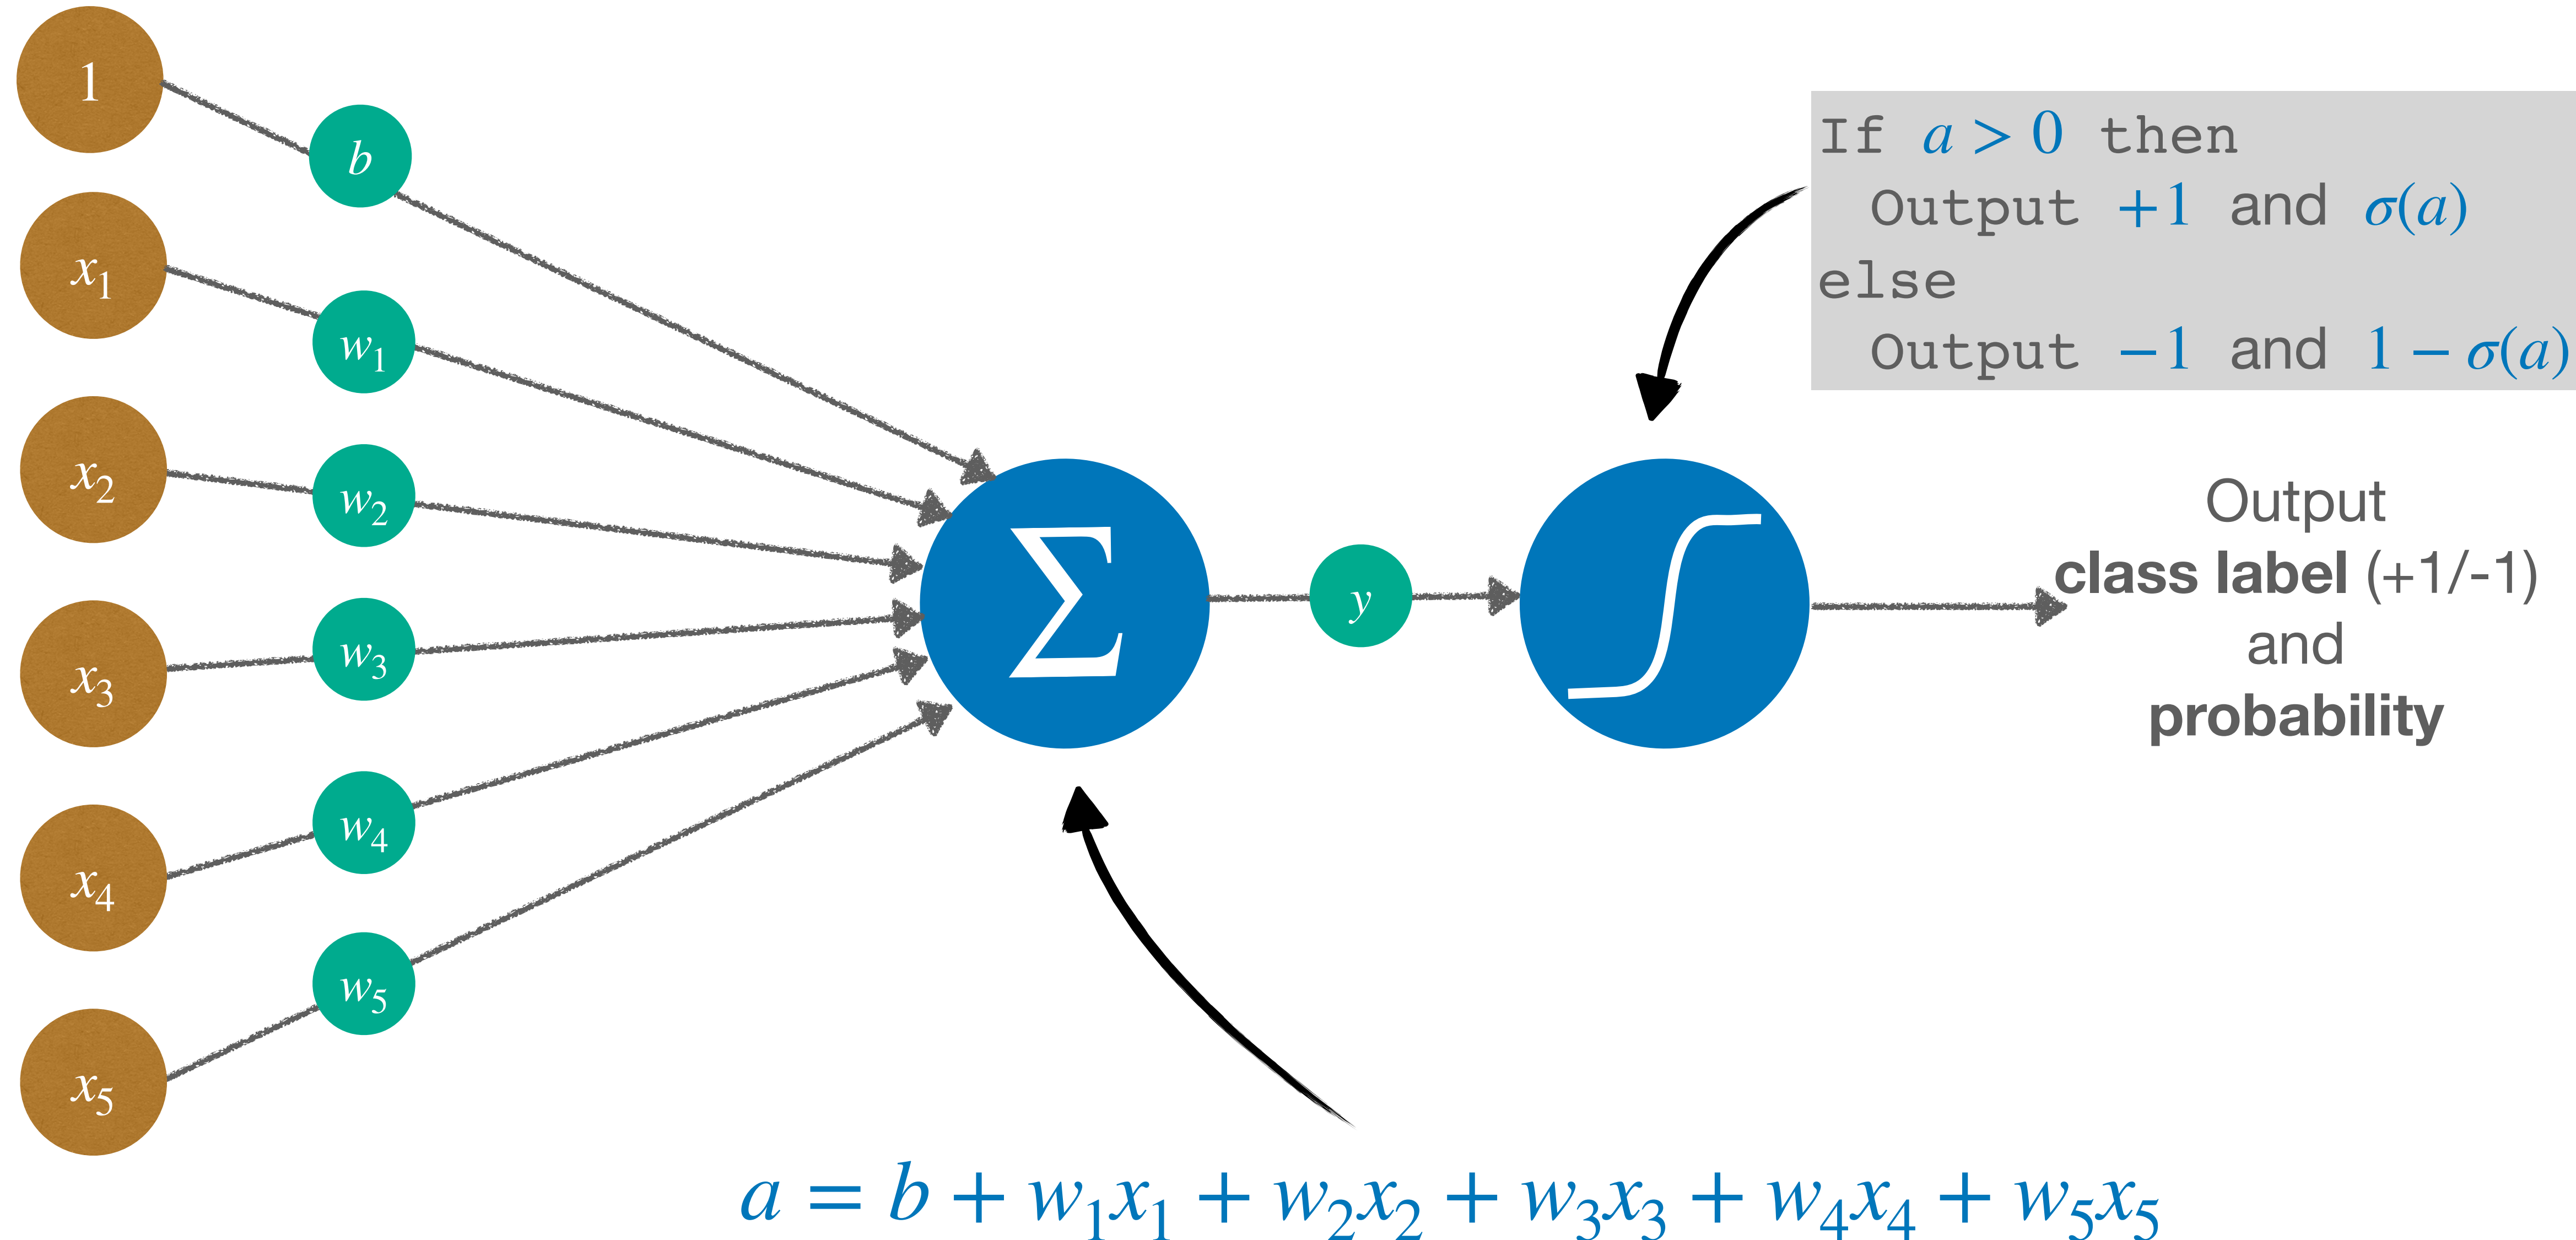
4:   probability that  $\bar{X}$  belongs to the positive class =  $\sigma(a)$  #confidence

5: else

6:   predicted label =  $-1$  #negative class

7:   probability that  $\bar{X}$  belongs to the negative class =  $1 - \sigma(a)$  #confidence

# Logistic regression: neuron interpretation



# L2 regularisation

- Let us denote by  $L(\mathcal{D}, \bar{W})$  the Loss of classifying a dataset  $\mathcal{D}$  using a model represented by a weight vector  $\bar{W}$
- We would like to impose L2 regularisation on  $\bar{W}$ .
- The overall objective to minimise can then be written as follows

$$J(\mathcal{D}, \bar{W}) = L(\mathcal{D}, \bar{W}) + \lambda ||\bar{W}||_2^2 = L(\mathcal{D}, \bar{W}) + \lambda \sum_{i=1}^d w_i^2.$$

Here  $\lambda$  is called the **regularisation coefficient** and is usually set via **cross-validation**.

- The gradient of the overall objective simply becomes the addition of the loss-gradient and the scaled weight vector  $\bar{W}$ .

$$\nabla_{\bar{W}} J(\mathcal{D}, \bar{W}) = \nabla_{\bar{W}} L(\mathcal{D}, \bar{W}) + 2\lambda \bar{W}$$

# L2 regularisation in logistic regression

L2 regularised logistic regression update rule for training object  $(\bar{X}, y)$  with  $a = b + \bar{W}^T X$

No regularisation:  $\bar{W} \leftarrow \bar{W} + \mu \cdot y \cdot \sigma(-y \cdot a) \cdot \bar{X}$

With regularisation:

$$\bar{W} \leftarrow \bar{W} - \mu \cdot \left( -y \cdot \sigma(-y \cdot a) \cdot \bar{X} + 2\lambda \bar{W} \right)$$

$$= \bar{W} + \mu \cdot y \cdot \sigma(-y \cdot a) \cdot \bar{X} - 2\mu\lambda \bar{W}$$

$$= (1 - 2\mu\lambda) \bar{W} + \mu \cdot y \cdot \sigma(-y \cdot a) \cdot \bar{X}$$